

[v9 ai followup](#) | [v8 remediation report](#) | [README](#) | [ai provider and council](#) | [numerical models and growth](#) | [game backend and state machines](#) | [system data and evidence](#) | [gaps and next iteration](#) | [reproducibility](#)

Offline reading copy. Scientific figures, rendered state diagrams and available rendered equations are embedded. External source/evidence links require GitHub access.

V9 AI follow-up: semantic context and bounded answers

Date: 11 September 2026 UTC
State: candidate; immutable publication and actual-provider verification pending
Scope: AI context/validation corrections after the immutable V8 public trial
Data: synthetic demonstration; actual farm operations disabled

Abstract

The V8 public trial showed why successful HTTP transport, schema validation and exact reference membership are insufficient evidence of answer quality. Its planning Council completed seven roles but sometimes confused metric meanings, omitted role-relevant facts or misstated the selection policy. Its conversational invite stopped when a return turn exceeded the existing content limit. This follow-up addresses the observed causes: opaque aliases that removed semantic names, an incomplete comparison-field filter, unbalanced context admission, missing numerical evidence and verbose answers. It retains all failed V8 evidence and leaves that edition unchanged. Numerical planning, synthetic cohorts and the recorded execution engine retain the V8 implementation and its separate evidence.

1. Evidence and hypotheses


The V8 public postmortem is the failure record. The planning mission used seven actual requests and passed its mechanical gate; the conversation used four requests and remained partial. The combined scorer evaluated ten completed cases: six passed and four failed. Exact workflow completeness also failed because the invited return and conversational Council were absent. AI-assisted examination identified additional qualitative problems; that examination was not human expert agronomic review.

Observation	Identified mechanism	Bounded correction
Market prose called margin a booked-value total	Mission aliases replaced canonical semantic reference names	Retain exact aliases for output selection while supplying canonical references, semantic labels, units, strategy identity and role purpose
Weather and Supply discussed general feasibility without useful role evidence	Generic strategy context could substitute for role context	Role-specific context admission; numerical roles require typed facts; explicit unavailable-source state requires abstention

Observation	Identified mechanism	Bounded correction
Chair described withholding while the numerical selection was accepted	Conditional generic policy prose was ambiguous and the Council was not given the mission's required/advisory policy	Pass the actual mission policy and expose structured deterministic selection semantics
Numerical direct reply cited controls but no typed result	A qualitative context reference alone satisfied the local evidence gate	Require typed evidence for a non-abstaining numerical interpretation when relevant facts are supplied
Unrelated comparison metrics occupied the context	The filter recognized .metrics. but missed .baseline_metrics. and .scenario_metrics.	Recognize all comparison segments and apply the role metric set consistently
Relevant scenario values and deltas could be crowded out	Alphabetical reference order favored an early policy and segment	Interleave metric/policy/segment records under the existing bound; prioritize controls and relevant source status
Invited return exceeded 400 characters	Guidance encouraged explanation near the hard bound	Target fewer than 220 content characters, while retaining the 400-character schema limit and finite shared repairs

These are testable software hypotheses. They do not claim that arbitrary free-form prose is entailed by its references, nor that repeated success on this fixture establishes provider reliability.

2. Revised evidence path



Mission prompt/context is V6; conversation prompt/context is V5. The common validator is V4. Output schema V3 keeps the same 400-character content limit, three typed references, three qualitative references, one evidence reference, one highlight and one proposed action.

Reference maps and provider-returned aliases remain auditable; unknown aliases fail closed. Public facts use canonical references, not private model lookup aliases.

The actual mission `council_policy` now crosses the application/Council boundary. Required and advisory policy must remain distinct. The Council does not gain permission to modify quantities, set a new ranking rule, run the numerical model, change farm state or authorize operations. Only an explicit application action can create a synthetic execution world from an accepted numerical result.

Conversation context includes server-authored response requirements. A numerical non-abstention must select a supplied typed fact; a typed fact accidentally placed in qualitative references remains eligible for the existing bounded format-only repair, and the final response must still place it correctly. An explicit absence of relevant frozen weather or community-market observations requires abstention. General news does not become observed price, demand or weather evidence merely because a feed is connected. Arbitrary unanswerable questions may still abstain.

The context remains bounded to 48 ordinary typed facts, 24 ordinary qualitative references, 16 ordinary prior turns and six evidence records, with declared research/control and reply-target exceptions and a hard 120,000-character prompt ceiling. Balanced admission improves coverage within these bounds; it does not provide exhaustive context or prove that omitted records are irrelevant.

3. Validation design

Deterministic tests cover comparison-field exclusion, baseline/scenario/delta triplets, required typed evidence, absence-driven abstention, exact alias mapping, role evidence and required/advisory policy. Existing tests retain tenant isolation, receipt replay, request accounting, preserved rejected attempts and format-repair semantics. Tests use isolated PostgreSQL databases where required; background application workers must never share a regression database.

The conversation focus passed 43 tests. The final mission-focused and whole-repository results will be attached after the candidate is frozen. The earlier V8 full regression passed 589 tests with one skip and remains evidence for its recorded source scope. V9 does not claim new numerical accuracy from those software tests.

Actual-provider verification is a single bounded sequence on the new immutable edition: planning Council (normally seven requests, at most nine), direct/invite/ conversational Council (normally ten, at most fourteen), and a frozen research adviser (normally one, at most two). Stored replay and quality scoring make no provider calls. Failures are retained; neither runtime limits nor validators are relaxed to obtain a passing answer. Changes after publication require another immutable edition.

4. Accounting and result interpretation

The task ledger records 55 actual requests through V8, including failures and repairs. A revised internal experiment ceiling of 96 permits at most 25 further requests for this V9 sequence, for a reserved maximum total of 80. The application's shared production limit remains 48 requests per day. An internal reservation is not an actual call and is not a change to the production budget.

A PASS_AUTOMATED quality result covers the specified roles, counts, persisted validation, required topic/reference groups and known contradiction patterns. It cannot establish the meaning of every sentence. AI-assisted case review is labelled separately; human expert and real-farm outcome evidence remain absent. A pass does not promote fitted models, make synthetic values observed, or permit actual farm operations.

5. Immutable publication and final results

Publication, preservation, final live output and quality evidence are pending in this candidate. The deployed V8 and its failed quality evidence remain available and unchanged. The final report must attach the actual V9 source/image identity, all-edition health, prior-edition preservation, workflow counts and per-case result before describing this candidate as verified.

FarmTact V8 remediation: methods, implementation evidence, and validation limits

Audit date: 11 September 2026 UTC

System state: immutable V8 published; all edition health and prior-state preservation checks passed; V9 AI correction under development

Data status: synthetic demonstration only

Operational status: real farm operations disabled

Live Council quality status: **FAILED / remediation in progress**

Abstract

This report evaluates the V8 remediation of FarmTact as a software and numerical system, rather than as a validated agricultural decision system. The evaluated implementation joins a deterministic demand baseline, a whole-bed planning model, a persistent synthetic execution world, bounded scenario and research workflows, and optional DeepSeek interpretation. Planned schedules can be advanced through a recorded civil-date clock with append-only events, order and lot attribution, inventory conservation, cash accounting, and future replanning. `daily-bed-cpsat-v3` adds absolute crop-cycle identities, collision-checked harvest-lot IDs, a persistent cross-replan origin map, exclusion of all historically executed cycle IDs, and resource reservations against the largest declared scenario yield. The data package preserves the original fixture hash while adding disjoint, versioned synthetic training and evaluation cohorts. Demand evaluation uses rolling origins and one-, two-, and four-week horizons; crop-cycle evaluation uses independent whole-batch outcomes. The current evaluator is `rolling-origin-demand-v4` and includes a behavioral counterfactual availability probe. No fitted candidate is eligible for production, and every real-data promotion gate remains blocked. The Council now selects server-owned typed facts and keeps execution, evidence, and decision influence as separate statuses. Conversation projection is V4 and mission-reference selection is V5: the model selects short lookup aliases that resolve exactly to canonical server-owned evidence. Role projection has explicit ranking and size bounds, frozen research inputs, abstention rules and preserved rejected attempts. The isolated PostgreSQL regression passed 589 tests with one skip; a subsequent focused AI suite covers the narrow V5 alias change. A 56-day network execution trial survived an actual application restart, and a separate 84-day regression exercised two replans and cross-segment lot attribution. The clean-checkout zero-inference pipeline and V3 synthetic numerical report passed. The final local research adviser passed reference validation using 6,704 prompt tokens, compared with 64,756 in the earlier retained request. Forty-four actual provider requests have been consumed across the retained experiments. The pre-context-fix scorer passed 8 of 18 cases and remains historical failed evidence. V8 was published from `a00e546b1270a5c532e8eae6b8b64e568c344c13`; all edition health and prior-state preservation checks passed. The public mission passed its mechanical gate using

seven calls, but a public invited conversation stopped after its return turn exceeded the response limit. The public combined scorer passed six of ten completed cases and failed workflow completeness; no public conversation Council ran. Total consumption through V8 was 55 requests. These failures motivated the separate V9 follow-up; mechanical correctness and exact references do not establish qualitative truth.

1. Questions and claim boundary

The audit addressed six questions.

1. Can a frozen accepted mission become a persistent, inspectable synthetic world?
2. Do planning, replay, and execution share explicit schedule and inventory rules?
3. Can the model evaluation detect useful implementation differences without representing synthetic scores as farm accuracy?
4. Are model-generated statements traceable to immutable numerical facts while qualitative interpretation remains visibly unverified?
5. Do retries, cancellation, concurrency, and tenant isolation preserve one authoritative state transition?
6. Which report gaps are resolved in code, and which require live or real-world evidence before release or model promotion?

The strongest supported claim is an engineering claim: FarmTact V8 can calculate and persist bounded synthetic planning experiments with explicit provenance and replay controls. The evidence does not support a biological efficacy claim, a forecast-accuracy claim for a real farm, an economic-return claim, or permission to act on a farm. The system itself labels these boundaries in the private data contract, simulation service, model evaluation report, and Council remediation report.

2. Evidence hierarchy and methods

2.1 Evidence hierarchy

The audit used the following order of authority.

Level	Evidence	Interpretation
1	Current executable source and strict schemas	Defines implemented behavior
2	Machine-readable generated reports with hashes and versions	Defines the recorded result of a bounded run
3	Focused test logs and retained browser artifacts	Supports a particular tested path and environment
4	Specialist narrative reports	Explains methods and known limitations

Level	Evidence	Interpretation
5	Specifications and earlier technical chapters	Defines intended behavior and historical context
6	Provider announcements	Defines external model naming and migration facts only

Conflicts were resolved in favor of current source and newer generated artifacts. Historical reports were retained rather than rewritten. The machine-readable gap register was `in_progress` when inspected and was being updated concurrently, so this report states evidence dispositions without overriding its eventual integration status; see `gap-status.json`.

2.2 Static inspection

Static inspection covered the strict farm and outcome contracts, synthetic generators, feature construction, evaluation code, planner, accounting replay, simulation routes, SQL-backed stores, scenario and conversation workers, Council orchestration, provider gateway, client mutation transport, and the recorded-world UI. Primary entry points are listed in the reproducibility appendix. The review checked version constants, bounded collection sizes, tenant predicates, state-transition guards, frozen-input hashes, cancellation boundaries, and the absence of an automatic public-data-to-planner edge.

2.3 Dynamic evidence

Dynamic evidence came from existing local artifacts and focused test reports. No provider request was made for this document. The synthetic model report was regenerated with versioned seeds and contains its own report ID, manifest hashes, evaluation versions, and timestamp. The browser report was generated against a temporary isolated PostgreSQL database with the normal worker active and Council inference disabled. The concurrency reports used PostgreSQL rather than treating SQLite behavior as production evidence. Exact commands and the limits of each run appear in Section 12.


2.4 Status vocabulary

“Implemented” means a code path and contract exist. “Focused pass” means the cited bounded test completed successfully. “Synthetic result” means every target or execution outcome was generated by code. “Blocked” means a required external or operational condition is absent. “Pending” means the integration owner has not yet produced the final result. “Unsupported” means local evidence validation rejected a model statement even if the provider returned a syntactically usable response.

3. V8 system architecture

The browser speaks same-origin JSON and event-stream APIs to FastAPI. PostgreSQL holds tenant-owned farm versions, missions, branches, conversations, research sessions, simulation worlds, events, and mutation receipts. An in-process worker claims durable numerical and

provider jobs. Local Python owns forecasting, optimization, replay, state transition, and factual rendering. DeepSeek is called only for explicit interpretation paths.



The implementation preserves the design separation described in `system-data-and-evidence.md`: tenant farm state and explicit scenario controls can alter numerical inputs; public weather, news, and trade records cannot do so automatically. `public_features_used` is empty in both current numerical evaluation artifacts.

4. Persistent synthetic execution

4.1 Admission and world creation

A world can be created only from the current tenant's mission whose status is `ACCEPTED_FOR_SIMULATION` and whose input version is still current. The route rejects a stale mission, an unaccepted mission, overdue historical harvests without execution records, or a ninth world in the same session. Creation chooses the accepted strategy and freezes a central execution scenario with yield and demand factors equal to one. It stores the segment farm, allocations, deterministic replay trace, opening lots, opening cash, exact decimal revenue and cost accumulators, a civil-date horizon, and `real_operations_enabled=false`. The implementation is in `simulation.py`.

The public world contains its revision, clock, bed states, inventory, task IDs, totals, cash, event sequence, plan history, runtime provenance, and state hash. Its stage representation comes from the shared `schedule-state-v2` helper, not from a second UI-specific growth rule. The helper explicitly declares that progress is elapsed schedule fraction and is not physiological biomass.

4.2 Canonical schedule semantics

The schedule states are empty, nursery, growing, ready, harvested, and sanitation. An executed harvest requires a recorded task; a snapshot does not infer that a past scheduled harvest occurred. Projected views may assume scheduled tasks execute. A bed is reserved through the harvest date plus `sanitation_days`, inclusive, and becomes available on the following civil date. The API exposes the same recipe calendar metadata to the client, including nursery, grow, sanitation, shelf-life, and growth-model version fields. The UI preview range ends at `horizon_days - 1`, so 7-, 56-, and 84-day horizons end at indices 6, 55, and 83.

4.3 Daily state transition

For each included day, execution verifies that persisted opening lot mass equals the frozen trace's opening mass. It records due sow, transplant, harvest, and sanitation-complete tasks once. It appends order-service events from exact order allocations. Revenue accumulates only when a delivered line has a non-null effective price. Input, labour, packing, and disposal costs are accumulated with `Decimal`, and public money values round at the cent boundary. The closing lot snapshot becomes the next day's inventory. The following invariant is checked for each day:

$$[M_{\text{open}} + M_{\text{harvest}} - M_{\text{delivered}} - M_{\text{disposed}} = M_{\text{close}}.]$$


The cumulative cash identity is:

$$[C_t = C_0 + \sum_{d \leq t} \text{Revenue}_d - \sum_{d \leq t} \text{Cost}_d.]$$

The daily event and world update share the store's ambient transaction, which is important because a multi-day advance must commit all included days or none. The world revision advances once per mutation, while the event sequence advances once per appended event.

4.4 Replanning

Replanning is admitted after at least one executed day and only when seven or more unexecuted days remain. The route allows at most twelve replans per world. It carries forward inventory, future orders, bounded nonnegative cash, and growing allocations whose sanitation occupancy reaches the new segment. Those allocations enter the planner as executed locks and cannot be silently removed or charged again for already completed sowing inputs. New candidate sowing cannot precede the new segment start. The local calculation runs outside the tenant lock; commit reacquires the lock and compares the original revision. If another mutation has advanced the world, the computed continuation is discarded. Only a feasible Balanced continuation replaces the future segment. The prior plan hashes and through-date remain in `plan_history`.



4.5 Absolute cycle identity and lot provenance

Final review found that horizon-relative allocation IDs could refer to different absolute crop cycles after a replan. This could suppress a later task while its numerical trace still produced harvest. The V3 planner now derives allocation identity from bed, crop, recipe and absolute sow/transplant/harvest dates. Replanning also excludes every previously executed allocation ID from new candidates, covering caller-controlled imported batch IDs while preserving existing locks.

Harvest lots use a separate deterministic hash and collision ordinal, checked against opening and earlier generated lots. The recorded world retains a lot-to-allocation origin map across replans. Harvest events expose produced lot IDs, and subsequent demand-service events carry their source allocation even when stock crosses a planning segment. These fields are provenance, not additional yield. The new 84-day ASGI HTTP regression replans after days 7 and 42 and checks cycle identity, harvest tasks and cross-segment lot attribution. Its initial failure is preserved in the discovery record.

4.6 Execution evidence and limits


The dedicated PostgreSQL test issued six simultaneous same-key advances and found one committed civil day and one matching receipt. Two different keys competing at the next revision produced one success and one 409 conflict. Historical receipt replay returned the original response without changing the newer world. A second tenant received 404 for the world, events, and mutation route. A reconstructed store observed the same world and ordered events. This focused test passed once in 7.38 seconds; details are in `execution-validation.md`.

The archived intermediate 56-day HTTP trial in `v8_execution_trial.py` passed across an application restart and made zero provider calls. Its immutable intermediate artifact records the exact synthetic counts and invariants it checked. The separate final network restart trial passed with 56 task events and 64 demand-service events: `execution-trial.json`. The cash constraint remains narrower than the execution ledger: planning constrains new input cost, while execution deducts input, labour, packing, and disposal. A conserved ledger can consequently have a negative balance. That semantic gap must be resolved before cash is described as a full liquidity constraint.

5. Planner and accounting remediation

5.1 Weighted declared scenarios

`daily-bed-cpsat-v3` accepts an optional declared scenario set while preserving the original three-scenario default. The planner rejects empty sets, duplicate IDs, more than 100 rows, invalid factors, negative or non-finite weights, and non-positive total weight. Raw weights are normalized and converted to 1,000 deterministic integer objective units, with at least one unit for every positive weight. Returned scenario records expose raw, normalized, and integer weights. This is an engineering distribution chosen by the caller, not a calibrated probability model.



A constructed one-bed test changes the chosen action when the quiet/busy weights are reversed from 99/1 to 1/99. This behavioral check supports use of weights in the objective rather than merely their presence in metadata. The objective charges scenario-independent new-work cost once. Harvest-labour and conservative cash reservations use the maximum yield factor in the actual declared scenario set; custom factors above the default high-yield case cannot silently bypass those resource checks. The reported maximum quantization scale is approximately 0.001 of total objective weight.

5.2 Policy objectives and proof language

Lean and Balanced maximize weighted policy utility. Each scenario exposes revenue, packing cost, disposal cost, shortfall mass and coefficient, disposed mass and coefficient, terminal stock and coefficient, and new-work commitment cost. The resulting unit is described as SGD-equivalent policy score because penalty coefficients encode preferences; it is not literal expected profit.

Resilient first maximizes the minimum aggregate fill rate across all declared scenarios to one basis-point resolution. Only an OPTIMAL primary result freezes the proven floor for a secondary weighted utility solve. A merely FEASIBLE primary result is an incumbent and cannot support a claim of optimal downside protection. If the secondary stage times out, the proven primary schedule can be retained with an unfinished tie-break label. This is aggregate scenario service

protection; it is not a per-order guarantee, chance constraint, CVaR estimate, or empirical risk distribution.

5.3 FEFO, order priority, and missing price

The CP model uses remaining-life buckets keyed to civil expiry dates. Replay allocates the earliest expiry first, then harvest date, then lot ID. For common shelf life this resembles FIFO; differing explicit expiry dates make the implemented policy FEFO. A later-expiring bucket cannot be consumed while an earlier bucket is available. The test space enumerated 81 small two-lot/two-day combinations and reconciled the allocator with the best feasible delivery/disposal outcome.

Demand service is sequential: booked commitments precede inferred residual demand, known higher prices precede lower prices within kind, and stable IDs break ties. Confirmed lines expand to current order IDs after cancellation. `order_allocations` records requested, delivered, and shortfall mass plus each lot's ID, quantity, harvest date, and expiry date. Residual lines use stable synthetic IDs. The schema lacks buyer priority, grade, and lot/customer eligibility, so the model makes no such allocation claim.

Forecast contract 2.0 retains numeric `price_sgd_per_kg=0` for compatibility when no booked price exists, but adds `price_status=unavailable_no_booked_price`. The planner converts that state to a null effective price and records unpriced requested and delivered mass. Dates with booked prices use `booked_weighted_average`. Missing price therefore does not become evidence of observed zero revenue.

5.4 Terminal stock and horizon effects

Closing inventory remains usable or expiring stock at horizon close. It has zero declared salvage value and is not counted as waste merely because the horizon stops. Lean, Balanced, and Resilient apply provisional closing-stock penalties of 0.50, 0.20, and 0.05 SGD-equivalent per kilogram respectively. A later horizon can convert that same mass into actual expiry disposal. These penalties require observed storage, spoilage, future-order, and disposal economics before operational interpretation. Implementation and test evidence are summarized in `planner.md`, with source in `engine.py` and `accounting.py`.


6. Data lineage and reproducible fixtures

6.1 Registry coverage

Coverage is derived from the current crop and evidence registries rather than hard-coded report prose. The validated registry contains 12 crop profiles and 24 evidence documents, plus registry versions and source-file SHA-256 digests. A contradictory manifest fails validation. The older 8 September data-quality artifact remains historical and was not edited to imitate a current rebuild.

The clean-checkout fixture bundle contains six small project-authored CC0-1.0 provider-contract payloads. Its manifest fixes byte counts and SHA-256 values and rejects unsafe paths or changed content before normalization. It produced 24 normalized rows: six weather

observations, seventeen forecast rows, and one trade-contract row. Installation and rebuild produced identical dataset IDs, row counts, and normalized hashes in `data_pipeline_verification.json`. The bundle is labeled `synthetic_contract_fixture`. Archived SingStat and NASA response bodies were not redistributed because lawful reuse had not been established.



6.2 Compatibility fixture

The original `synthetic-farm-v1` remains byte-for-byte compatible. Its contract hash is `00286779541fa9ed1245ea2aff65d08c1a05241cdc6a399706aa6601fdb3f76`. It remains useful for UI, numerical, and regression checks. It is not used as the sole evidence for V8 model behavior.

7. Synthetic model evaluation

7.1 Cohort design

Training and evaluation use different generator names, seeds, entity IDs, farms, buyers, batches, and date ranges.

Partition	Version	Seed	Rows	Design
Demand training	<code>synthetic-demand-train-v2</code>	731104	7,488	6 farms, 18 buyers, 104 weeks from 2023-01-02
Demand evaluation	<code>synthetic-demand-evaluation-v2</code>	982451	3,072	4 unseen farms, 12 unseen buyers, 64 weeks from 2025-01-06
Cycle training	<code>synthetic-crop-cycle-train-v1</code>	445901	576	6 farms × 24 cycles × 4 crops
Cycle evaluation	<code>synthetic-crop-cycle-evaluation-v1</code>	771283	256	4 unseen farms × 16 cycles × 4 crops

Demand generation includes stable, growth, and compression regimes; a 13-week pattern; bounded heterogeneity; surge and drop shocks; booking lead times; cancellations; and separate booking, cancellation, and outcome availability. Crop generation uses independent whole-batch observations with planned and actual stage dates, area, recipe baseline, measurement method, operating regime, delay shocks, and quality shocks. It does not use weather inputs or contain a weather-response coefficient.




Figure 1. Synthetic cohort design generated from the current evaluation artifact. Distinct identifiers and time ranges test pipeline separation, but both partitions remain products of declared code and can share generator-design bias.

7.2 Rolling-origin demand method

rolling-origin-demand-v4 evaluates farm/crop and buyer/crop series independently. For each forecast origin it supplies only dependencies whose availability timestamp does not exceed that origin. It evaluates one-, two-, and four-week horizons against last-week, 13-week seasonal-naive, fixed alpha-0.35 EWMA, and training-tuned EWMA baselines. Metrics are MAE, RMSE, WAPE, and signed bias. MAE uncertainty uses a deterministic series bootstrap with 500 replicates.

Alpha tuning uses only the training partition and freezes a cutoff of 2025-01-06T12:00:00+00:00. Candidate alphas were 0.10, 0.20, 0.35, 0.50, 0.65, 0.80, and 0.90. The selected alpha was 0.80 based on one-step training MAE. The evaluation partition was not used for selection.

The initial dependency counter design had a structural weakness: it attempted to count late timestamps only after selecting rows already known to be on time. V4 now reports explicit included dependency auditing and adds an independent behavioral counterfactual. The counterfactual perturbs a permanently unavailable historical target and compares every resulting forecast for a held-out farm/crop and its buyers. It compared 492 farm forecasts and 1,476 buyer forecasts and changed zero. The probe status is passed. This supports the tested invariance; it does not prove the absence of every future feature-leakage defect.

7.3 Demand results

The held-out farm cohort contains 16 series and 656 predictions per model/horizon. The buyer cohort contains 48 series and 1,968 predictions per model/horizon.

Cohort	Horizon	Best WAPE	Model	Tuned EWMA WAPE
Farm	1 week	0.048919	tuned EWMA	0.048919
Farm	2 weeks	0.096992	fixed EWMA	0.100631

Cohort	Horizon	Best WAPE	Model	Tuned EWMA WAPE
Farm	4 weeks	0.121254	fixed EWMA	0.139717
Buyer	1 week	0.061945	fixed EWMA	0.065245
Buyer	2 weeks	0.093836	fixed EWMA	0.102266
Buyer	4 weeks	0.140546	fixed EWMA	0.163816

The tuned candidate wins only the farm one-week WAPE comparison. The fixed EWMA performs better in the five other cohort/horizon comparisons shown. This is a useful negative result: training selection does not justify a general promotion claim even inside the synthetic benchmark. Complete metrics and bootstrap intervals are preserved in `synthetic_model_evaluation.json`.

Synthetic demand holdout · generated outcomes, not farm accuracy




Figure 2. Current V4 rolling-origin synthetic demand results. The figure covers the two synthetic cohort levels and declared horizons; it contains no observed customer or farm demand.

7.4 Whole-batch crop-cycle method and result

The candidate adds crop-specific mean training residuals to recipe maturity and marketable-mass baselines. Features are crop ID, recipe cycle days, recipe marketable kilograms per square metre, and occupied area. Public weather, news, trade, and causal weather coefficients are excluded. The evaluation unit is an unseen farm's whole batch, not repeated measurements that could leak the same batch across splits.

Target	Recipe baseline MAE	Residual candidate MAE	Better result
Fresh marketable mass	3.722184 kg	3.907328 kg	Recipe baseline
Maturity	0.957031 days	1.035536 days	Recipe baseline

Training-residual 10th-90th percentile ranges covered 84.375% of held-out mass outcomes and 87.5% of held-out maturity outcomes. These are descriptive residual ranges and are not calibrated prediction intervals.




Figure 3. Independent synthetic crop-cycle evaluation. The fitted residual candidate fails to improve the primary MAE results and is not promoted.

7.5 Promotion decision

The policy is `farmtact-model-promotion-v1`. Schema and split integrity passed. The point-in-time gate passed based on zero included dependency violations, split availability assertions, and the counterfactual invariance described above. Three mandatory gates remain blocked:

- no authorized real farm/customer temporal holdout exists;
- no external farm cohort exists;
- no production drift, alert, rollback, or retraining trial exists.


The report therefore records `eligible_for_production=false` and status `blocked`. No fitted demand or crop-cycle candidate has been promoted. The forecast used by the application remains the declared engineering baseline `recipe-ewma-v1`.

8. Council evidence architecture and actual-call inventory

8.1 Exact roster and workflow

The current Council is `seven-agent-council-v3` with workflow `sequential_specialists_then_chair`. Six specialists independently receive role-projected slices of the same frozen context. Only the planning chair receives the prior claims, their evidence statuses, validation issues, and eligibility flags. The implementation does not claim a free-form debate or challenge round.

Persona	Role identifier	Bounded responsibility
Ravi	demand_analyst	Booked orders, forecast demand, required crop quantities and delivery dates
Hana	weather_analyst	Observed context and freshness; no unsupported yield adjustment
Idris	market_analyst	Prices and sourced reactions; reactions are not measured demand
Mei	production_analyst	Recipes, lead times, space, and modeled feasibility
Lina	supply_chain_analyst	Inventory, expiry, inputs, and delivery timing
Ben	profit_analyst	Computed costs, labour, cash, and margins under declared assumptions
Asha	planning_chair	Reconcile eligible claims and explain the deterministic selection policy



The normal mission budget is seven calls with up to two bounded structural repairs, for a maximum of nine. Direct and invitation workflows reserve their exact maxima separately. `RunBudget.cancel()` prevents subsequent transmission and stops bounded response consumption between transport chunks. It does not promise immediate interruption inside an already active synchronous HTTP call.

8.2 Typed facts and qualitative limits

V8 replaces numeric coincidence with selected `fact_refs`. The server builds a typed catalogue from frozen values and renders the value, unit, entity, period, and snapshot hash after model output. Model-authored quantities and dates are rejected even when they happen to match a value elsewhere in the context. `tool_result_refs` are restricted to qualitative context. Wrong entity, wrong unit, wrong meaning, citation laundering, and numeric prose have adversarial tests in `test_ai_injection_boundaries.py` and `test_ai_quality.py`.

`RESPONSE_LIMITS` is the shared prompt/schema authority: 400 content characters, one evidence reference, three context references, three typed facts, one highlight, and one proposed action. Conversation prompt/projection is V4; mission prompt/projection is V5. Mission outputs select exact short F001/C001 aliases, resolved locally to canonical references with both forms preserved

in the audit. Unknown aliases are rejected; alias resolution does not weaken numeric, membership or evidence checks. Role projection admits at most 48 typed facts, 24 qualitative references, 16 prior turns and six evidence records, subject to a hard 120,000-character serialized prompt limit. Exact research inputs referenced by the frozen result are included even when ordinary rank caps have been reached. Projection truncation is a declared limitation, and questions that cannot be answered from admitted context must receive an abstention. Contract versions identify prompt, output schema, validator, frozen context, and source context. New records include a public-context SHA-256 value. Earlier stored records remain readable through conservative additive defaults.

A structural format repair preserves the rejected attempt as an auditable record. Its instruction must preserve lexical meaning, relationships, proposed actions and evidence/highlight references; it may only remove or relocate an invalid number or date form. This correction can improve contract conformance but cannot retroactively make the rejected attempt valid or prove factual quality.

Execution status answers whether a provider workflow ran. Evidence status answers whether the returned references and typed facts passed the local checks. Decision influence answers whether policy can consider the result or must withhold. A successful HTTP response can therefore be completed, unsupported_all, and advisory_only at the same time. Qualitative interpretation remains qualitative_unverified; selecting a correct fact reference does not verify causal or agronomic prose.

8.3 Provider caller inventory

The fail-closed inventory in `deepseek_runtime.json` contains four callers.

Caller	Integration	Routes	Maximum	Failure semantics
Mission Council	active product	seven Council roles	9	visible withheld or partial
Persistent conversation	active product	direct, invite, Council, research interpretation	9	visible unsupported, partial, or failed
Synthetic label observer	active API	visual observer	1	visible missing visual evidence
Authenticated gateway trial	diagnostic only	demand, production, visual observer	16	preserved failed trial

The gateway requires the compiled inventory and configuration to agree, restricts models to the reviewed set, verifies that the returned model ID exactly matches the requested ID, bounds request/response/image sizes and timeouts, limits concurrency to one, and records token categories where the provider supplies them. Helper or evaluator names do not imply an active caller.

8.4 Model migration

On 10 September 2026, DeepSeek announced V4.1-Flash, stated that it was live on the API with native multimodal support, instructed API users to select deepseek-flash, and retired the previous V4 Flash and Vision Experimental names; see the official DeepSeek announcement. The runtime's canonical model is consequently deepseek-flash for every current text and native-vision route. The discovery evidence is retained in model-discovery.json. This migration statement establishes naming and routing policy; it does not prove FarmTact output quality.

8.5 Retained actual-provider evidence

The evolving experiment ledger preserves all attempts and reconciles actual paid requests; see live-call-ledger.json.

Experiment	Requests	Observed result	Audit interpretation
Initial local planning Council	1	retired-model identity mismatch	Failed closed; run retained as incomplete
Canonical model with vision	9	native-vision path passed; five claims unsupported	Transport/capability evidence only; Council quality failed
V3 context/schema run	7	six valid outputs; one response rejected for the word three	Improved contract adherence, still incomplete
V3 format-repair run	7	six valid outputs; Supply Chain selected one unknown reference	Reference membership rejected locally; mission withheld

The four mission counts sum to 24. Before the final V4 mission run, the ledger additionally records 11 V4 conversation requests and one V4 research request, for 36 actual paid requests. Conversation inputs totalled 805,878 prompt tokens and produced ten persisted adviser messages; the research request used 64,756 prompt tokens. These are usage measurements, not a quality score or a recommended production budget. The first mission failure demonstrates why exact returned-model validation matters. The canonical vision run shows that native visual input traversed the route, but five unsupported outputs prevent a quality pass. The next two runs each produced six locally valid role outputs and still failed complete coverage.

The retained pre-context-fix scorer covers 18 cases: eight PASS and ten FAIL, so its overall automated status is **FAIL**. Exact workflow-integrity counts passed for direct, invite, Council, planning and research sequences. Those count checks show that the intended routes ran exactly once in the tested workflow; they do not establish entailment or usefulness. Conversation V4 and mission V5 address context admission, exact aliases and repair semantics, but no final live evidence yet proves that it fixes quality.

Historical combined status: FAILED. Final V8 public quality status: FAILED. No deployed-quality claim is made. The retained planning artifacts are planning-local.json, planning-local-canonical.json, planning-local-v3.json, and planning-local-v4.json. The final

mission budget reserves at most nine requests; unused reservation is not counted as consumption.

9. Scenarios, conversations, and research state machines

9.1 Scenario attempts

Scenario roots are immutable tenant-owned numerical experiments. Each scenario keeps at most three append-only attempt summaries with queued, started, completed, cancellation, terminal status, and safe exception type. Retries preserve the frozen farm, baseline root, model versions, prior attempts, and original run idempotency key. A changed or stale key receives 409. An unsupported frozen version fails rather than being silently recomputed by the current model.

Cancellation atomically moves DRAFT, QUEUED, or RUNNING to CANCELLED. Queued cancelled jobs cannot be claimed. Because CP-SAT is synchronous, a running solve is observed as cancelled after the current local solve boundary; the compare-and-write guard prevents its stale result from overwriting the terminal cancellation. Scenario calculation makes no provider request.

Lists use SQL ordering and bounded `limit=1..30` with an exclusive before cursor. The thirty-first persisted scenario is rejected before numerical work. Idempotency lookup precedes the quota check so receipt replay still works at the limit.

9.2 Conversations

Conversations are independently capped at 30 per tenant and 120 messages per conversation. Admission rechecks capacity under the conversation lock, including the user message and every requested advisor turn. Queued and running requests can be cancelled through an atomic store primitive. Council execution checks cancellation before numerical context assembly, budget reservation, each initial or repair provider call, message persistence, and final completion. Cancellation is cooperative at call boundaries.


Replay preserves the original workflow type, model-call status, evidence state, decision influence, contract versions, and calculation provenance. It does not claim a new calculation or provider call. Historical records missing V8 fields receive conservative compatibility values.

9.3 Scripted research and challenge

Council Research is explicitly `scripted_research` with `inference_origin=local_rules` until a user requests a distinct advisor interpretation. Its dialogue can propose bounded reservations, unconfirmed-order treatment, and a 50–100% labour ceiling. Weather challenges state that outdoor rainfall is context rather than a numerical yield input for the sheltered synthetic farm. Farm time does not advance while research calculations run.

Each numerical result is tied to the research schema, forecast version, planner version, input version, and content hash. Revision history stores snapshots before mutations and exposes bounded incremental history pages. `cancel_calculation` and `retry_calculation` are separate from `stop`, which stops scripted dialogue only. A failed or cancelled current calculation can be

retried up to three attempts; queued and running states cannot be duplicated. Actions require a current revision and an idempotency key.



The research workflow supports comparison and provenance, but it does not establish that a proposed control is causally effective. Qualitative suggestions remain subject to the same local evidence labels when a provider is explicitly invoked.

10. Concurrency, tenancy, and idempotency

10.1 Server transaction model

The core store hashes the anonymous session bearer before persistence and scopes farm, mission, event, receipt, scenario, conversation, research, and simulation queries by tenant. PostgreSQL tenant transactions lock the tenant row to serialize state allocation. An ambient connection context lets nested event, world, and receipt writes join the same transaction. SQLite uses a serialization lock only as a test adapter; it is not the production concurrency claim.

Dedicated PostgreSQL tests used six concurrent same-key scenario retries after a forced transient failure and observed exactly one second attempt. Six concurrent scenario cancellations and six conversation-request cancellations each produced exactly one transition. The simulation concurrency result is described in Section 4.6. One pre-existing shared-database test observed a queued job become RUNNING because another live worker claimed it; the isolated V8 PostgreSQL test passed. This is why the final suite should run without unrelated workers polling the same test queue.

A later focused integration group passed 10 tests covering the SQLite background-worker serialization adapter and both queued and running research-calculation cancellation paths. This result supports those regressions in the test adapter; the production concurrency evidence remains the isolated PostgreSQL runs.

10.2 Client retry identity

The shared client mutation transport writes an idempotency key before transmission. Its session-storage identity includes immutable edition, HTTP method, API path, and serialized body. The key survives navigation and reload after a network error, 5xx, 408, 425, or 429 because the server may have committed even when the client lacks a response. Definitive success and definitive 4xx responses clear the pending operation. After a mutation or receipt replay the client fetches the current world, because an old exact receipt intentionally contains an older revision. Implementation is in `api.ts` and `research.ts`.

The browser harness injected a 503 followed by a 429 and then success across page reloads. All three requests carried the same key. A later successful operation cleared it, and a deliberate 409 import response also cleared its separate pending key. This test exercises uncertainty semantics rather than merely checking storage keys.

10.3 Retention mechanism and remaining persistence limits

Anonymous-session authentication expires after 24 hours. An explicit operator CLI can identify inactive tenants and delete tenant-owned rows in dependency order within one transaction. It defaults to dry-run and 30 retained days, rejects fewer than seven retained days, considers at most 500 tenants per invocation, and skips any tenant with queued or running mission, scenario, conversation or research work. It covers the registered schema while preserving shared budgets and security controls; 12 isolated tests cover dry-run, full child deletion, recent/active retention and bounds. The command has no automatic schedule and Store initialization may create missing schema objects even in dry-run mode, so operators must target the intended database and inspect its JSON report.

Cancellation cannot preempt a CP-SAT solve or abort an already active synchronous provider call immediately. Concurrent identical replans may compute the same local solution more than once, although revision and receipt serialization allow only one committed state. There is not yet an approved production retention schedule or operating history for the CLI.

11. Browser behavior and visual evidence

The Outcomes room distinguishes a static schedule preview from recorded synthetic execution. The recorded panel can create a world from the accepted current mission, advance one or seven days, display clock and revision, show bed stage and progress, list task and demand events, show inventory and delivery/harvest/cash totals, and replan future work when at least seven days remain. It states that inference was not triggered and real operations are disabled.

Capability presentation separates configured routes, historical probe evidence, current verification, and last observed tenant execution. Planning does not treat “credential configured” as “provider verified.” Council cards separately expose execution, evidence, and decision

influence, and label qualitative prose unverified. Council planning remains the primary action, while numerical-only planning is an explicit alternative and was used for this browser run.




Figure 4. Mobile browser evidence at 390 × 844 CSS pixels. The Playwright run used an isolated local PostgreSQL database, normal background worker, numerical-only mission, one-day advance, future replan, then seven-day advance. It checked the visible execution totals, event rendering, shared sanitation calendar, mission persistence across refresh, and absence of body overflow. It made zero Council inference requests.




Figure 5. Desktop browser evidence at 1280 × 900 CSS pixels. This screenshot comes from the same numerical-only recorded world and the same 27-check run. The overflow assertion was repeated at 360, 390, 430, and 1280 CSS pixels. The screenshot does not show a live provider result or a real farm event.

The machine-readable browser report records 27 of 27 checks passing. It also records expected local /api/releases 404 responses plus deliberately injected 503, 429, and 409 responses. Those responses were excluded from the “unexpected failures” assertion because they are the

intended retry-contract stimuli. The visible synthetic totals after the tested interactions were 54 kg inventory, 136.004 kg delivered, 190 kg harvested, 2,805*cash*, 1,030 revenue, and \$425 cost. These values belong only to that deterministic test world.

12. Verification results

12.1 Focused test evidence

Subsystem	Command or artifact	Recorded result	Scope limit
Data/model/planner	focused pytest group in reports/v8/data-ml.md	74 passed in 33.38 s	Synthetic and unit/integration scope
Planner plus adversarial contracts	focused group in reports/v8/planner.md	36 passed in 61.05 s	Includes deterministic fallback paths
Council/gateway/offline quality	package group in reports/v8/council.md	227 passed, 1 skipped, 1 PG test deselected	No live-provider quality pass
Council frozen subset	package group in reports/v8/council.md	96 passed in 35.54 s	Bounded gate/gateway/cancel subset
Scenario reliability	V8 reliability unit suite	5 passed in 3.30 s	Local service mechanics
Scenario/conversation concurrency	isolated PostgreSQL test	1 passed in 2.81 s	Dedicated database/worker conditions
Scenario combined group	gameplay reliability group	34 passed in 86.55 s	Focused, not whole repository
Simulation execution	execution unit suite	3 passed in 7.76 s	Synthetic world behavior
Simulation concurrency	isolated PostgreSQL test	1 passed in 7.38 s	Same-key, competing-key, tenant, restart-store
Worker/research cancellation	latest focused integration group	10 passed	SQLite worker serialization plus queued/running cancellation
Root identity/retention/cancellation focus	final-root-focused.xml	23 passed in 20.369 s	Includes 84-day ASGI world and replans on days 7 and 42
Final V8 AI implementation focus	focused pytest group	250 passed, 1 skipped, 1 PostgreSQL case deselected	Covers mission V5/conversation V4; no live-quality pass
Clean checkout	clean-checkout.json	PASS, zero inference calls	Rebuild/install/fixture/API path in recorded isolated environment

Subsystem	Command or artifact	Recorded result	Scope limit
Numerical evaluation	numerical_evaluation.json	PASS, planner V3	Synthetic-only; no fitted candidate promoted
Intermediate restart execution	archive/execution-trial.intermediate.json	PASS: 56 days, one restart, zero provider calls	Superseded by the final network PASS below
Web build	TypeScript no-emit plus Vite	PASS	Build only
Browser	v8_simulation_ui.mjs	27/27 PASS	Local PG, zero inference, exact journey above
Intermediate full regression	archive/full-regression.intermediate.xml	556 passed, 1 skipped in 398.002 s	Predates final identity/context changes
Final full regression	full-regression.xml	589 passed, 1 skipped in 437.76 s	Predates only the narrow V5 alias and trial-harness changes, covered by focused regressions
Final network execution	execution-trial.json	PASS: 56 days, one restart, zero provider calls	Exact receipts and cash/mass/order-lot reconciliation
Final research adviser	research-advisor-final.json	PASS: references verified, one actual call	Exact references do not prove all qualitative prose

These results are not additive because test sets can overlap. They should not be summed into a repository pass count.

12.2 Repository-wide gate

An earlier repository-wide run reported 499 passing tests and 34 failures while multiple packages were still editing generated manifests. The integration review attributed 32 failures to concurrent manifest churn and identified two actual defects, which were fixed. The next run completed with 557 collected tests: 556 passed, one skipped, zero failures and zero errors in 398.002 seconds. That result is preserved as intermediate evidence, because planner identity, retention and V4 context changes followed it.

The later focused root result records 23 of 23 tests passing in 20.369 seconds. It includes the 84-day ASGI world, replans on days 7 and 42, cycle identity, lot origins, tenant retention, cancellation and SQLite worker serialization. The clean-checkout result separately records a PASS with zero inference calls. These results narrow specific risks but do not replace the final repository-wide run.

Final repository-wide regression gate: PASS. The isolated-PostgreSQL run completed with 589 passed and one skipped in 437.76 seconds: full-regression.xml. The subsequent narrow V5 alias and trial-harness changes have separate focused evidence. Release and live AI quality still require their own recorded checks.

12.3 Test interpretation

Unit tests support schema boundaries and small deterministic cases. Exhaustive tiny planner tests support equivalence only inside their integer-gram, single-grade, single-harvest domain. PostgreSQL tests support the exact concurrent schedules they execute, not arbitrary distributed failure histories. The browser test supports the specified sizes, controls, and local topology. Actual-provider runs support observed transport and validation outcomes but do not replace human agronomic quality assessment. Synthetic model metrics compare implementations under generator assumptions and cannot estimate field accuracy.

13. Gap dispositions

The authoritative machine-readable register was being revised during report preparation. The following table records evidence disposition for this report without changing or predicting that register's final status.

Gap family	Implemented evidence	Required residual evidence
DATA-01	Dynamic validation derives 12 crops and 24 documents	Maintain registry/source review on future changes
DATA-03	Versioned report 2.0 metadata and archived pre-V8 report	Clean final regeneration after frozen revision
DATA-04	Self-contained hashed CC0 synthetic source fixtures	Authorized archival store for any real source snapshot
DATA-05	Empty public feature set; no weather/news/trade coupling	Pre-registered real-data ablation before adding any edge
DATA-06	Typed whole-batch outcome records and split checks	Authorized real farm batches and measurement protocol
NUM-03	Weights affect objective behavior and are reported	Calibrated scenario probabilities if probability language is desired
NUM-04	Resilient lexicographic maximin with proof status	Stakeholder-approved service definition and risk calibration
NUM-05	FEFO solver/replay alignment and order/lot attribution	Grades, buyer rules, multi-cut and eligibility schema
NUM-06	Terminal stock kept distinct from waste	Observed salvage, spoilage, disposal, and post-horizon economics
NUM-07	Explicit missing-price state and unpriced accounting	Current observed prices and commercial validation
NUM-08	Rolling-origin multi-horizon V4 evaluator	Real temporal and external-farm holdouts

Gap family	Implemented evidence	Required residual evidence
NUM-09	Promotion gates block synthetic candidates	Monitoring, rollback, drift, and approved thresholds
GAME-08	Bounded attempt history, cancellation, retry, pagination and active-job retention exclusions	Wider soak testing and approved operational retention schedule
GAME-11	Tenant locks, same-key receipts, revision conflicts and explicit bounded retention CLI	Multi-process fault injection, production monitoring and operator execution evidence
GAME-12	Exact computed order delivery impacts	Buyer/grade/customer-priority contract expansion
Execution world	Durable clock, events, ledgers, replanning	Completed restart HTTP artifact and negative-cash resolution
AI-01/02	Typed facts and independent local quality dimensions	Successful full live quality suite plus expert scoring
AI-03/12	Execution/evidence/influence split	Validate comprehension with planned human study
AI-04	Configured, historical, and observed statuses separated	Fresh successful explicit execution after quality fixes
AI-05/06	Fail-closed caller inventory and bounded budgets	Dated cost controls and deployment reconciliation
AI-07	Canonical harvest references and value status	Real measured outcome provenance
AI-08	Truthful six-specialists-then-chair workflow	Measure incremental utility before adding challenge calls
AI-09/10/11	Shared limits, rejection projection, version/hash provenance	Complete seven-role live conformance result
UX-01	Recorded/static distinction and responsive controls	Planned human usability/accessibility study
OPS-01	Local PG and build evidence	Final suite, immutable edition, deployment and post-deploy checks

14. Scientific and operational limitations

14.1 Synthetic data validity

Every demand and crop-cycle target in V8 evaluation is generated. Independent seeds and cohorts prevent direct identity overlap, but they do not remove shared design assumptions. A

model structurally similar to the generator can look favorable for the wrong reason. No score in this report estimates performance on an operating farm.

14.2 Biological scope

The schedule helper is a civil-date state machine, not a physiological growth model. Marketable mass uses recipes and declared scenario factors rather than measured light, nutrient, temperature, pest, cultivar, or operator effects. Crop intervals are descriptive synthetic residual ranges. No weather causal coefficient exists.

14.3 Demand and commercial scope

The private schema lacks buyer identity on orders, substitution, stockout-censoring state, grade, contract priority, and realized loss records. The demand forecast therefore cannot be described as buyer behavior modeling for a real operation. Missing current prices are made visible, but the system has no real pricing feed. Policy utility is not expected profit.

14.4 Optimization scope

The planner uses whole beds, integer grams, one grade, single-harvest lots, and declared scenario factors. It lacks split-bed, multi-cut, water, energy, packaging-capacity, rotation, substitution, partner-supply, and lot/customer eligibility constraints. Bounded CP-SAT FEASIBLE status does not prove global optimality. Scenario weights are not probabilities.

14.5 Execution scope

Execution replays the deterministic plan and central synthetic scenario. Events do not actuate equipment, assign a worker, send a buyer message, purchase an input, or record an observed harvest. The cash hard constraint does not yet cover every execution cost. The final 56-day HTTP protocol passed with a real application restart, exact historical receipts and reconciled cash/mass/order-lot totals. It remains synthetic engineering evidence.

14.6 AI scope

Typed rendering prevents a model from altering selected values, but it cannot make qualitative interpretation true. The current live quality status is failed: the pre-context-fix scorer passed 8 of 18 cases. Vision transport success is not agronomic visual validation. V4 role projection can omit relevant information under its 48/24/16/6 and 120,000-character bounds; it therefore requires abstention when admitted context is insufficient and must not be described as universal context coverage. Forty-four actual requests were consumed through the final local research trial. The final local mission still failed one reference check; the subsequent V5 alias change awaits public verification.

14.7 Operations and privacy scope

The session is an anonymous tenant, not an authenticated farm user. The explicit retention CLI supplies a bounded deletion mechanism, but it has no automatic schedule, approved production cadence or production execution evidence. Focused local PostgreSQL concurrency does not prove behavior under every network, process, database failover, or deployment condition. The

final repository suite and network restart trial passed; deployment and final public AI verification remain pending.

15. Evidence required before stronger claims

Real demand evaluation requires authorized order-line history containing farm, buyer/SKU, booking, modification, cancellation, due, delivered, stockout, substitution, and availability timestamps. The evaluation protocol must freeze farm and buyer sampling, temporal splits, censoring policy, metrics, acceptance thresholds, and subgroup checks in advance.

Real crop evaluation requires independent batch records containing recipe and cultivar versions, sow/transplant/harvest and sanitation events, occupied area and system, fresh harvested mass, marketable mass by grade, packout, losses, measurement method, and availability timestamp. The current outcome collection protocol is a protocol proposal, not a collected dataset.

Public weather or market context may enter a numerical model only after units, farm exposure, spatial mapping, and point-in-time availability are resolved and a pre-registered ablation improves an authorized real holdout. That analysis must distinguish prediction from causation.

AI promotion requires a complete current-model trial across direct, invitation, Council, vision where applicable, and research interpretation workflows. The retained report must include every request, repair, usage record, local issue, and failed output. Agronomic usefulness and harm must then receive independent human scoring under the human study protocol.

Operational release requires a successful final build and regression, edition immutability checks, database migration checks, a final network restart trial, deployment health checks and a documented rollback target. Passing technical gates advance autonomously; there is no human release-approval checkpoint in this development phase. None of these steps enables real farm operations under the present policy.

16. Reproducibility map

Purpose	Primary source or artifact
Private contracts and recipe calendar	contracts.py
Canonical civil-date stages	growth.py
Compatibility and benchmark generators	fixtures.py
Public fixture installation and validation	packages/ingestion
Feature/forecast contract build	build_features.py
Data rebuild entry point	build_dataset.py
Demand and crop evaluation	evaluation.py

Purpose	Primary source or artifact
Synthetic evaluation command	evaluate_synthetic_models.py
Synthetic evaluation result	synthetic_model_evaluation.json
Numerical evaluation result	numerical_evaluation.json
Planner	engine.py
Replay accounting	accounting.py
Simulation worlds and events	simulation.py
Scenario state machine	scenarios.py
Core transactional store	store.py
Conversation persistence	conversation_store.py
Tenant retention mechanism	retention.py
Retention operator command	prune_expired_tenants.py
Research state machine/history	council_research.py
Council orchestration	council.py
Council roster and gate	agents.py
Typed fact contract	ai_contracts.py
Offline quality scorer	ai_quality.py
Provider policy and transport	deepseek_gateway.py
Provider caller manifest	deepseek_runtime.json
Client mutation transport	api.ts
Recorded execution panel	SimulationPanel.tsx
Browser method	v8_simulation_ui.mjs
Browser result	ui-browser.json
Intermediate restart execution	execution-trial.intermediate.json
Final restart execution target	execution-trial.json
Focused root regression	final-root-focused.xml
Clean-checkout result	clean-checkout.json
Intermediate full regression	full-regression.intermediate.xml

Purpose	Primary source or artifact
Final full regression target	full-regression.xml
Pre-context-fix AI score	ai-quality-before-context-fix.json
Execution concurrency method	test_v8_simulation_postgres.py
Scenario/conversation reliability report	game-reliability.md
UI evidence report	ui.md
Data/ML evidence report	data-ml.md
Planner evidence report	planner.md
Council evidence report	council.md
Execution evidence report	execution-validation.md

The numerical report used a deterministic zero-second fallback and records `daily-bed-cpsat-v3`, `recipe-ewma-v1`, fixture hash, settings, cutoff, source revision, dirty state, report schema, and public feature exclusions. Bounded CP-SAT wall time, objective, and best bound can vary with solver build and runtime when a positive time limit is used. The archived pre-V8 report remains at `numerical_evaluation.pre-v8.json`.

17. Conclusions and pending decisions

V8 materially improves the fidelity of FarmTact's synthetic planning experiment. The implementation now connects an accepted plan to a durable, revisioned execution world whose clock, tasks, lots, orders, costs, and replans can be inspected. Weighted objectives, maximin proof labels, FEFO replay, explicit missing-price state, terminal inventory and exact order attribution remove several important ambiguities from the numerical layer. Planner V3 further makes cycle identity absolute, collision-checks harvest lots, persists their origin across replans, excludes all historically executed IDs and reserves resources against the declared maximum yield.

The data and model work also improves the honesty of evaluation. Disjoint generators and cohorts, rolling origins, multiple horizons, whole-batch outcomes, fit cutoffs, and a counterfactual availability probe provide useful engineering tests. Their current result supports no fitted-model promotion, and the promotion policy correctly blocks production eligibility.

The Council architecture has stronger provenance and safer local validation. Its current actual-provider evidence nevertheless fails the quality gate. The canonical `deepseek-flash` migration, conversation V4 projection and mission V5 aliases are implemented, but the retained pre-fix score is only 8 of 18 cases passing. Exact workflow counts pass; answer quality remains unvalidated until the final live result is recorded.

The remaining integration decisions are evidence-driven and explicit:

- **PASS:** clean-checkout pipeline, V3 numerical report, 589-test PostgreSQL regression (one skip), focused AI/root regressions, final 56-day restart execution and final local research adviser;
- **ARCHIVED PASS:** intermediate 556-pass regression and intermediate 56-day restart-safe execution trial;
- **PENDING:** immutable V9 follow-up publication and live verification;
- **FAILED:** public V8 combined AI quality and exact workflow completeness (6 of 10 completed cases passed);
- **BLOCKED:** any production model promotion without authorized real outcomes;
- **DISABLED:** all real farm operations.

This report is therefore a reproducible V8 engineering assessment and limitation record, not a deployment certificate, agronomic validation, or operational authorization.

18. Final local integration results and release handoff

The final full PostgreSQL regression completed with **589 passed, one skipped, zero failed** in 437.76 seconds. The recorded 56-day HTTP world again passed after an actual application restart, with 56 day-close events, 56 task events, 64 demand-service events, a future replan, exact historical receipts and reconciled cash/mass/order-lot totals. The separate 84-day ASGI HTTP regression covers replans on days 7 and 42 and cross-segment lot attribution. These are synthetic engineering results.

The final research-adviser request passed `references_verified` with exact reservation dates, a bed-area fact and the unconfirmed-order context. Its prompt used 6,704 tokens, versus 64,756 in the earlier retained request. These are observed counts from separate bounded trial instances, not a cost or reliability benchmark. The final research report and read-only replay retain the response and evidence.

A fifth local mission consumed seven requests and again withheld acceptance because Supply Chain shortened a complex order reference. This repeatable failure motivated a structural V5 mission-reference change: the model selects short F001 and C001 lookup aliases; the server resolves them exactly to the original canonical references before applying the existing type and evidence checks. Unknown aliases fail closed. Audit rows preserve both the mapping and the returned aliases, and public claims retain canonical references. This does not prove qualitative entailment. The failed final local mission remains preserved.

The local total is 44 actual provider requests. The internal experiment ceiling was explicitly revised as defects were discovered; all revisions and request accounting remain in the ledger. Runtime and shared-production daily limits were not changed. A public mission and public direct/invite/Council trial remain to be run after immutable V8 publication. Their results must be reported separately from these local passes.

The full suite predates the narrow V5 alias change; the focused AI suite then passed 250 tests, with one skipped and one PostgreSQL-only case excluded because the full PostgreSQL run covers storage. Root also checks the alias and interrupted-harness regressions separately. Final publication and automated live-quality outcomes remain pending in this source snapshot.

19. Immutable V8 publication and public quality result

V8 is published from a00e546b1270a5c532e8eae6b8b64e568c344c13, image sha256:8ed6fc3b1a6632f88bbc2b6a20b200b44bba13d7583c70f40e1c31f85c1e8357. All eight health checks matched source pins and the prior-seven preservation audit passed. Concurrent v7/v8 numerical scenarios completed in 22.768 and 36.169 seconds; 70 timed requests had read p95 0.2392 seconds. Cross-edition cookie rejection and unchanged main-farm snapshots passed. Available memory after the trial was 2,325,268 KiB. This supports the tested light load, not general scalability: capacity evidence.

The public planning mission completed all roles, passed reference/schema checks, was accepted for synthetic simulation and passed its deterministic replan and zero-inference replay. It consumed seven requests with no repairs. AI-assisted inspection nevertheless found missing role relevance, metric mislabelling and a misleading policy explanation. These are qualitative defects that reference membership alone cannot detect.

The public conversation consumed four requests: a direct reply, a Demand invite and its schema repair, then a Production return rejected for excess content length after the shared repair allowance was spent. The last request is PARTIAL; the conversation Council was never started. The retained direct response also omitted typed facts. The combined public V8 scorer passed six of ten completed cases, failed four, and failed workflow completeness. The postmortem records exact outputs, usage, failed checks and the limits of the AI-assisted semantic review.

V8 remains immutable. The subsequent V9 follow-up corrects the identified context/response-contract defects and has its own source, image and live verification. The final-local observations in Section 18 are historical checkpoints; this section is the final V8 public outcome.

FarmTact: scientific implementation report

Historical baseline: This chapter records the pre-remediation audit of source a96025e and published v7. Its implementation findings and measured results describe that baseline. Read the V8 remediation report for simulation and synthetic model remediation, and the V9 AI follow-up for corrections prompted by the V8 public trial. Source links below are navigation aids into the maintained repository; use the frozen baseline to reproduce the original inspection.

AI orchestration, numerical planning, crop development and evidence boundaries

Audit date: **11 September 2026** · Source baseline: a96025e · Published v7 source: 6efbbfd · Study type: **source-code and engineering-evidence audit**, not an agronomic trial.

Original audit reading copies: PDF · offline HTML. The Markdown chapters below preserve the original audit with explicit current-status notices.

V8 follow-up: remediation report · combined PDF · offline HTML.

V9 follow-up: AI correction report. Final reading copies and live verification are pending.

Abstract

FarmTact combines a visual synthetic-farm application, local statistical baselines, a constrained scheduling solver, deterministic scenario simulation and optional pretrained DeepSeek interpretation. This report traces the implemented pathways rather than inferring capabilities from role names or design requirements. Demand uses an exponentially weighted moving average (EWMA) combined with booked orders. Crop development uses fixed recipe durations and marketable yield assumptions; the farm view derives progress from calendar dates. OR-Tools CP-SAT selects whole-bed planting candidates, and a local inventory simulator evaluates three fixed stress scenarios. Public observations, News and crop literature provide context without automatically becoming numerical forecast inputs.

Acceptance stores a projected outcome and worklist; no farm-time or task-execution engine advances sowing, harvesting or inventory state. The council has distinct planning, conversation and research-study implementations. Default v7 research dialogue is scripted; real numerical jobs and optional adviser calls are separately triggered. Actual LLM calls use a server-only DeepSeek gateway with role/model allowlists, typed output validation, finite reservations and recorded replay. No application-specific ML training, crop physiological model, disease classifier or autonomous farm execution is implemented. Archived tests support specific engineering properties. In particular, both recorded v7 adviser responses remain unsupported; those probes do not establish answer quality. The report provides equations, architecture figures, source traces, reproducibility guidance and a prioritized requirements register for the next immutable application edition.

Contents

1. Scope, method and evidence
2. System overview
3. AI provider and council mechanics
4. Numerical models and plant growth
5. Complete game backend and state machines
6. System, data lineage and evidence
7. Results and interpretation
8. Limitations and next-iteration requirements
9. Reproduction and verification

The linked chapters form one report. Each includes implementation references; the documentation index covers the wider repository.

1. Scope, method and evidence

1.1 Research questions

- Which user actions can cause provider inference, through which roles and models?
- How do the three council workflows obtain context, validate answers and persist results?
- What is calculated locally, and what does “plant growth” mean in the application?
- Which inputs are observed, synthetic, curated, inferred or simulated?
- What do the tests establish, and which requirements remain unimplemented or unvalidated?

1.2 Audit protocol

The audit inspected the committed source, runtime configuration, fixture generator, contracts, numerical solver/simulator, API workers, frontend state mappings, release manifests and dated reports. Three bounded specialist audits covered AI, numerics and system/data architecture; the integrating author reconciled findings into this report, the README and the specifications. These are code-review roles, not application council calls or human study participants.

Read-only GitHub and Fly inspection established repository access, source state, the active Machine and public edition registry. Offline tests and deterministic figure generation are recorded separately in the documentation audit record. No new metered adviser experiment, deployment, farm-state mutation or full browser usability study was performed for this documentation task. Existing reports were read as historical observations, not relabelled as newly reproduced results.

1.3 Evidence hierarchy used in this report

Label	Meaning	What it cannot establish
Source-inspected	Behavior traced to a function, contract, configuration or UI mapping	That every branch was executed successfully
Fresh offline check	Named command run during this audit against local code/fixtures/mocks	Current external account capability or production performance
Archived experiment	Dated committed report from an earlier execution	Present availability, a different source commit or unseen inputs
External documentation	Provider/tool documentation, with retrieval context	Account-level feature access or successful application integration
Requirement / recommendation	Desired future behavior with an acceptance criterion	Implemented or validated capability

A schema-valid response is not automatically factually supported. A solver's feasible result is not automatically globally optimal or commercially desirable. A repeatable synthetic experiment is not external validation against real crops, orders or human decision-making.

2. System overview

FarmTact: numerical authority and optional AI interpretation




Figure 1. Source audit, 11 September 2026. Solid arrows show implemented information flow; no physiological model is implied.

Figure 1. Source-derived information flow. Numerical tools create the authoritative quantities. Frozen results and bounded context reach optional DeepSeek interpretation. Local checks and the relevant backend decision path control simulation acceptance. Scripted dialogue is a distinct zero-inference interaction. Source: app.py, council.py, council_research.py and engine.py. PNG · SVG.

2.1 Mechanism classification

Mechanism	Inputs → outputs	Learned from FarmTact records?	Provider call?
Fixture generation	Fixed parameters → fictional orders/batches/history	No	No
EWMA	Available historical orders + bookings → weekly residual demand	Recursive statistical state; no trained regressor	No
Harvest calendar	Batch dates/recipe assumptions → harvest dates and marketable kg	No	No
Farm artwork	Snapshot planning date + batch dates → stage/progress	No	No
CP-SAT	Candidate schedules/resources → whole-bed allocation choices	No	No
Inventory simulation	Allocations + dated demand + stress cases → ledgers/metrics	No	No
Scripted research dialogue	Bounded intent/context → proposal or clarification	No	No
DeepSeek advisers	Frozen context + role prompt → typed interpretation	Pretrained external model; no app fine-tuning	Yes, explicit paths
Synthetic vision probe	Generated image → checked batch-label observation	No FarmTact training	Yes
News collector	Allowlisted RSS bytes → validated metadata cache	No	No LLM inference

A registered `satellite_visual_reviewer` or `test_evaluator` route grants a bounded capability to trusted server code; it does not establish a deployed satellite pipeline or an LLM evaluation service. See the call-site inventory in the AI chapter.

2.2 Plant development and temporal meaning

Plant development is a recipe calendar

Illustrative common sow date = day 0. Existing batches retain their recorded dates.




Figure 2. Continuous spans illustrate durations. Daily bed occupancy includes harvest + sanitation through that day. The earliest next transplant is harvest + sanitation + 1 day. Shelf-life expiration is a separate inventory rule.

Figure 2. The four fixture recipes shown on a common illustrative sow date. Nursery and grow-out durations determine harvest. Sanitation affects bed reuse; shelf life affects harvested inventory. Yield coefficients are **synthetic engineering assumptions**. The chart is not a measured biological growth curve. Source: fixtures.py, contracts.py and engine.py. PNG · SVG.

For a batch with sow date s , transplant date t , harvest date h and current snapshot date d , the backend view computes $\text{progress} = \text{clamp}((d - s) / (h - s), 0, 1)$ and chooses nursery/growing/ready from those dates. Artwork adds presentation categories, not biomass measurements. A long calculation does not move d . Research dialogue does not secretly accelerate crop maturity. See the numerical chapter for inclusive-day occupancy, expiry, yield shocks and the limits of simulated execution.

2.3 Demand baseline and interpretation

Synthetic demand: observed fixture values and one-step EWMA

Caixin · alpha = 0.35 · fabricated repeating history · no public/weather features




Figure 3. Latest EWMA level = 31.001183 kg/week. Future expected demand = bookings + max(0, level - weekly bookings).

Figure 3. Freshly generated illustration from the source fixture. The blue one-step forecast uses only preceding observations; the shaded observations match the six held-out weeks used by the archived numerical evaluator. The newest EWMA level is used across future weeks, with already booked quantities preserved and only positive residual demand added. A repeating synthetic history can make a baseline appear accurate without establishing performance on real orders. Source: `models, evaluate_numerical.py`. PNG · SVG.

All figures have source SVGs, raster exports, accessible titles/descriptions and textual explanations. Generation code and underlying fixture/forecast data support reproduction. Architecture annotations are authored from source inspection; no generative-image model or inferred experimental data was used to create these figures.

3. Results and interpretation

3.1 Numerical evidence

The archived numerical evaluation compares EWMA with a last-week baseline over six held-out weeks per crop. Caixin EWMA MAE is approximately **2.383 kg/week**, versus **2.667 kg/week** for last-week naive; WAPE is approximately **7.77%** versus **8.70%**. These are synthetic targets from a deterministic four-week pattern. The four crops share the pattern with level offsets, so they do not constitute four independent real-world datasets.

The same archive records constraint-checked strategy outputs and time-limited solver statuses. Its FEASIBLE statuses should not be translated into “the optimal farm plan.” Whole-bed schedules may leave both shortages and surplus because crops, dates and shelf life constrain allocation. The locally calculated minimum across three stress cases is a scenario downside, not a statistical lower quantile. Financial outputs use assumed prices and costs; they are not audited profit.

3.2 Adviser evidence

The v7 local and public probes each made one actual request. Their context/replay/isolation checks passed, but both responses remained **unsupported**. This distinction is central: transport correctness and budget control are necessary engineering properties; they do not validate agronomic interpretation. The older authenticated trial summary retains its own tested roles, models, budget and dates.

The app's capability display also relies on archived trial records rather than a fresh provider probe. The gap register requires the next iteration to distinguish historical verification, current configuration and recent successful execution without generating paid background calls merely for badges.

3.3 Interaction and infrastructure evidence

The v7 report records 111 public browser checks and separate novice-agent checks. Automated completion times, screenshots and AI reviewer opinions do not establish farmer learning, enjoyment, actual screen-reader support or physical touch/keyboard comfort.

The archived shared-capacity experiment completed two concurrent v6/v7 numerical jobs in approximately 22 seconds each and sampled at least 2.34 GiB available memory. This is evidence for that bounded workload, not sustained scalability. The 11 September read-only inspection found one running Fly Machine, 1/1 health checks and seven published editions. No new concurrency experiment was performed.

4. Discussion and next iteration

FarmTact's strongest implemented separation is between numerical authority, contextual interpretation and simulation-only decisions. Its most consequential open issues are adviser evidence quality, truthful capability labels, temporal and metric precision, realistic out-of-sample validation, and human evaluation. Adding more role names or source cards cannot resolve these issues by itself.

The gap register assigns stable IDs, implementation evidence, proposed changes and measurable acceptance criteria. The specifications now reference those IDs and explicitly distinguish a current-state correction from a requested future capability. The existing four-recipe demonstration remains valid as a bounded engineering fixture; new crop recipes, observed growth response, learned demand models and real operations each require their own data and validation.

5. Reproducibility and references

See reproducibility for commands and environment prerequisites, and the audit record for commands actually executed. The detailed chapters provide function-level source links and external references where relevant. Source links target repository-relative files; the baseline commit at the top defines the inspected implementation. Frozen release manifests identify the distinct deployed source/image.

The report does not use archived pass counts as a substitute for new scientific experiments. It supplies the implementation account and acceptance criteria needed to design those experiments responsibly in the next iteration.

AI provider and council implementation report

Historical baseline: This chapter records the pre-remediation audit of source a96025e and published v7. Its implementation findings and measured results describe that baseline. Read the V8 remediation report for the current simulation engine, synthetic model evaluation, Council contracts, tests and remaining limitations. Source links below are navigation aids into the maintained repository; use the frozen baseline to reproduce the original inspection.

Audit date: 11 September 2026 (UTC)

Audited revision: a96025e49226d8ca641cf07b49192bd85b464805

Scope: the DeepSeek provider boundary, mission and conversation councils, the playable council study, vision, numerical forecasting and simulation interfaces, evaluation evidence, persistence, replay, budgets, admission control, and egress.

This chapter is part of the FarmTact: scientific implementation report. Consolidated issue IDs, priorities and next-iteration acceptance criteria are maintained in the implementation gap register; this chapter supplies the detailed call-path evidence behind its AI-* entries.

Safety and interpretation boundary. FarmTact is a synthetic planning game in `execution_mode=test`. Its generated text is advisory, its crop parameters are demonstration fixtures, and accepted worklists are for simulation only. Nothing in the audited implementation authorizes a physical farm action.

Abstract

FarmTact is a hybrid decision system. A local statistical baseline forecasts unbooked demand; a local CP-SAT planner constructs three whole-bed policies; a deterministic daily simulator checks timing, inventory, labour, cash, waste and service; and, only after an explicit user request, DeepSeek models interpret a frozen copy of those results. The language models neither calculate strategy totals nor directly mutate the farm. A backend evidence gate checks model output before the development policy can accept a strategy for simulation.

The current council has seven roles. Five use the configured `deepseek-v4-flash` route and two use `deepseek-v4-pro`. A complete main-mission council normally makes seven provider requests and permits at most two additional schema-repair requests. Optional vision adds one request, but that workflow reads an artificial batch label; it does not assess a plant. Persistent conversations make one request for a direct question, two for an invitation exchange, or seven for a full council, before bounded repairs.

There is no learned crop-growth, yield, disease, nutrient, or climate-response model in the current application. "Growth" on the farm board is a linear schedule progress indicator derived from declared sow, transplant and harvest dates. Marketable yield is a fixed recipe coefficient

multiplied by bed area, then perturbed only by declared synthetic scenarios. These boundaries are central to interpreting both the interface and evaluation results.

Questions and method

This audit answers five questions:

1. What can trigger paid inference, which model does each route use, and how many HTTP requests can it create?
2. How do council outputs affect a planning decision?
3. Which parts are statistical estimation, mathematical optimization, deterministic simulation, scripted dialogue, or language-model inference?
4. What is persisted, replayed, validated, rate-limited, and allowed to leave the process?
5. Which repository claims are stale, incomplete, or stronger than their evidence?

The method was static source tracing from each public API trigger through job creation, frozen inputs, gateway construction, response validation, persistence, and replay. It also inspected checked-in trial artifacts and compared provider claims with DeepSeek's official documentation. No paid inference request, Fly deployment, or farm operation was made for this audit. Implementation facts take precedence over prose and historical reports. A checked-in PASS file establishes only what that particular artifact records; it is not a current provider health check.

System boundary



This separation is visible in the implementation: missions call the local `plan()` pipeline before considering the credential or reserving an inference budget. The council receives calculated strategies and exact references, and the development policy performs its own eligibility and evidence checks (`services/api/app.py`).

Terms used in this report

Term	Meaning in FarmTact	Provider call?
Forecast	A per-crop EWMA of historical demand, converted to residual demand after confirmed orders	No
Optimization	A bounded CP-SAT search over whole-bed crop allocations	No
Simulation	Daily stock-flow accounting under three declared yield/demand scenarios	No

Term	Meaning in FarmTact	Provider call?
Schedule preview	Date interpolation and stage artwork in the UI	No
Computed perspective	A deterministic text template over scenario metrics	No
Scripted council study	Rules, buttons and predefined dialogue around local calculations	No
Advisor message/council	Structured natural-language interpretation returned by DeepSeek	Yes
Replay	Reading previously stored messages/events	No

The numerical implementation has its own detailed report in `numerical-models-and-growth.md`. Data lineage and the evidence register are covered in `system-data-and-evidence.md`.

Provider contract

Current repository allowlist

The gateway admits one origin, two endpoints, three model identifiers and fourteen named routes. Configuration is accepted only when it exactly matches the compiled policy; unknown or missing fields, a different origin, an extra route, a larger limit, or a model/route mismatch fail closed (`gateway policy`, `GatewayConfig.load`, `config/deepseek_runtime.json`).

Route	Configured model	Capability	Deployed caller in audited app
demand_analyst	deepseek-v4-flash	text	Mission and conversation councils; direct conversations
weather_analyst	deepseek-v4-flash	text	Mission and conversation councils; direct conversations
market_analyst	deepseek-v4-flash	text	Mission and conversation councils; direct conversations
production_analyst	deepseek-v4-pro	text	Mission and conversation councils; direct conversations
supply_chain_analyst	deepseek-v4-flash	text	Mission and conversation councils; direct conversations

Route	Configured model	Capability	Deployed caller in audited app
profit_analyst	deepseek-v4-flash	text	Mission and conversation councils; direct conversations
planning_chair	deepseek-v4-pro	text	Mission and conversation councils; direct conversations
visual_observer	deepseek-v4-flash-vision-exp	vision	Optional synthetic-label observation
evidence_extractor	deepseek-v4-flash	text	No deployed HTTP call site found; gateway/trial surface only
crop_alias_resolver	deepseek-v4-flash	text	No deployed HTTP call site found; gateway/trial surface only
runtime_researcher	deepseek-v4-pro	text	No deployed HTTP call site found; gateway/trial surface only
document_vision	deepseek-v4-flash-vision-exp	vision	No deployed HTTP call site found; gateway/trial surface only
satellite_visual_reviewer	deepseek-v4-flash-vision-exp	vision	No deployed HTTP call site found; gateway/trial surface only
test_evaluator	deepseek-v4-pro	text	No deployed HTTP call site found; local evaluation does not use it

“Configured” is not synonymous with “used.” The route registry is broader than the current application call graph. The generic two-request tool continuation, model-list request and streaming path are implemented in the gateway and exercised by standalone trials, but the deployed mission/conversation workflows call structured chat and vision only (gateway methods).

Comparison with official provider documentation

DeepSeek's official documentation, checked on the audit date, documents the API origin and current model aliases in its Quick Start, the POST /chat/completions request and JSON response mode in Chat Completion, model discovery in List Models, inline image content in Vision, and the continuation treatment of reasoning_content in Thinking Mode. The provider also documents tool calls as model requests whose functions are executed by the client, not by the model itself (Tool Calls). These documents support the protocol choices and aliases in the repository as of this audit.

They do **not** prove that the FarmTact deployment's account currently has access, credit, or healthy connectivity. The application does not perform a fresh model or chat probe when it renders capability status. See “Evidence freshness” below.

Request construction and privacy

Every inference request is constructed by DeepSeekGateway with:

- an environment-delivered bearer credential; no credential is accepted from an HTTP client;
- `trust_env=False`, so ambient proxy configuration is not inherited;
- redirects disabled and the origin fixed to `https://api.deepseek.com`;
- the model selected from the server-side role table;
- `execution_mode=test` and `data_mode` limited to `synthetic_demo` or `historical_replay` (mode checks);
- a stable provider `user_id` derived by hashing the internal tenant identifier, prefixed with `farmtact_` (pseudonymization).

DeepSeek's official Rate Limit & Isolation page permits an alphanumeric, hyphen or underscore `user_id` and says not to send private information in it. FarmTact's hashed pseudonym complies with that shape; the source tenant ID is not submitted as `user_id`. This is pseudonymization, not an anonymity claim: frozen prompts still contain synthetic farm and conversation context needed for the task.


Gateway bounds

Bound	Hard ceiling
Requests in one gateway budget	16
Reserved output tokens	16,384
Concurrent requests	1
Images per request	4
Bytes per source image	4 MiB
Maximum image dimension	4,096 px
Serialized request	16 MiB
Response	4 MiB
Gateway wall time	300 s
Connect/read/write/pool timeouts	10 / 90 / 30 / 5 s

These are code-enforced maxima, not statements about provider-wide limits (limit constants). The budget reserves the declared maximum output tokens immediately before each HTTP request. A failed or shorter request still consumes the local request count and reserved-token

allowance; usage is not reconciled downward (`RunBudget.reserve`). That is a conservative token budget, but it is not a currency-cost estimator.

Validation lifecycle



For text JSON, the gateway asks for `response_format={type: json_object}`, accepts exactly one choice with an allowed finish reason, checks the returned model, parses the content, and validates it with the supplied Pydantic type (`chat_json`, response validation). Provider errors are converted to bounded messages and only safe request identifiers are retained (sanitization). Raw request bodies and credentials are not copied into public error records.

The generic tool method permits exactly one allowlisted tool call, validates its arguments locally, runs a local callable, then submits a second request with the tool result. Private `reasoning_content` is retained only in the local continuation and removed from public return values (`run_tool_json`). This method is implemented and trialled, but it is not part of the current deployed council.

Inference triggers and exact request counts

All counts below are provider HTTP requests, not UI turns, agents shown on screen, or local tool executions.

User/API trigger	Normal calls	Maximum calls	Models/routes	Notes
Browse, bootstrap, open advisor, open saved conversation	0	0	—	Read or create frozen state only

User/API trigger	Normal calls	Maximum calls	Models/routes	Notes
Local mission, council=false	0	0	—	Forecast, optimize, simulate and validate locally
Main mission council	7	9	5 Flash + 2 Pro; same role may repeat for shared repair	One call per current role; at most two schema repairs across the council
Main mission council + with_vision=true	8	10	council models + vision-exp	One synthetic label image call before council
Direct conversation message	1	2	selected advisor route	At most one schema repair
Invite advisor	2	3	invited advisor, then original advisor	At most one shared schema repair
Persistent conversation council	7	9	5 Flash + 2 Pro	One sequential turn per role; at most two shared repairs
Direct + invite + council trial sequence	10	14	role-dependent	Aggregate implemented by deepseek_conversation_trial.py; trial budget ceiling remains 16
Playable council-research action or numerical run	0	0	—	Scripted dialogue and local calculation
“Ask DeepSeek” inside completed council research	1	2	selected advisor route	Creates a separate frozen research conversation; one question can require one repair
Planning or conversation replay	0	0	—	Stored events/messages only
Standalone tool-call trial	2	2	configured text route	One model tool selection + one continuation; no deployed call site
Standalone model-list probe	1	1	/models	Diagnostic script only
Standalone stream trial	1	1	configured text route	Diagnostic script only

The main mission reserves exactly $7 + 2$, plus one when vision is requested (reservation). Conversation jobs reserve nine for a council and five for either a direct or invite job; unused capacity is returned after the job. Thus direct and invite reservations are more conservative than their true maxima (conversation worker). The main web client currently submits the council flag but does not expose `with_vision`; synthetic-label vision is an optional API capability rather than part of the ordinary board mission flow (client mission request).


The seven-person council

The authoritative roster and order are defined in `packages/agents.py`:

Order	Person	Role	Model	Frozen responsibility
1	Ravi	Demand analyst	Flash	booked orders, forecast residuals, quantities and dates
2	Hana	Weather analyst	Flash	observed context, freshness and environmental uncertainty
3	Idris	Market analyst	Flash	prices and sourced buyer/grower reactions, without treating reaction as demand
4	Mei	Production analyst	Pro	crop recipes, lead times, space and feasible production
5	Lina	Supply-chain analyst	Flash	inventory, expiry, inputs and delivery timing
6	Ben	Profit analyst	Flash	calculated cost, cash, labour and margin
7	Asha	Planning chair	Pro	reconcile evidence-backed trade-offs among feasible strategies

There is no independent critic in the current roster. The former critic belonged to the archived six-role council. Current automatic checking is ordinary local code, separate from every persona (`council_review_issues`).

Mission council data flow



The council receives each strategy's metrics, constraints, forecast cutoff and uncertainty, the lead-time rule, weather scope, market summary, frozen news evidence, and six selected research-register entries (context assembly). Every role is called sequentially and returns exactly one `Claim` object. The prompt asks for a concise, role-specific finding; thinking is explicitly disabled. Only Asha receives the six preceding claims as `prior_claims`. The six specialists do not see one another's claims (council loop). This is a sequential roster with a final synthesis context, not a multi-round debate or two challenge rounds.

The timing of validation matters. Each specialist claim is schema-checked, then reference-checked and augmented with status and rejection_reasons before it is appended. Because Asha runs last, her `prior_claims` includes both accepted and rejected specialist findings and the reasons for a rejection. The chair can respond to that state, but her response cannot retroactively validate an earlier claim. There is no short-circuit: all seven roles are attempted unless cancellation, budget, transport or response failure stops the job.

The word “tool” in council events can be misleading. The model is not offered a function schema and does not decide to run the forecast, planner, source retrieval or database queries. FarmTact runs those operations first and serializes selected results into the model's user message. The `tool_result_refs` in a claim are keys into that precomputed map. This council therefore implements retrieval-augmented structured interpretation, not agent-directed tool execution.

What local validation proves

For each mission claim the backend checks:

- every `tool_result_ref` is in the supplied reference map;
- every `evidence_id` is in the supplied six-document context; and
- each numeric literal in the statement exactly equals a numeric value in at least one cited scalar tool reference.

It then marks the claim validated or rejected (claim checks). This is useful structural grounding. It does not prove that a qualitative sentence is entailed by the cited record, that a paper is scientifically valid for the current farm, or that the recommendation is agronomically correct. A schema-valid qualitative error can pass unless another deterministic rule catches it.

The two repair allowances apply only when provider output fails JSON/Pydantic schema validation. They do not retry a schema-valid claim with an unknown reference, uncited number or unsupported recommendation. Such a claim is saved as rejected and the council continues. The same distinction holds in persistent conversations: format failure can consume a repair, while reference or action failure produces a saved unsupported message without another provider request.


What the council decides

The model does not pick an arbitrary allocation. The backend first keeps strategies that are locally FEASIBLE and pass allocation validation, then chooses Balanced when available. With council claims present, acceptance is withheld unless all seven roles are present in order, all claims passed the checks above, and every recommendation is `proceed_simulation`. With no council claims, the local numerical baseline can be accepted on its own. The result is `ACCEPTED_FOR_SIMULATION`, never an operational commitment (acceptance policy).

This gives natural-language advice a limited but material effect: it can veto an otherwise feasible simulation through a rejected or non-proceeding claim. It cannot make an infeasible strategy eligible, choose an unsupported new strategy, or execute farm work.

Persistent advisor conversations

Persistent conversations are a second inference surface. Creating or reading a conversation freezes and returns state without calling DeepSeek. Explicitly posting a message, inviting an advisor, or convening a council enqueues a job (route definitions).



Frozen context

The snapshot kind is farm, scenario or research. The freezer copies the applicable farm/version/hash, terminal calculation, tool-result map, sources, market signals, news and evidence. A research conversation requires a completed calculation for the current research version; scenario context requires a completed scenario (snapshot freezer). Subsequent model turns operate on these saved values rather than silently reading a newer farm.

The prompt includes at most the last 32 messages and an explicit reply graph, along with frozen tool and evidence values. Retrieved material is labelled untrusted (context construction). An invite first asks the invited advisor and then asks the original advisor to respond. A conversation council calls the seven current roles in order; this council also is not a multi-round debate.

Unlike the mission council, every conversational speaker after the first receives earlier persisted public turns, not only the chair. A reply is locally validated and saved before the next prompt is assembled. Asha therefore sees earlier turns and their `validation_status`, including unsupported, but the context projection does not include the detailed `validation_errors` array. The system prompt says earlier rejected turns are visible; in practice the rejected status is visible while its specific reason must be inferred from the response and references. Each turn's `reply_to` points to the previous advisor reply, forming a chain even though all turns also address the same council question (worker loop).

Conversation models also receive no executable function tools. Their `tool_refs` refer to the frozen lookup table. `proposed_actions` are JSON suggestions validated after generation and displayed as hypotheses; they are not invoked by the worker.

Reply validation and proposed actions

An advisor reply can contain narrative, references and bounded hypothetical action controls. The validator rejects unknown tool, evidence or highlighted-message references. It also rejects numeric literals, spelled-out numbers, dates and relative dates in prose, directing exact quantities into separately rendered reference cards (reply validation). This conversation rule is stricter than the main mission rule, which permits exact cited numeric values.

Action controls are locally bounded to:

- batch delay: 0–14 days;
- batch yield: 50–100%;
- crop demand: 50–150%; and
- labour or cash: 50–150%.

Targets must exist in the frozen snapshot. Actions remain `hypothesis_only` or are blocked; the reply cannot execute them. The prompt requests two short sentences and 400 characters, but the response schema permits 900 characters and the validator does not enforce the sentence or 400-character limits. The prompt also asks for at most three tool refs, one evidence ref, one highlight and one action, while the schema allows 10, 8, 8 and 3 respectively. Reference membership is checked, but those lower prompt cardinalities are not. These are open deterministic-validation gaps (reply schema, `_system_prompt`).

Partial results and restart semantics


Each validated or unsupported advisor message and its audit are persisted after its turn. A later failure can therefore leave a `PARTIAL` conversation with useful saved replies. On process restart, a conversation job that had reached `RUNNING` is marked interrupted and is not automatically requeued, because the process cannot know whether a paid request completed before interruption (restart handling). This is the safe default for uncertain side effects. By comparison, purely numerical scenario and research jobs can be requeued.

`COMPLETED` is the request lifecycle status: it means all requested provider turns were received and persisted. It is not an answer-quality verdict. Each reply carries its own `validation_status` (`references_verified` or `unsupported`) and remains an `unverified_advisor_interpretation` even when references pass. A seven-turn request may therefore be `COMPLETED` while one or more of its saved replies are unsupported (message persistence and request completion).


Replay explicitly reports `execution_mode=replay`, `inference_origin=stored_messages`, and `inference_triggered=false` (public replay).

Playable council research is scripted

The page labelled “Council research” is primarily a deterministic research prototype, not a seven-call LLM council. It uses typed actions and a regular-expression interpreter to reserve a bed, change an order's confirmation state, change a labour limit, challenge an assumption, run a local calculation and choose an eligible simulation. A new session records `inference_calls=0`, `dialogue_mode=scripted_research` and `operational_execution=false` (session creation).



The numerical worker validates a frozen hash and version, calls `calculate_research()` locally, and creates predefined turns only for inputs that changed, followed by a Planner turn. An initial run normally shows Production and Demand plus Planner; changing labour can add Profit. Roles with no changed input do not speak (research calculation). Agreement between those scripted speakers is explicitly not evidence.



Limits are 30 sessions per tenant, 400 actions per session, 12 numerical versions per study, and one queued/running research calculation per tenant workspace (research routes). Choosing a result records a research-participant simulation choice and does not change the main farm.

After a current result completes, the user may deliberately open “Optional actual DeepSeek interpretation.” That UI creates a separate persistent conversation with `snapshot_kind=research` and posts one direct question (component). It does not run the full seven-role council. “One

question” normally means one provider request, but one bounded schema repair can make it two.

Forecasting, growth and numerical simulation

What is actually learned or estimated

There is no trained machine-learning model. The only fitted statistical state is a single exponentially weighted moving average per crop:

$$E_t = \alpha D_t + (1 - \alpha)E_{t-1}, \quad \alpha = 0.35 \text{ by default.}$$

Only demand records available before the cutoff and before the planning date enter the recursion. Each future week receives the same EWMA residual after booked demand:

$$R_w = \max(0, E - B_w).$$

That residual is placed on the end date of the week. Confirmed orders remain dated commitments. The forecast labels uncertainty as the fixed synthetic scenarios below; it does not calculate a prediction interval (forecast implementation).

Scenario	Yield multiplier	Residual-demand multiplier	Statistical meaning
Low yield / high demand	0.85	1.15	Declared stress case
Central	1.00	1.00	Declared reference case
High yield / low demand	1.10	0.90	Declared favourable case

The cases each enter the optimizer once and carry an equal declared weight of one third. The solver does not explicitly read the weight field, and the seed argument does not generate samples. They are not learned probabilities, confidence intervals, or calibrated quantiles.

Plant growth is a schedule, not a biological model


For an existing batch the server derives a stage from recorded dates:

- ready when `harvest_date <= planning_date`;
- nursery before transplant; otherwise
- growing.

Its progress is the clamped linear fraction

$$p = \text{clip}_{[0,1]} \left(\frac{\text{planning day} - \text{sow date}}{\text{harvest date} - \text{sow date}} \right).$$

The farm view implements exactly this calculation (`farm_view`). The frontend preview performs a similar date interpolation and swaps stage artwork. It does not integrate sensor readings, thermal time, light, nutrients, water, humidity, disease, cultivar traits, biomass or observed growth. The UI describes it as schedule preview rather than an observation.



For proposed crops, marketable yield is fixed by the synthetic recipe:

$$Y_{bed} = A_{bed} y_{recipe},$$

where A is bed area and y is declared marketable kilograms per square metre. The simulator applies a scenario multiplier on the scheduled harvest date. There is no growth curve between sow and harvest. Existing-batch forecasts simply copy their declared harvest dates and expected kilograms.

The board and preview also have a boundary inconsistency worth retaining in tests (previewBed): the server labels an existing batch ready on its harvest date, while the frontend's preview renders an existing bed empty at or after its scheduled harvest; new accepted allocations render ready on their harvest date. The two views answer slightly different questions but do not state that difference clearly.


Optimization and daily accounting

The planner builds whole-bed candidates by counting backward from permitted harvest dates through fixed nursery and grow durations. CP-SAT applies occupancy, nursery, weekly labour and conservative cash constraints and solves Lean, Balanced and Resilient policy objectives. It uses one worker, a fixed seed and a four-second limit; a feasible solution need not be proven optimal (planner). A deterministic backward scheduler is used when CP-SAT produces no solution, then passed through the same validation (plan).

The daily simulator:

1. adds scheduled harvests on their declared day;
2. expires lots by shelf life;
3. serves confirmed orders and residual forecast demand from FIFO stock;
4. records unfilled demand without retroactively filling it from a later harvest;
5. calculates fixed input, labour, packing and disposal costs; and
6. checks mass balance and hard constraints.

This is a stock-flow model, not crop growth simulation (daily simulation). Scenario controls can directly move an existing batch harvest date, scale its yield, scale demand and change resource ceilings before a new local plan is run (scenario controls). The crop-demand control scales both historical demand and matching current orders, so it changes the EWMA baseline and commitments together; it should be interpreted as a broad hypothetical branch, not a causal market forecast.



Synthetic disruption and “replan”

The replan API does not ingest a measured plant observation. It takes the first batch in the current synthetic farm, adds seven days to its harvest date and multiplies expected marketable yield by 0.8, saves a new farm version, and replans (replan route). This is a useful deterministic demonstration of versioning and response to a disruption. Calling it observed growth, a detected delay, or a predictive crop model would be inaccurate.

Vision

The only deployed vision path generates an 800×360 PNG in memory containing the words “SYNTHETIC BATCH” and a run-derived FT- label. It sends that image to `visual_observer`, checks that the returned label is exact, and records `agronomic_measurement=false` (`observe_fixture`). The resulting structured description can be included in the text council's frozen references.

The gateway independently validates image count, encoded size, dimensions and supported static formats, rejects animated images and decompression bombs, converts the image to normalized PNG, and transports it inline (image normalization, `vision_json`). No current UI path uploads a crop photograph. The vision feature therefore demonstrates bounded multimodal transport and label extraction only; it gives no evidence of plant health, growth, maturity, disease or yield measurement.


Persistence, provenance and replay

Mission creation stores a frozen farm snapshot and hash, input version, cached source snapshot, research-register version hash, market signals, frozen news, council version, requested mode and parent/disruption metadata before work begins (mission creation). During a mission it persists strategies, claims, safe provider audits, usage summaries and ordered events. The acceptance transaction rechecks the current farm version and content hash; a changed farm invalidates acceptance and starts a fresh local, no-council calculation.

Conversation records likewise preserve their snapshot and input hash, messages, reply graph, references, job status, model, usage and validation status. Replay returns these stored records without new inference. The mission replay endpoint also returns stored events and identifies `inference_origin=stored_events` (mission replay).

The records do not currently carry an explicit prompt-template version on each provider request or output. `council_version` identifies the roster on missions, but it is not sufficient to reconstruct every prompt, schema and validator revision. The generic event `schema_version='1.0'` identifies the event envelope, not the Pydantic advisor/claim schema or prompt revision. Conversation records similarly store model and usage without those semantic versions. Exact reproducibility therefore also depends on the source commit, frozen edition image and checked-in configuration.

Admission, budgets and egress



AI-triggering paths have durable limits of 6 requests per IP per minute, 20 per IP per hour and 12 per tenant per hour. The broader API, write, session, preview and event-stream limits are separate (rules and route matcher). The AI route matcher conservatively includes all planning-run creation and replan posts, including requests whose body later selects `council=false`; these can consume admission allowance even though they cause zero inference.

The persistent daily ledger has a hard ceiling of 48 provider calls. Published editions delegate that ledger to the shared edition-control service; local instances use their own database (daily reservation). Main mission and conversation workers reserve worst-case counts before opening the gateway and release unused requests against the original UTC day afterward.

The service endpoint loads a protected local credential when necessary, installs the process egress policy, then imports the API server (scripts/serve.py). The Python socket policy allows api.deepseek.com:443, the configured PostgreSQL host, validated private edition upstreams when running as a gateway, and a validated .flycast control host (services/api/egress.py). It checks hostname resolution and socket-connect audit events.

This is defense in depth inside the Python process. It is not an operating-system, container, subprocess or infrastructure firewall. Production isolation therefore cannot be inferred from the hook alone. The gateway's fixed HTTPS origin and disabled redirects independently constrain provider HTTP traffic.

Evaluation and checked-in provider evidence

Numerical evaluation

scripts/evaluate_numerical.py runs locally. It performs rolling held-out checks on the synthetic generator, compares EWMA with a last-week naive forecast, and verifies that reported strategies have no recorded violations. It makes no LLM call. Its own limitations correctly say that outcomes and recipes are synthetic, scenario weights are not calibrated intervals, and no real-farm forecast has been promoted. Passing it demonstrates internal consistency on generated data, not agronomic accuracy or external predictive validity.

Provider artifacts and what they establish

Artifact	Recorded time / scope	Recorded requests	What it establishes	Limitation
reports/provider_egress_probe.json	8 Sep 2026 08:29 UTC	1	A model-list request succeeded under the process hook then	Not a current health or chat test
reports/deepseek/latest.json	8 Sep 2026 08:41:30-08:41:44 UTC	15 cumulative	Flash/Pro JSON, generic two-call tool continuation, vision, streaming and an old council trial passed then	Council section is the archived six-role roster, including independent critic
reports/seven_agents/deepseek.json	Seven-agent local candidate; artifact contains no timestamp	10	Direct reply, two invited replies, seven current council replies, stored replay without inference	Cannot be independently dated from the JSON; outputs include unsupported claims

Artifact	Recorded time / scope	Recorded requests	What it establishes	Limitation
reports/deepseek_conversation_live.json	Deployed conversation aggregate; artifact contains no timestamp	11	Saved direct/invite/council messages, grounding states and no-inference replay were observed	Does not by itself identify current account health
reports/v7/advisor_local.json	9 Sep 2026 23:55 UTC event	1	One local v7 advisor boundary run completed	Reply was not reference-verified
reports/v7/advisor_public.json	10 Sep 2026 00:58 UTC event	1	One public v7 advisor run completed and remained isolated	Reply was marked unsupported
reports/fly_integrated_demo.json	Historical integrated script artifact	—	Some component details are inspectable	Overall state is INCOMPLETE; it is not release proof

On a fresh full run, current `scripts/deepseek_trial.py` makes seven base capability requests: model list, Flash JSON, Pro JSON, two-request tool continuation, vision and streaming. `--with-council` adds seven current roster calls, for fourteen normally and a hard ceiling of sixteen. However, the `reports/deepseek/latest.json` file consumed by the capability badge came from the prior six-role version and records fifteen cumulative requests. The script's council scenario is a transport test using toy partner-supply alternatives and a report-local decision event. It is not the deployed planner, which deliberately does not invent partner purchases. A claim that the current seven-role standalone DS-G2 trial has fresh authenticated proof would therefore be unsupported.

Evidence freshness and the capability badge

`capabilities()` labels DeepSeek text verified when two fields in `reports/deepseek/latest.json` are PASS and a credential exists in the current environment. It labels vision verified solely from the archived vision PASS. It does not inspect the artifact timestamp, check the current roster, call `/models`, perform a current account health request, or verify that the credential can use the configured aliases.

Badge field	Actual inputs	Missing from the decision
DeepSeek text status	Archived Flash JSON PASS + archived Pro JSON PASS + current process key presence	Trial age/roster, key validity, account credit/access, current model availability, live request success
DeepSeek model list	Models discovered in the archived report	Current configured-vs-discovered comparison and fresh discovery
DeepSeek overall trial	Archived report's <code>overall_status</code>	Current provider/account state

Badge field	Actual inputs	Missing from the decision
Vision status	Archived vision PASS only	Even current process key presence, plus every freshness/account check above

Consequently the UI badge means approximately “a historical checked-in probe passed, and for text a key is present.” It must not be described as a fresh provider or account verification. A future capability contract should expose at least:

- `historical_trial_status` and `historical_trial_finished_at`;
- `configured_models` separately from models discovered by that trial;
- `credential_present` as a server-side boolean only;
- `current_availability` as unknown unless a deliberately authorized, timestamped health probe has run; and
- roster/trial version so old six-role proof cannot imply current seven-role proof.

Findings and specification corrections

The following items should be treated as corrections or explicit future-work gaps, not silently rewritten as capabilities that already exist.

Priority	Finding	Evidence	Required correction or next iteration
Critical	“Plant growth simulation” currently means fixed date/yield scheduling plus daily inventory flow. There is no biological or learned growth model.	<code>packages/models</code> , <code>farm_view</code> , <code>planner</code>	Rename user-facing/spec claims to schedule and production simulation. Define validation data and equations before adding any empirical growth model.
High	Current council is seven turns plus at most two shared repairs; docs that say eight turns are stale.	<code>packages/agents.py</code> , <code>services/api/council.py</code>	Replace “eight advisor turns plus two repairs” with “seven advisor turns plus at most two shared schema repairs.”
High	The independent critic no longer exists.	Current roster above; historical artifact has six old roles	Remove current-tense “simulated critic” language. Say deterministic backend evidence gate, and label the old critic as archived.
High	Specifications describing two challenge rounds do not match either deployed council.	Mission and conversation loops each call roles once	Specify seven sequential specialist/chair turns. If debate is desired, design and budget explicit rounds rather than implying them.

Priority	Finding	Evidence	Required correction or next iteration
High	A runtime requirement recommends enabled thinking for deep council work, but both current councils explicitly send thinking='disabled'.	services/api/council.py, conversation worker	Either change the requirement to disabled thinking for bounded structured claims, or implement and evaluate an enabled-thinking variant with privacy/cost limits.
High	Latest general capability artifact is an old six-role trial, although current trial code has seven roles.	latest.json timestamps and all_six_roles_present; current script	Separate archived transport evidence from current roster acceptance. Do not claim freshly authenticated seven-role DS-G2 until a deliberately authorized run is recorded.
High	Capability verified status has no freshness or account-health semantics; vision does not even require present credential.	views.capabilities	Publish archived timestamp/status and current availability separately; use unknown when no fresh probe exists.
Medium	Local checks validate schema, reference membership and numeric equality, but not qualitative entailment.	Council and conversation validators	Describe claims as structurally grounded or unsupported. Add deterministic entailment rules or reviewed eval sets before calling them fact-verified.
Medium	Persistent request records lack an explicit prompt/schema/validator version.	Mission and conversation persistence	Store prompt-template version, output-schema version, validator version, route and source commit with every request.
Medium	Conversation prompt says two sentences/400 characters and lower reference/action counts, while the schema allows 900 characters, 10 tool refs, 8 evidence refs, 8 highlights and 3 actions. The validator does not enforce the lower prompt limits.	AdvisorReply, validator	Align the schema and deterministic validator with the documented limits, or remove the stronger prompt claims.

Priority	Finding	Evidence	Required correction or next iteration
Medium	Main council reservations count requests and maximum output tokens, not total monetary cost; cancellation is checked between calls, not an external hard kill.	gateway budget and mission worker	State current semantics exactly. Add price-versioned cost accounting and an independently enforceable job deadline only if required.
Medium	Optional vision reads a synthetic label, not a plant.	services/api/vision.py	Rename it synthetic label verification in capability text. A real crop-imaging feature needs consent, acquisition provenance, task-specific validation and agronomic uncertainty.
Medium	UI/server stage semantics differ around harvest for existing batches.	server farm view and frontend preview	Define whether preview means crop present, scheduled work, or post-harvest bed state; encode one date-boundary contract and test it.
Medium	Demand scenario control scales historical demand and booked orders together.	services/api/scenarios.py	Split forecast-history, unbooked-demand and confirmed-order interventions so causal meaning is visible.
Low	Four configured text helper/evaluator routes and two extra vision routes have no deployed caller; generic tools/streaming are trial-only.	route registry and call-site audit	Mark them reserved/experimental or remove them from product capability claims until an owned workflow exists.
Low	The Python egress hook is described too broadly as the production network boundary.	services/api/egress.py	Call it process-level defense in depth and document infrastructure egress controls separately.

Council assurance and acceptance tests

The council boundary should be reviewed as a chain of independently testable properties. Existing tests cover many structural properties; the remaining quality claims need concrete fixtures rather than provider-success assertions.

Property	Acceptance test	Present evidence / gap
Exact roster and order	Submit a mocked council and assert Ravi, Hana, Idris, Mei, Lina, Ben, Asha exactly once in that order.	Current roster and mission-gate tests cover order; retain as a release gate.
Request ceiling	Force invalid schema responses at different roles and assert no more than two repairs and nine total council calls; add vision and assert ten.	Gateway and conversation repair tests exist; preserve an integrated main-mission maximum test.
Chair context	Return one locally rejected specialist claim and assert Asha's outgoing prompt contains its status and rejection_reasons; assert specialist prompts do not contain earlier claims.	Behavior follows current mission code; add a direct regression because it defines the intended synthesis chain.
Conversation context	Mark an early turn unsupported and assert every later prompt contains its public validation_status; decide whether to include bounded validation reasons and test the chosen contract.	Status is present, reasons are omitted. Current prompt language is stronger than the projection.
No model tool execution	Capture provider payloads and assert council requests have no tools; verify every cited key already exists in the frozen context.	Architecture currently satisfies this; encode it explicitly to prevent the generic gateway tool path being introduced accidentally.
Numerical authority	Have every model recommend an allocation absent from local results and assert it cannot be accepted or added. Have models approve an infeasible strategy and assert local validation still blocks it.	Backend chooses only locally eligible recorded strategies. Keep adversarial integration coverage.
Veto semantics	Give one otherwise valid claim exclude_unsupported and assert the mission becomes REVIEW_WITHHELD; run the same frozen input without council and assert the local baseline remains inspectable/eligible.	Current council_review_issues has this behavior. A paired test makes the model's exact influence measurable.
Numeric grounding	Generate zero, negative, comma-formatted and decimal literals with and without exact scalar refs. Assert only exact cited values pass; include dates and identifiers so the regex boundary is known.	Mission checks floats by tolerance; conversation blocks all quantitative prose. More adversarial fixtures would document edge cases.
Qualitative grounding	Use a valid reference ID with a statement that contradicts the referenced value and assert rejection.	Not currently implemented. This is the principal grounding gap. Build a typed claim/evidence relation or reviewed entailment fixture before claiming fact verification.

Property	Acceptance test	Present evidence / gap
Research-page truthfulness	Exercise all scripted actions and numerical runs with provider transport set to fail if called; assert zero inference. Then use the explicit DeepSeek action and assert only a separate direct conversation is queued.	Existing council-research tests cover local behavior; retain a transport tripwire in browser/API integration.
Replay	Record completed and partial jobs, replay them with provider transport set to fail if called, and assert byte-stable messages/events plus <code>inference_triggered=false</code> .	Existing replay trials and tests support the invariant. Add it to every edition's release gate.
Restart uncertainty	Interrupt a paid conversation after one saved turn and assert INTERRUPTED/partial preservation with no automatic resubmit; interrupt local research calculation and assert safe requeue.	Current stores intentionally differ; retain database-backed concurrency tests.
Prompt reproducibility	Persist and assert prompt, schema, validator, council and source versions on every generated record.	Gap: only some surrounding versions are stored.
Capability freshness	Feed an expired six-role report plus a credential and assert UI says historical/unknown, not currently verified; feed a timestamped authorized current probe and assert its exact scope.	Gap: current capabilities() reports verified from stale archive fields.
Scientific validity	Evaluate forecasts and any future growth model on held-out real, provenance-controlled farm data with predeclared metrics and uncertainty calibration.	Gap by design: current evaluation is synthetic and cannot establish real-farm accuracy.

Reproducibility checklist

A no-inference review can be reproduced from a clean checkout with these classes of checks:

```
# Static call graph and route inventory
rg -n "DeepSeekGateway|chat_json|vision_json|run_tool_json|list_models|stream_chat" \
  services runtime scripts

# Verify the current roster and exact compiled/configured gateway policy
python -m pytest -q tests/seven_agents tests/deepseek

# Exercise local forecasting, optimization and scenario behavior only
python -m pytest -q tests/numerical tests/council_research
```

Tests should run with provider transport mocked or the provider credential absent. Authenticated scripts under `scripts/deepseek_*trial.py` and the egress probe make real provider requests and require deliberate execution within the authorized budget; they are not part of this documentation audit and are unnecessary to inspect the code paths or replay stored evidence.

Limitations of this report

This is a source and artifact audit at one revision. It does not measure model answer quality, latency, current price, current account access, Fly network policy, physical farm outcomes, or real crop growth. Checked-in screenshots demonstrate presentation at their named viewport and release context; they are not evidence of provider activity or scientific validity. DeepSeek may change its service documentation after the audit date. Repository configuration remains the application authority until a reviewed change is made.

The most defensible description of the present application is: **a synthetic, versioned farm production-planning simulator with a bounded optional seven-role DeepSeek interpretation layer and deterministic safety gates.**

Numerical models, planning, and crop-growth representation

Historical baseline: This chapter records the pre-remediation audit of source a96025e and published v7. Its implementation findings and measured results describe that baseline. Read the V8 remediation report for the current simulation engine, synthetic model evaluation, Council contracts, tests and remaining limitations. Source links below are navigation aids into the maintained repository; use the frozen baseline to reproduce the original inspection.

Implementation audit date: 11 September 2026

Current numerical versions: synthetic-farm-v1, recipe-ewma-v1, daily-bed-cpsat-v1, council-research-v1


This report describes the numerical behavior in the current repository. It separates implemented calculations from requirements and future research. FarmTact currently combines a deterministic fictional farm, a constant-level exponentially weighted moving average (EWMA), fixed recipe dates and yields, whole-bed constraint optimization, and a daily inventory ledger. It has **no trained or fine-tuned FarmTact model**, no learned plant-growth model, and no validation on commercial farm outcomes. The four crop recipes and all scenario coefficients are engineering fixtures with `validation_status=demo_only` and `origin=synthetic`.

The principal implementation sources are the private-data contracts, fixture generator, forecast baseline, planner and simulator, research wrapper, scenario service, farm view, and world preview. The higher-level inventory in `data-and-models.md` is consistent with this core boundary, but some specification phrases are broader than the implementation; those differences are listed below.

1. What each component actually does

Component	Implemented method	Output	Scientific status
Fixture generator	Closed-form deterministic records	Fictional recipes, beds, batches, orders, history, resources, inventory	Engineering test data; no measured farm observations
Demand forecast	Simple level EWMA per crop	Weekly confirmed and residual kilograms	Statistical baseline evaluated only on synthetic patterned history
Harvest forecast	Copies scheduled batch date and declared marketable mass	Batch/date/kg records	Schedule projection; no fitted maturity or yield model

Component	Implemented method	Output	Scientific status
Growth display	Date thresholds and linear interpolation from sow to harvest	nursery, growing, ready, and percent progress	Calendar visualization; no physiological state estimate
Planner	OR-Tools CP-SAT over Boolean whole-bed candidates	Lean, Balanced, and Resilient schedules	Deterministic optimization over declared assumptions; no ML
Outcome simulator	Daily crop-specific lot ledger, with expiry and FIFO delivery	Harvest, delivery, shortfall, waste, stock, margin, labour	Synthetic counterfactual accounting; no actual outcome prediction validation
Research wrapper	Deep-copies a farm, applies bounded controls, reruns the planner	Versioned, isolated policy comparison	Local numerical experiment; no inference and no main-farm mutation
Runtime advisers	Receive frozen numerical results through the separate DeepSeek gateway	Bounded interpretations and claims	Pretrained model output; advisers do not calculate or validate the plan



Public NEA, SingStat, NASA POWER, and News records are displayed as context but do not enter the current EWMA, harvest projection, scenario factors, or solver coefficients. The feature manifest explicitly records `public_features_used=[]`; see the feature builder and feature manifest.

2. Units, calendar, and numerical conventions

The farm cutoff is a timezone-aware instant. Planning dates are Singapore civil dates obtained by converting the cutoff to Asia/Singapore. The reference cutoff is 2026-09-08T00:00:00Z, which is 8 September 2026 in Singapore. The default horizon is 56 calendar days, represented internally by day indices 0 through 55 and eight week bins.

The contracts use `Decimal` for physical and financial inputs. Several planner calculations convert values to Python floating point before rounding. CP-SAT coefficients use:

- whole beds and Boolean selection variables;
- grams for mass inside the optimization;
- cents for prices and costs;
- minutes for weekly labour capacity; and
- daily occupancy intervals.

The solver uses one search worker and the farm's `fixture_seed`, which improves repeatability. Its default limit is four wall-clock seconds per policy, so the incumbent can still vary with OR-Tools version, machine speed, or load when search ends before optimality. A returned FEASIBLE result is not claimed to be optimal; the output retains the objective, best bound, elapsed time, and solver status.

3. Reference fixture generation

The fixture generator constructs Little Plot Collective, a fictional sheltered-hydroponic farm in Lim Chu Kang. It contains 16 beds of 20 m², so the declared effective growing area is 320 m². It creates eight active batches, 32 future order lines, 48 historical weekly demand records, and one opening inventory lot. The same data is materialized in `synthetic_farm.json`.

3.1 Synthetic recipes

All four recipes assume 16 plants/m², 0.08 sow-labour hours/m², 0.07 harvest-labour hours/kg, four days of shelf life, and two sanitation days after harvest.

Crop	Nursery (days)	Grow-out (days)	Total lead time (days)	Marketable yield (kg/m ² /cycle)	Input cost (SGD/m ²)	Full 20 m ² bed (kg)
Caixin	7	21	28	2.2	2.00	44
Pak choi	7	28	35	2.6	2.50	52
Kailan	10	32	42	2.4	2.80	48
Lettuce	10	25	35	2.3	3.00	46

These coefficients were declared to exercise scheduling logic. They were not estimated from the papers in the evidence registry and are not claimed to be locally optimal crop recipes.

3.2 Existing batches and inventory

For zero-based batch index i , the scheduled harvest date is

$$h_i = d_0 + 14 \left\lfloor \frac{i}{4} \right\rfloor + 6 \text{ days},$$

where d_0 is the planning date. The first four batches therefore harvest on day 6, and the next four on day 20. Sow and transplant dates are obtained by subtracting the recipe's total and grow-out durations. Every batch occupies a whole 20 m² bed and its expected mass is

$$Y_i = 20 \text{ m}^2 \times y_{r(i)}.$$

The opening stock is 5 kg of caixin, harvested one day before the planning date and explicitly expiring two days after it. The lot ledger uses the supplied expiry date; it does not replace it with a recipe-derived date.

3.3 Orders and historical demand

For crop index $c \in \{0, 1, 2, 3\}$ and future week $w \in \{0, \dots, 7\}$, the unmodified order fixture is

$$Q_{c,w}^{order} = 22 + 4(c \bmod 2) + 2(w \bmod 3) \text{ kg},$$

with price

$$P_c = 6 + c \text{ SGD/kg}.$$

Every order is booked seven days before the cutoff and due on day $7w + 6$. There is one order per crop per week in the default fixture.

For historical week $u \in \{0, \dots, 11\}$, the Data Explorer generator uses

$$Q_{c,u}^{hist} = [27 + 2c + 2(u \bmod 4)a]m_h \left(1 + \frac{gu}{11}\right),$$

where a is pattern_amplitude, m_h is history_multiplier, and g is history_trend. The multiplication in code is equivalently $\text{base} \times \text{multiplier} \times \text{trend}$; the bracket above includes the amplitude-adjusted four-week pattern. With default settings $(a, m_h, g) = (1, 1, 0)$, the four-week pattern repeats exactly. This periodicity is designed, not discovered seasonality.

The bounded playground controls are:

Setting	Minimum	Default	Maximum	Affected records
Historical multiplier	0.50	1.00	1.50	Historical kg
Historical trend	-0.30	0.00	0.30	Historical kg
Pattern amplitude	0.00	1.00	2.00	Historical kg
Order multiplier	0.50	1.00	1.50	Order quantity kg
Price multiplier	0.50	1.00	1.50	SGD/kg

The default resources are 2,560 simultaneous nursery sites, 32 labour hours per week, and SGD 2,200 for the planning horizon. These are declared capacities, not measured availability.

4. Demand and harvest baselines

4.1 Point-in-time filtering

For each crop, history is eligible only when both

$$available_at \leq cutoff \text{ and } week < planning_date.$$

Orders are eligible when $booked_at \leq cutoff$. The Farm contract already rejects future-booked orders, and the forecast repeats this cutoff check. Net booked demand is $quantity_kg - cancelled_kg$.

4.2 EWMA formula

Let $q_{c,1}, \dots, q_{c,n}$ be the eligible historical weekly kilograms for crop c . The implementation initializes the level with the first value and folds subsequent observations:

$$e_{c,1} = q_{c,1}, \quad e_{c,t} = \alpha q_{c,t} + (1 - \alpha)e_{c,t-1}.$$

The same terminal level $e_{c,n}$ is used for every future week. There is no seasonal index, trend extrapolation, public covariate, buyer model, quantile model, or fitted uncertainty distribution. `ForecastSettings` accepts finite floating-point $\alpha \in [0.05, 0.95]$; the default is 0.35.

For crop c and planning week w , let $B_{c,w}$ be all net booked kilograms in that week. Residual forecast demand is

$$R_{c,w} = \max(0, e_{c,n} - B_{c,w})$$

when history exists, and zero otherwise. Confirmed lines retain their actual due dates. The entire residual $R_{c,w}$ is added once, on the final day of the planning week. Thus confirmed bookings are not double-counted, but the weekly residual is given an artificial week-end due date.

When several order lines share a due date, their forecast price is the quantity-weighted mean price. A week-end residual without a same-date booking uses the first order price found elsewhere in that week. A week with residual demand and no booked order line receives a price of zero. This is a defined software behavior, not an economic estimate.

4.3 Harvest projection

The harvest baseline simply emits each active batch's recorded `harvest_date` and `expected Marketable_kg`, with endpoint `fresh Marketable_kg`. No inference is made from plant observations. Future planned harvest is similarly

$$Y_j = A_j y_{r(j)},$$

where A_j is the entire bed area and y_r is the fixed recipe yield. The marketable endpoint is assumed to include survival and packout once; there is no second reduction.

4.4 Versioning and lineage

The forecast stores:

- a hash of the complete farm input;
- a configuration hash of model version plus EWMA settings;
- a numerical-input hash combining those two hashes; and
- `validation_status=demo_only` plus an explicit uncalibrated-uncertainty statement.

The feature builder emits 32 crop/week rows for the reference fixture and records order/history dependency IDs, cutoff, input hash, and version. Saved Data Explorer snapshots additionally freeze the generated farm, generator and forecast settings, forecast payload, reference snapshot, and corresponding hashes. The explorer service rejects a saved snapshot if its schema version, farm contract, settings, dataset hash, frozen forecast hash, or reference hash no longer agrees.

5. Crop-growth representation

FarmTact currently represents a crop cycle as scheduled milestones:



The backend and frontend do not implement the entire diagram as a shared state machine. The planner accounts for sanitation in bed occupancy. The visual farm has no sanitation stage.

The backend farm view computes an existing batch's progress as

$$p(d) = \text{clip}_{[0,1]} \left(\frac{d - \text{sow}}{\text{harvest} - \text{sow}} \right),$$

then labels it nursery before transplant, growing from transplant to the day before harvest, and ready on or after harvest. The interactive preview applies the same linear ratio as a percentage to accepted allocations. It displays nursery, growing, or ready while the preview date is within the allocation's sow-to-harvest interval. For an existing bed without an accepted allocation overlay, it displays the bed as empty at or after its scheduled harvest.

Consequently, “progress” means elapsed calendar fraction between declared dates. It is not measured development, thermal time, accumulated light, biomass, leaf area, stress, quality, disease, or probability of maturity. The crop SVG changes are representative artwork tied to date labels. The interface itself correctly states that the view is a schedule preview and does not claim observed growth.

The scenario system can alter one existing batch by adding 0–14 days to its harvest date and multiplying its declared marketable mass by 50–100%. These are direct counterfactual controls. The main replan endpoint applies a fixed synthetic disruption of +7 days and $\times 0.8$ marketable mass to the first batch. Neither operation derives the delay or loss from weather, sensor data, an image, or a biological model.

6. Candidate schedules and constraints

6.1 Candidate construction

Each candidate j is one full-bed planting for a compatible recipe and harvest date. Candidate harvest dates are every week-end day in the horizon plus any exact booked order due dates. The planner schedules backward:

$$\text{sow}_j = \text{harvest}_j - (\text{nursery}_r + \text{grow}_r),$$

$$\text{transplant}_j = \text{sow}_j + \text{nursery}_r.$$

A candidate is discarded if sowing would precede the planning date, if transplanting would occur before that bed is released by an existing batch and its sanitation period, or if its grow-out/sanitation occupancy intersects a research reservation. New sowing can therefore never be offered as supply for a due date earlier than its recipe lead time.

The decision variable is Boolean:

$$x_j \in \{0, 1\}.$$

There is no fractional bed, partial area, or arbitrary planting-site decision in this version.

6.2 Capacity constraints

For day d , candidate nursery usage is the ceiling of bed area times declared density throughout [sow, transplant). With fixed existing usage N_d^0 ,

$$N_d^0 + \sum_j [A_j \rho_j] I(j \text{ in nursery on } d) x_j \leq N^{max}.$$

Grow-out bed occupancy is inclusive from transplant through harvest + sanitation_days. For each bed b and day d , existing and selected occupancy may not exceed one batch.

Sow labour is charged on the sow-date week at $A_j l_j^{sow}$. Harvest labour is reserved on the harvest-date week using the maximum declared yield factor 1.1 and rounded upward to the next minute:

$$L_j^{harvest} = \lceil 1.1 Y_j l_j^{harvest} \times 60 \rceil / 60.$$

Weekly existing plus candidate labour cannot exceed the farm's declared hours/week.

The CP-SAT cash bound is a conservative planning allowance composed of input cost, labour valued at SGD 12/hour, and a combined SGD 0.45/kg packing/disposal reserve at $1.1 \times \text{yield}$. It is then limited to a policy-specific fraction of declared cash. Existing sunk input cost is not recharged. The post-solve validator also checks the simulated total cost against the full farm cash amount.

The independent allocation validator checks duplicate actions, immutable executed work, dependencies, crop/system compatibility, exact whole-bed area, minimum biological lead time, past sowing, declared yield, reservations, nursery capacity, weekly labour, overlapping bed occupancy, and input cash. Its violations are retained as structured hard-constraint records.

The current data contract does not supply seed availability, water, energy, HVAC, packing capacity, grade, packaging, rotation rules, multiple cuts, partner supply, or crop-specific harvest windows. Although the build specification requires these for a fuller operational planner, the current solver cannot enforce absent fields.

7. CP-SAT objective and strategy families

All policies use the same candidate set and three static scenarios. The policies change four hand-set coefficients:

Policy	Shortage reward/penalty coefficient	Expiry/waste coefficient	New-commitment markup	Cash fraction
Lean	1	9	0.20	0.65

Policy	Shortage reward/penalty coefficient	Expiry/waste coefficient	New-commitment markup	Cash fraction
Balanced	4	5	0.04	0.85
Resilient	25	1	0.00	1.00

For each scenario, crop, and day, the model creates age-bucket inventory, delivery, and remaining-stock variables. Supply consists of eligible opening inventory, scenario-scaled existing harvest, and scenario-scaled selected harvest. Delivery is bounded above by due demand. Delivering a gram earns a coefficient based on selling price, shortage preference, and SGD 0.30/kg packing cost. Stock that reaches its final shelf-life age before the last horizon day incurs the policy waste coefficient plus SGD 0.15/kg disposal cost. Selected candidates incur input and reserved-labour cost, multiplied by the policy's commitment markup.

This is best described as a **policy-weighted utility objective**, not a literal expected contribution-margin maximization. The shortage, waste, and commitment weights are engineering preferences. Confirmed delivery is encouraged but is not a hard service requirement. The three scenario weights are stored as 1/3, but the solver loops over the scenarios equally and does not read the weight field. With the current equal weights, this is numerically equivalent to an unnormalized equal-weight expectation; changing stored weights alone would not change the optimization.

The Resilient policy does not solve a maximin or chance-constrained problem. It applies a much larger shortage coefficient across all three scenarios. “Downside” fill rate and margin are calculated after optimization as the minimum result across the three declared cases.

If CP-SAT returns neither FEASIBLE nor OPTIMAL, the planner invokes backward-scheduling-v1. That fallback orders demand rows by week, descending confirmed kilograms, and crop ID; it greedily adds matching-date candidates that pass the independent validator and remain under simulated cash in the high-yield/low-demand scenario. The fallback primarily covers confirmed demand and is explicitly recorded in solver metadata.

8. Scenario simulation and accounting

The fixed scenario set is:

Scenario	Yield factor	Residual-demand factor	Stored weight
Low yield / high demand	0.85	1.15	1/3
Central	1.00	1.00	1/3
High yield / low demand	1.10	0.90	1/3

Yield factors apply to both existing and newly planned harvest. Demand factors apply only to EWMA residual demand:

$$D_{c,d,s} = B_{c,d} + f_s^D R_{c,d}.$$

Confirmed commitments never shrink in the low-demand scenario. The farm seed is stored as `scenario_seed` and is also used as the CP-SAT random seed. The `scenarios(seed)` function does not use its argument to generate a draw, and `scenario_set_id` hashes the same static three-case list rather than the seed.

The outcome simulator evaluates a frozen daily accounting ledger; it does not advance farm time or execute events. For each dated ledger row, expired lots are disposed first, scheduled harvest is then added, due demand is served crop by crop from the earliest-expiring eligible lot, and remaining stock carries to the next calculated row. Harvest-day product can serve harvest-day demand. A newly harvested lot with shelf life k is eligible on the harvest day and the next $k - 1$ civil dates.

For each day, the simulator asserts

$$opening + harvest - delivery - disposal - closing = 0$$

within 10^{-6} kg. It reports weekly demand, harvest, delivery, and same-date shortfall. Later supply cannot repair an earlier day's shortfall.

Central-scenario financial metrics use:

$$revenue = \sum delivered_{c,d} \times price_{c,d},$$

$$cost = input + 12 \times labour + 0.30 \times delivered + 0.15 \times disposed,$$

$$margin = revenue - cost.$$

`margin_sgd` is therefore a synthetic contribution-like metric under a partial cost fixture. It excludes fixed overhead, tax, depreciation, financing, customer-specific fulfilment cost, and any unmodelled resource cost. `area_m2` is the sum of the areas of all new allocations, so repeated use of one bed in separate cycles is counted repeatedly; it is planted cycle-area, not unique land footprint or peak occupied area.

The accepted `simulated_outcome` is the selected strategy's **central** ledger, metrics, and proposed future work. Acceptance freezes a projection and bookkeeping record. It does not run a farm-time clock, execute the dated ledger as events, mutate batches as work occurs, or create observed biological or commercial outcomes.

9. Isolated scenarios and council-research wrapper

9.1 Scenario Lab

The scenario service deep-copies a frozen farm and supports:

- harvest delay: 0–14 whole days;
- expected batch yield: 50–100% of the frozen value;
- crop demand: 50–150%, applied to both order quantities/cancellations and historical demand for that crop;
- weekly labour: 50–150%; and

- planning cash: 50-150%.

Each branch stores its input snapshot, hash, baseline root, numerical versions, output, same-policy deltas, and zero inference calls. A changed crop-demand control scales both history and orders, so it changes both booked kilograms and the EWMA level. Completed branches are comparable only when their frozen baseline root, baseline hash, baseline forecast settings, and planner version agree. A branch result never replaces the main farm or its accepted worklist.

The backend selects a feasible Balanced branch for simulation when available, otherwise the first feasible policy. This branch acceptance is local and simulation-only.

9.2 V7 council research

The research wrapper accepts inclusive in-horizon bed reservations, a unique list of unconfirmed order IDs, and a whole-number labour ceiling from 50-100% of declared labour. It rejects reservations that intersect an executed batch from transplant through harvest plus sanitation.

Calculation occurs on a deep copy:

1. listed unconfirmed orders are removed from the booked-order set;
2. labour capacity is multiplied by the selected percentage;
3. the ordinary EWMA and three-policy planner run with the reservation windows; and
4. the result is bound to the base-farm hash, numerical-input hash, controls, and council-research-v1 version.

Removing an order does not erase historical demand. Its quantity can reappear partly or fully as EWMA residual demand, which is displayed separately from booked commitments.

The research service freezes each job's farm, input version, controls, model versions, and hash. It permits at most 12 numerical versions per study and only one queued/running research calculation per workspace. A late result for an older version is archived and cannot become current. If a later version runs before the unchanged version-1 reference was calculated, the worker calculates and stores that baseline too. Scripted research dialogue and calculations make no model request; an explicit adviser conversation is a separate action.

10. Evidence and evaluation

10.1 Archived repository evidence

The checked-in numerical evaluation is an archived run for the original fixture hash 002867...f76. It reports six held-out synthetic weeks per crop. EWMA MAE is about 2.383 kg for every crop, compared with 2.667 kg for a last-week naive baseline. Crop WAPE ranges from about 6.50% to 7.77% for EWMA. This result only shows that the implementation can replay a predictable synthetic four-week pattern. It is not evidence of demand accuracy on an SME farm.

The same archived report records feasible, non-optimal four-second solver results for all three policies and daily mass-balanced central ledgers. Those strategy values belong to that exact fixture, software state, time limit, and OR-Tools run. They are examples, not stable benchmarks

or farm-performance claims. Other archived release reports record wider API/browser/persistence checks; see `completion_audit.md` and the V7 implementation report.

The evaluation script performs a growing-window one-step synthetic comparison from history indices 6–11, runs the planner, and compares its strategies with the deterministic backward scheduler on the same declared scenarios. It does not hold out independent farms or crop cycles, estimate harvest error, test calibration, or perform statistical significance analysis.

10.2 Tests executed for this documentation audit

The command and outcome from this audit should be read separately from the archived evidence above. On 11 September 2026, the following current-tree command completed with **35 passed in 57.09 seconds**:

```
/tmp/farmtact-docs-venv/bin/python -m pytest -q \
  tests/numerical/test_planner.py \
  tests/models/test_explorer_numerics.py \
  tests/planner/test_research_constraints.py
```

The targeted suite covers fixture determinism and bounds, EWMA setting/hash behavior, planning feasibility and strategy distinction, impossible early sowing, daily mass balance, cost reconciliation, executed-work locks, independent bed/nursery/labour checks, yield semantics, timeout fallback, zero demand, downside aggregation, research reservations, unconfirmed orders, and a same-environment research repeatability check. That repeatability check is not a guarantee that a time-limited solver returns an identical schedule across environments. The suite does not constitute real-world agronomic validation.

11. Limitations and factual gaps

11.1 Scientific limitations

1. There is no fitted growth, maturity, yield, quality, disease, or weather-response model. Scheduled dates and fixed marketable kg/m² are assumptions.
2. No farm-specific crop cycle or harvest outcome has been used for training, calibration, or external validation.
3. Public weather, forecasts, trade records, and News do not affect numerical demand, yield, timing, or prices.
4. The static three-case uncertainty set has no empirical probability or correlation calibration. Its seed does not produce random scenarios.
5. The EWMA is a constant-level forecast. Calling it seasonal would overstate the implementation.
6. The synthetic evaluation's periodic generator favors a smoothing baseline and cannot establish accuracy under promotions, stockouts, trend breaks, buyer churn, or genuine seasonality.
7. A batch's yield changes only through a direct scenario multiplier; a harvest delay does not mechanistically affect mass, quality, shelf life, or later crop development.

8. The visual progress percentage is linear calendar time. It should never be interpreted as percent biomass, maturity, or completion confidence.

11.2 Planning and accounting limitations

1. Decisions are whole-bed Boolean choices. Partial beds, plant-site quantities, continuous area, mixed beds, and harvest-window choices are unavailable.
2. The current farm contract permits at most one existing batch per bed, regardless of whether historical intervals would be non-overlapping. It does not represent a complete sequence of historical/current commitments on one bed.
3. CP-SAT has age buckets but does not explicitly constrain deliveries to FIFO order. The outcome simulator does use FIFO. This can create a difference between the optimized inventory flow and the reported replay, even though expiry penalties usually favor old-stock use.
4. Final-horizon remaining stock is not assigned a terminal salvage value or a general closing-stock penalty in the CP objective. Results can depend on the arbitrary horizon boundary.
5. The objective is not pure expected margin and does not enforce a minimum service target. It combines revenue-like delivery reward with hand-set shortage, expiry, and commitment preferences.
6. Stored scenario weights are not consumed by the solver. They happen to be equal today.
7. "Resilient" is a high-shortage-penalty policy, not robust optimization, CVaR, maximin protection, or a service guarantee under the low-yield case.
8. The solver cash constraint uses a conservative approximation; reported scenario cost uses actual simulated delivery and disposal. These are related but not identical quantities.
9. Demand is fulfilled in aggregate by crop/date. The system does not attribute delivered kg to individual orders, grades, or customers.
10. Residual demand without any booked line in a week has zero selling price. That understates its possible revenue and should remain an explicit missing-price state.
11. Water, energy, HVAC, seed, packaging, buyer-grade, rotation, multi-cut occupancy, supplier lead time, and partner purchases are specified future constraints but absent from the implemented contract and solver.

11.3 Interface and reproducibility gaps

1. The forecast emits harvest mass as `marketable_kg`, while the conversation reference builder looks for `expected_kg` in harvest forecast rows. Batch-level references still expose `expected Marketable_kg`, but the forecast-harvest mass reference is omitted.
2. The backend marks an existing batch ready on and after harvest, whereas the standalone frontend preview makes that bed empty at and after harvest when no accepted allocation overlay is present. Neither view exposes the sanitation interval that still blocks planner occupancy.
3. The preview slider allows the date at `planning_date + horizon_days`, one day beyond the planner's final indexed day (`horizon_days - 1`).
4. The checked-in numerical report contains an `evaluation_scope` field that the current evaluation script does not write. Regenerating the report from the current script would remove that field unless the script is corrected.

5. The historical integrated-demo script still checks for ten crop profiles, while the current catalogue contains twelve. That script cannot report PASS under its current condition even when the current catalogue is correct.

12. Recommended next numerical iteration

The most useful next iteration is to preserve the transparent baseline while replacing implicit assumptions with measured states and testable contracts.

1. **Unify crop-cycle state.** Define one backend state function covering nursery, grow-out, harvest-ready, harvested, sanitation, and available. Return both scheduled and observed states, and make the UI consume it. Keep progress unknown unless there is an explicitly defined estimator.
2. **Collect farm-specific batch outcomes.** Record cultivar, system/zone, sow/transplant dates, repeated observations, harvest window, gross kg, marketable kg, reject reasons, and independent crop-cycle ID. Require enough independent cycles before fitting a challenger.
3. **Build honest harvest baselines first.** Compare recipe median duration/yield, recent farm/system medians, and a partially pooled model using whole-batch temporal holdouts. Report harvest-date MAE, marketable-yield MAE, interval coverage, and sharpness by crop and system.
4. **Strengthen demand evaluation.** Use rolling-origin multi-horizon backtests, cancellation and availability states, buyer/product hierarchy, and explicit missing-price handling. Compare EWMA with last-week and seasonal-naive baselines before considering gradient boosting or quantile models.
5. **Calibrate scenarios.** Estimate residual blocks from independent cycles and demand periods; preserve shared zone/weather shocks. Store effective sample size and realized coverage. Until then, name the current cases “declared stress cases,” not probabilities.
6. **Align optimization and replay.** Enforce FIFO or prove an equivalent inventory formulation, add terminal stock value/cost, consume declared scenario weights, and expose objective terms in physical/financial units. Test the CP objective against the simulator on small enumerated cases.
7. **Make policy semantics contractual.** Add explicit service targets or risk constraints where desired. If Resilient is intended to protect the worst case, implement maximin, CVaR, or scenario-specific service constraints and report the resulting trade-off.
8. **Expand resources only when measured.** Add seed/seedling lots, packing capacity, water/energy/HVAC, grade, packaging, rotation, repeated cuts, and supplier quotes as versioned inputs with units and availability times. Do not insert placeholder capacity as if observed.
9. **Separate observed outcomes from simulated ledgers.** Store actual harvest/delivery records independently, reconcile forecast error versus execution shortfall, and prohibit a simulated worklist from creating actual observations.
10. **Create a reproducible evaluation manifest.** Pin fixture hash, dependency versions, solver parameters, generated report schema, command, timestamp, and git commit. Fail report generation when the archived schema and generator disagree.

Any learned challenger should remain demo_only until it beats simple baselines on temporally and biologically independent farm cycles, passes leakage tests, and has uncertainty calibration adequate for its stated use. Even then, the optimizer and its constraints remain separate numerical systems; model accuracy alone does not validate a planting decision.

Game backend and state machines


Historical baseline: This chapter records the pre-remediation audit of source a96025e and published v7. Its implementation findings and measured results describe that baseline. Read the V8 remediation report for the current simulation engine, synthetic model evaluation, Council contracts, tests and remaining limitations. Source links below are navigation aids into the maintained repository; use the frozen baseline to reproduce the original inspection.

This chapter audits the implemented FarmTact game and simulation backend as of 11 September 2026. It follows persisted state from anonymous workspace creation through planning, automatic simulation acceptance, disruption replanning, scenario branches, quests, conversations and the playable Council research study. It also separates backend facts from browser-derived presentation.

The central finding is simple: FarmTact has a durable **planning and experiment game**, but it does not have a farm-time execution engine. Numerical routines calculate dated ledgers over a frozen horizon. Missions, branches, quest badges and research choices persist workflow progress. None of them advances the farm's `planning_date`, consumes an accepted allocation, changes crop growth stage day by day, fulfils an individual order, or records completion of sowing, transplanting and harvest work. The only implemented main-farm mutation after bootstrap/import is a deliberately synthetic disruption that edits the first batch and creates a new farm version (replan route).

For the numerical equations and their scientific limits, see Numerical models and growth. For provider routing and claim validation, see AI provider and Council. This chapter concentrates on authority, transitions, version binding, concurrency and recovery. It is part of the main scientific implementation report; remediation priorities are consolidated in the next-iteration gap register.

The implemented game in one view



Orange marks the mutable farm/research state and its explicit update transitions. Of those transitions, replan changes the main farm and apply changes only an isolated research input. ACCEPTED_FOR_SIMULATION, the scenario acceptance record, quest progress, a research choice and advisor messages are durable decisions or game progress. They are not agricultural execution. Scenario and research snapshots are isolated copies and never replace the main farm (scenario isolation test, research calculation).

What “simulation” means here


The same word covers several different mechanisms. They should not be presented as interchangeable.

Mechanism	What is actually calculated or persisted	Does farm time advance?	Does it mutate the main farm?
Planner simulation	Daily FIFO inventory ledger and aggregate crop/date demand, harvest, delivery, expiry, stock, cost and margin over the frozen horizon (simulator)	No	No
Automatic mission acceptance	Service records a chosen eligible strategy, hashes, policy version and a copy of its unexecuted allocations as <code>simulated_outcome.work</code> (acceptance)	No	No
Worklist	CSV rendering of the accepted strategy, guarded by current farm version and labelled <code>SIMULATION_ONLY</code> (worklist)	No	No
Replan disruption	Adds seven days and applies a 0.8 yield factor to the first current batch; persists a new farm version and child mission (replan)	No	Yes , but only as a synthetic observation update
Scenario branch	Applies bounded controls to a copy, calculates baseline and branch, and records a simulation-only acceptance if an eligible strategy exists (scenario execution)	No	No
Quest	Records that a matching branch completed and that the user explicitly inspected it; derives two fixed badge labels (quest progress)	No	No
Council research	Applies typed edits to a private study input version, calculates three policies and records an explicit simulation choice (actions)	No	No

Mechanism	What is actually calculated or persisted	Does farm time advance?	Does it mutate the main farm?
Conversation action	Stores a hypothesis_only scenario control proposal after local validation; the user must explicitly create a branch (reply validation)	No	No
Browser mission	Locally screens dated orders against eligible harvest and inventory and stores the selected mission in edition-scoped sessionStorage (derivation)	No	No
Audio	Plays an ephemeral client effect when a result becomes terminal; no sound event is persisted (audio result)	No	No

The planner steps through calendar dates inside one function call. “Day 12” is a row in a hypothetical ledger, not a scheduled backend wake-up. Worker waiting, advisor calls, scripted dialogue and UI animation do not change the biological calendar.

Persistence and ownership



The session cookie maps to a tenant through a SHA-256 hash. Authentication accepts only tenants created during the previous 24 hours; the cookie is HTTP-only, same-site strict and secure on HTTPS (session store, bootstrap). Rows themselves are not deleted when authentication expires.

Farm versions are append-only by (tenant_id, version). Version allocation locks the tenant row, calculates $\max(\text{version})+1$, hashes the whole payload and inserts a new row (farm persistence). Planning runs, scenario branches, conversations and research sessions then freeze full JSON snapshots. Most relationships also have tenant-aware foreign keys; conversation messages, requests and events have their own monotonically increasing sequences (conversation schema).

There is no domain table for executed tasks, crop-stage observations, inventory movements, order fulfilment, simulated clock ticks or game turns. work is nested JSON in a run, and quest progress is a compact per-tenant/per-quest JSON record.

HTTP surface and transition authority

All routes below require the tenant session except bootstrap and the fixed health endpoint. Mutating requests pass same-origin, body-size and admission checks before route parsing (middleware, admission).

Main farm and planning missions

Method and route	Transition / result	Guard and idempotency
GET /api/v1/bootstrap	Creates tenant + synthetic farm v1 when no valid cookie; otherwise returns latest farm view and latest run	New-session IP/global limits; cookie lifetime 24 h
POST /api/v1/imports	Validates exactly one fixture or Farm, always appends a farm version; after a terminal latest run also creates a local refresh mission	No idempotency key ; tenant transaction serializes the version
POST /api/v1/planning-runs	Freezes latest farm and creates CREATED mission	Required key, tenant-wide key/hash equality, and at most one CREATED/RUNNING planning run
GET /api/v1/planning-runs/{id}	Returns public run and current acceptance_stale flag	Tenant ownership
GET /api/v1/planning-runs/{id}/events	SSE after query/header cursor; stops after terminal run	Per-IP open and active-stream limits
POST /api/v1/planning-runs/{id}/cancel	Sets cancel_requested only while created/running	Cooperative; response always says cancel_requested even if already terminal
GET /api/v1/planning-runs/{id}/replay	Returns stored run/events with replay labels	Run must be terminal; no inference
POST /api/v1/planning-runs/{id}/replan	Writes the fixed synthetic first-batch disruption, new farm version and child mission	Parent must be accepted, current farm version must equal parent input, required tenant-wide run key; body/key conflict is 409
POST /api/v1/strategies/{id}/accept-for-simulation	Always 403	Browser cannot accept; backend policy owns acceptance
GET /api/v1/planning-runs/{id}/worklist.csv	Renders accepted allocation list	Accepted strategy must exist and latest farm version must still equal run input version
GET /api/v1/reports/{id}	Alias of public run	Tenant ownership

Method and route	Transition / result	Guard and idempotency
GET /api/v1/strategies/{id}	Searches all owned run JSON for the strategy	Tenant ownership; no run/version parameter

The definitions are in the main route installer. Import's lack of idempotency is especially relevant: an identical retry creates a new version, makes accepted worklists stale and may enqueue another refresh.

Scenario branches and quests

Method and route	Transition / result	Guard and idempotency
GET /api/v1/scenarios	Returns every owned full branch, newest first in Python	No pagination or branch count limit
POST /api/v1/scenarios	Freezes one source and creates DRAFT branch	Required tenant-wide key; source is explorer or parent/conversation; completed parent required; quest must change its named assumption
POST /api/v1/scenarios/{id}/run	DRAFT -> QUEUED; later worker claims RUNNING	First run_key binds the branch forever; same key reuses, any other key gets 409
GET /api/v1/scenarios/{id}	Returns full saved branch/results	Tenant ownership
GET /api/v1/scenarios/compare?ids=...	Returns one to three completed branches and shared baseline	Distinct IDs; same comparison root, baseline hash, baseline forecast settings and calculation version
GET /api/v1/quests	Combines four fixed quest definitions with progress rows	No mutation
POST /api/v1/quests/{id}/inspect	Adds branch ID to inspected set, marks quest completed and grants fixed badges	Scenario must be completed and have the matching quest ID; semantically deduplicated, no request key

See scenario request and controls and scenario routes. Controls are bounded: batch delay 0..14 days, marketable yield 50..100%, demand, labour and cash 50..150%. Delay/yield require a known batch; changed demand requires one of the four fixture recipe crops. Demand control scales both historical ordered kg and all booked/cancelled order quantities for that crop (control application). A child starts from its parent's already-modified input, so percentages compound across branches.

Persistent conversations

Method and route	Transition / result	Guard and idempotency
GET /api/v1/conversations/advisors	Returns current seven-person roster	Read only

Method and route	Transition / result	Guard and idempotency
GET /api/v1/conversations	Lists public conversation metadata without messages	Tenant ownership; no pagination
POST /api/v1/conversations	Freezes farm, completed scenario or current completed research calculation; creates OPEN	Required tenant-wide conversation key; same body reuses
GET /api/v1/conversations/{id}	Public state, all messages, frozen tool/evidence context	Tenant ownership
GET /api/v1/conversations/{id}/replay	Same messages with replay labels	Read only, no provider call
POST .../{id}/messages	Appends user message + direct one-role QUEUED request atomically	Required key scoped to conversation; one active request per conversation
POST .../{id}/invite	Queues invited role followed by author of referenced advisor message	Must reply to an advisor and invite a different advisor
POST .../{id}/council	Queues seven ordered roles	Optional reply must belong to same conversation
GET .../{id}/events	JSON cursor page or SSE	Cursor nonnegative; stream stops when latest request is terminal

The route implementation is in `conversations`. A farm conversation always freezes the latest farm; its optional `snapshot_id` checks the farm ID but cannot select an old farm version. Scenario conversations require `COMPLETED`. Research conversations require a completed result whose version is the current research input version, although the error text says “complete and select” and the implementation does not require chosen (snapshot freezer).

Conversation request admission counts one user message plus all requested roles against a hard 120-message conversation limit. User message and job creation are one locked transaction, with (conversation, tenant, key) uniqueness (request persistence).


Playable Council research

Method and route	Transition / result	Guard and idempotency
GET /api/v1/council-research	Lists up to 30 session summaries	Query limit occurs before Python sort; no explicit SQL order
POST /api/v1/council-research	Creates isolated synthetic study v1	Required tenant-wide session key; 30-session tenant cap
GET /api/v1/council-research/report	Returns bounded local report excerpts, fixed citations and screenshot URLs	Read only; paths are fixed in code

Method and route	Transition / result	Guard and idempotency
GET /api/v1/council-research/{id}	Returns complete mutable session JSON	Tenant ownership
POST .../{id}/actions	Applies one of 12 action types and normally increments revision	Required key unique across all research actions in the tenant ; body revision must equal current; 400 actions/session

The route implementation and exact action switch are in Council research routes. say, select, propose, apply, discard, challenge, resolve, choose, configure, stop, next and run are one endpoint because the request model is a tagged action with optional fields. Unknown action names and extra fields fail schema validation; action-specific guards validate required context. Known optional fields are not a fully discriminated per-action schema and may be unused on a different action (action schema).

Main planning state machine



Creation freezes the whole farm payload and its hash, source cards, evidence file hash, market summary, news context, Council version and parent/disruption metadata before queueing (mission creation). Store.create_run locks the tenant, reuses an equal key/hash, rejects a changed body under the same key, and permits only one active planning mission for the tenant (run persistence).

The worker calculates all three strategies locally first. Strategy eligibility is exactly:


1. stored strategy status == FEASIBLE;
2. a fresh validate_allocations(frozen_farm, allocations) returns no violations;
3. choose Balanced when present, otherwise the first eligible strategy in the planner's current list order;
4. if Council claims exist, all seven roles must appear in exact roster order, all must have status=validated, and all must recommend proceed_simulation (Council gate);

- immediately before acceptance, the latest farm must still have the same version **and** content hash.

An accepted record binds input version/hash, strategy ID/hash, validation report hash, policy version, service actor and `simulation_only=true`. The public manual acceptance endpoint always returns 403, so no browser or advisor can bypass that gate.

There is a significant asymmetric rule. If a requested Council fails or is blocked before producing any claim, `claims=[]`; the evidence gate uses numerical-baseline and the strategy can still be accepted. If even one claim was persisted before failure, the incomplete seven-role set produces review issues and acceptance is withheld. This is intentional in the current test contract (mission gate test), but the API has no explicit `council_required_for_acceptance` policy field.


Planning sequence, stale input and replan



Stale-input refresh intentionally disables the Council, avoiding automatic repeat paid inference. Process restart is different: a planning job recorded as running becomes FAILED and is never replayed because the system cannot know whether a paid request completed (restart rule).

Replan is only available from `ACCEPTED_FOR_SIMULATION`. It requires the parent farm version still be current, preserves existing sow/transplant dates, changes the first batch's expected harvest date/yield, appends a farm version and creates a child. Its request schema exposes only `crop_delay`; callers cannot choose the batch, delay or yield factor (schemas). This is a fixed demo transition rather than a general disruption engine.

Scenario branches, comparisons and quests



A scenario source is one of:

- a saved explorer snapshot;
- a completed parent scenario's modified input;
- a conversation's frozen snapshot, with a scenario conversation resolving back to its scenario parent;
- otherwise the current farm.

The branch stores its modified input, root baseline, hashes, optional generator and forecast settings, news context and calculation version. The worker refuses to substitute a current forecast/planner implementation when the frozen version does not equal the installed version (version check). On success it calculates same-policy baseline/branch deltas, validates allocations again and writes its own simulation-only acceptance. A COMPLETED branch may still have `simulation_status=NO_FEASIBLE_PLAN`; completion means the experiment calculated, not that the requested farm plan was feasible.

Comparison eligibility is strict enough to prevent mixing unrelated roots: one to three distinct, completed branches must share comparison root, baseline hash, baseline forecast configuration and calculation version. Child control factors apply to the parent's already changed snapshot. For example, a 90% cash child of a 90% cash parent represents 81% of the root cash, not another branch at 90% of root.

Quest state has no prerequisite graph or scoring engine. Four definitions are hard-coded. Creating a matching changed branch is permitted in any order; worker completion adds it to `experiment_ids`; explicit inspect adds it to `inspected_ids`, changes status, and assigns Experiment explorer plus Tradeoff discovered. Even infeasible results count as completed learning (infeasible quest test).

`affected_bed_ids` compares allocation hashes grouped by bed across every policy. `affected_deliveries` identifies changed order inputs or orders on dates where aggregate ledger demand/delivery changed. It explicitly is not per-order fulfilment attribution (impact calculation).

How the three Councils work


1. Mission Council: seven provider-backed findings

The current roster is Ravi/demand, Hana/weather, Idris/market, Mei/production, Lina/supply chain, Ben/profit and Asha/planning chair. The mission Council runs one structured provider request per role with at most two shared schema repairs. Each role receives the same frozen numerical strategy summary and bounded reference map. Only the planning chair receives the claims already produced by earlier roles (Council loop).

This is sequential multi-role review, not a deliberative debate. Specialists do not see or rebut other specialists. There are no implemented challenge rounds, votes, tool calls initiated by roles, or user intervention points. The chair synthesizes prior claims once. Local validation checks schema, role, known reference membership and literal numeric correspondence; it does not prove that qualitative prose follows from the cited item.

Council output may block automatic simulation acceptance but cannot alter a strategy or execute work. All seven findings must validate and recommend proceed. Any dissent, rejection, duplicate, reorder or partial set with at least one claim withholds acceptance.

2. Conversation Council: a persisted reply chain




Direct mode has one role. Invite mode has two: the invited advisor responds to a specific advisor message, then that original advisor responds to the invited reply. Council mode produces seven turns in roster order; each turn after the first replies to the immediately preceding generated message. The planning chair must use the final conclusion relationship to be marked a planner conclusion (conversation execution).

Every generated message is stored before the next provider call, so a later failure preserves a useful partial transcript. Restart marks uncertain running requests INTERRUPTED, conversation PARTIAL, and never automatically retries possibly billed work (conversation recovery). A new explicit request is allowed after a terminal partial/failed/blocked request.

COMPLETED is transport/workflow completion. A reply can be stored with `validation_status=unsupported`, blocked proposed actions and validation errors, yet the request still becomes COMPLETED. Thus conversation status does not mean that the Council's conclusion passed a decision gate. Conversation proposals are always hypotheses and require an explicit user-created scenario branch.

3. Council research: scripted study and local tool



The study begins from a new synthetic fixture plus an 18 kg lettuce order due in 48 days. It records `dialogue_mode=scripted_research`, `inference_calls=0` and `operational_execution=false` (session creation). Its text interpreter recognizes a small set of reservation, order-status, labour and challenge phrases. It never executes free text. A proposal has a `base_version`; only explicit apply validates the controls and creates the next research input version (proposal/application).

Each input version can have only one research job. One queued/running research job is allowed per tenant across all studies, and one study stores at most 12 result versions. A restart requeues a running job because calculation is local and has no paid side effect. When a later version calculates before v1 has a completed baseline, the worker reconstructs v1 from the original controls. A result that finishes after inputs changed is archived and cannot overwrite current input state (research worker).

After a result, the backend creates fixed scripted contributions only for control categories that changed, plus a planner message. continuous and checkpoints change how pending_turns are

revealed; they do not change numerical logic. The presentation concept, selection mode and animation setting are also game/UI configuration only.

A challenge stores one mutable object and clears any choice. Rainfall gets a specific supported correction: outdoor rainfall is context, not a sheltered hydroponic yield input. Other challenge categories have no verified correction. resolve=evidence leaves status unchanged; corrected is allowed only for rainfall; unresolved and reject continue to block choose. choose requires a completed result for the current version, a feasible violation-free selected policy, no proposal, and no challenge whose status differs from corrected (challenge, choose). It writes a simulation-only research choice and never affects the main farm.

The optional “actual DeepSeek interpretation” in this screen creates a separate research-snapshot conversation and sends a direct one-role message. It is not the scripted dialogue and it is not a full Council (UI boundary).

Worker concurrency, admission and recovery

One daemon thread polls all queues every 300 ms. In each pass it exhausts pending planning runs, then scenarios, research calculations and conversations, in that fixed order. Each job runs synchronously on the same thread (worker loop). This provides simple atomic claiming but no fairness, parallelism, per-tenant scheduling or isolation between paid provider latency and local numerical play.

The application-level concurrency rules are:

Work type	Active limit	Claim rule	Restart behavior
Planning mission	One CREATED/RUNNING per tenant	Conditional status update	Running -> FAILED; explicit new mission required
Scenario	No tenant-wide active cap	Conditional QUEUED -> RUNNING	Running -> QUEUED
Research calculation	One queued/running per tenant	Conditional job status update	Running -> QUEUED
Conversation	One queued/running request per conversation	Locked request row	Running -> INTERRUPTED; conversation -> PARTIAL
Provider calls	One at a time inside each synchronous job	Gateway/run budget	Never automatically repeated after uncertain planning/conversation interruption

Planning and conversation inference attempts have durable request admission: 6/IP/minute, 20/IP/hour and 12/tenant/hour. All writes also consume 60/IP/minute; all API requests consume 600/IP/minute and 3,000 global/minute. The shared provider reservation ceiling is 48 calls per UTC day. SSE openings have their own 12/IP/minute rule plus four active streams/IP, 32 globally and a 360-second deadline (rules, stream slots). Scenario and research calculation actions do not consume the AI-specific admission bucket because they make no provider calls.

Cancellation exists only for main planning. It is cooperative: a queued mission will stop before planning; a running mission checks after local planning and the Council receives a callback it can check between bounded work. There is no external kill of a current solver or provider HTTP request. Scenarios, research jobs and conversations have no cancel route. Research stop clears pending scripted turns but explicitly does not cancel a calculation.

Browser-derived game state and audio

The UI derives several useful views that are not backend domain state:

- `DecisionMission` groups future booked orders by crop/date, subtracts harvests and inventory that appear eligible by date/shelf life, and selects the earliest positive provisional gap. It warns that it does not allocate those lots across earlier orders. It is edition-scoped `sessionStorage`, not a server entity.
- `NumericalPerspectives` renders seven named perspectives deterministically from scenario comparison metrics and violations. These cards are not provider messages even though they use the same advisor personas.
- selected scenario, explorer filters and current research/actual-conversation pointers live in edition-scoped session storage. The durable objects remain in the backend; the selection/navigation state does not.
- main mission events use SSE plus 2.5-second polling fallback. Scenario and research screens poll. Conversation supports JSON cursors and SSE.

Audio defaults are edition-specific and preferences are stored in local storage, but every visit begins inactive until direct user activation. Background music pauses while hidden. Effects are rate-limited in memory and stop each other so only one plays at a time (audio state, playback).

Main-run sound outcome is derived from terminal status and the accepted, Balanced or first strategy's `shortfall_kg`; scenario sound uses the same fallback order (sound mapping). A run only becomes audible if this browser instance started/replanned it and kept its ID in an in-memory set. `playSimulationResult` suppresses duplicate sounds in that page lifetime. These cues are presentation feedback, not durable evidence that a backend event was consumed, and are appropriately absent from replay provenance.

Gaps and acceptance tests

The IDs below are stable references for remediation. Severity reflects decision integrity and reliability, not visual polish.

ID	Severity	Implemented gap and consequence	Recommended acceptance test
GAME-01	Critical	There is no simulated clock or execution state machine. Acceptance and worklist export cannot record task start/completion, inventory movement, batch stage, order fulfilment or observed outcome. Gamified progress can therefore be mistaken for farm progression.	Given an accepted version, advance a synthetic clock through sow/transplant/harvest/delivery; assert idempotent task events update new versioned state, conserve mass, never touch operational integrations, and replay exactly. Until built, assert every UI/API result retains <code>simulation_only</code> and no “completed farm work” wording.
GAME-02	High	One synchronous worker causes head-of-line blocking across all tenants and job types. A slow Council/provider request can delay unrelated scenario and research calculations for up to its bounded wall time.	Queue a blocked 300-second provider job and local jobs from another tenant; require bounded local-job start latency and independent worker pools/leases. Kill a worker mid-claim and prove each job class follows its documented recovery rule once.
GAME-03	High	Requested Council acceptance is asymmetric: zero claims after block/failure passes as numerical baseline, while one-to-six claims withhold. The request has no field stating whether Council is advisory or required.	Parameterize missing credential, budget block, failure before first claim, failure after each claim and seven valid claims. Assert the chosen explicit policy (advisory or required) and persist it in the acceptance record.
GAME-04	High	Neither provider-backed Council implements the build specification's two challenge rounds. Mission specialists are independent; only the chair sees prior findings. Conversation Council is a linear reply chain. Scripted research has one user challenge object but no agent debate.	Record the exact expected topology: participant visibility, round number, challenge target, rebuttal and final synthesis. Test request ceilings and restart behavior for every round, or change the specification/UI to “seven sequential reviews plus chair synthesis.”

ID	Severity	Implemented gap and consequence	Recommended acceptance test
GAME-05	High	Research challenge version binding is incomplete. challenge.input_version is recorded, but apply leaves the challenge in place and resolve/choose do not require its version to match current inputs. A correction to an old version can authorize choosing a later version. A new challenge also overwrites the prior one.	Challenge v1, apply v2, resolve the old challenge and attempt choose v2; require rejection or explicit challenge migration/revalidation. Preserve multiple challenge records with version and resolution evidence.
GAME-06	High	API idempotency exists for most expensive transitions, but browser helpers generate a fresh UUID on every call and do not retain it across an uncertain network retry. Import has no key at all. Duplicate imports append versions and stale accepted work; a lost mission response can be retried as a second run after the first finishes.	Simulate “server commits, response lost” for import, mission, scenario creation/run, conversation send and research action. Reload/retry from UI and assert exactly one logical mutation/provider job using a persisted operation key.
GAME-07	High	Conversation request COMPLETED only means all turns were stored. Every reply, including the final chair, may be unsupported; the conversation returns to OPEN and looks terminal-successful. No separate Council validity/result state exists.	Make one and all seven replies fail semantic/reference validation. Assert transport status and decision-validity status are distinct, the final result is visibly withheld, and no proposed action can be applied implicitly.
GAME-08	High	The scenario “run once” key is permanent. A failed calculation cannot be retried on the same immutable branch, including after a transient local failure. The UI must create a new branch, losing a direct retry/recovery relation.	Inject a transient calculation failure after claim; retry same run key and a dedicated retry endpoint. Require either safe same-input local recomputation with attempt history or a child retry that explicitly links to the failed branch.

ID	Severity	Implemented gap and consequence	Recommended acceptance test
GAME-09	Medium	Frozen calculation version support is equality with the currently installed code, not a runner registry/container reference. Old uncalculated branches or jobs fail after deployment; exact source commit is not embedded in every game object.	Create DRAFT/QUEUED objects on version N, deploy N+1, then prove either version-N execution remains available or a migration creates a separately identified result without overwriting frozen provenance.
GAME-10	Medium	Research messages and events are mutable JSON arrays truncated to 120 and 240. The separate action table retains only key/hash identity, not action payload, response, resulting revision or challenge history. Long studies lose early audit detail.	Execute 400 actions and reconstruct every state transition from append-only records. Assert old messages/events remain queryable with pagination and action retries return the original response revision.
GAME-11	Medium	Scenario storage/listing has no per-tenant branch cap or pagination. Conversation listing also lacks pagination. Expired tenant data is not reclaimed by the application. A 24-hour anonymous session can create substantial persistent JSON and every scenario list returns it all.	Fill to documented quotas, verify 429 before expensive work, paginate with stable cursors, expire a tenant and prove retention/deletion policy without affecting active edition data.
GAME-12	Medium	Only planning can be cancelled, and its cancellation cannot interrupt current solver/provider work. Scenario, research and conversation jobs lack cancel endpoints.	Cancel before claim, during solver, between Council turns and during provider I/O for every job type. Assert bounded terminal latency, no new paid request after cancel, unused reservation reconciliation and unambiguous terminal status.

ID	Severity	Implemented gap and consequence	Recommended acceptance test
GAME-13	Medium	Main replan is hard-coded to the first batch, +7 days and $\times 0.8$ yield. The request's disruption enum does not bind a target or magnitude, so it is a demo rather than a general game action.	Submit typed disruptions for explicit batch/version with bounds; reject stale/unknown targets; prove executed commitments remain identical; hash the observation and deterministic child input.
GAME-14	Medium	Strategy fallback is Balanced, then first eligible in current list order. No explicit ranking result explains why the fallback won if Balanced is absent.	Permute strategy order with Balanced infeasible; require an explicit deterministic rank/tie-break record and stable selected ID across permutations.
GAME-15	Medium	Scenario affected deliveries and the planner ledger are aggregate crop/date flows, not per-order allocations. Quest completion can establish that tradeoffs were inspected, but not which buyer order is covered.	Add two same-crop orders on the same/earlier dates and competing lots. Require an allocation ledger with lot-to-order quantities, expiry, shortfall and mass balance before displaying order-level coverage.
GAME-16	Medium	Browser mission, persona perspectives, navigation pointers and sound completion memory are local derived state. They are not revisioned backend progress and do not synchronize across tabs/devices; sound does not acknowledge event consumption.	Change farm/branch in another tab and verify derived mission invalidates by hash. Reload during a terminal event and assert UI truth stays correct with or without sound. Label deterministic persona cards separately from actual Council messages.
GAME-17	Low	Cancel returns <code>{"status": "cancel_request"}</code> for terminal runs even though it changed nothing. This can mislead client state.	Cancel every terminal status and require an explicit no-op/409 plus unchanged persisted record; cancel active run and return the actual current cancellation state.

ID	Severity	Implemented gap and consequence	Recommended acceptance test
GAME-18	Medium	One conversation contract test assumes cached public source D04 is present, while <code>source_views()</code> legitimately returns an empty list when the ignored public cache is absent. A fresh offline checkout therefore fails a game test for environmental data availability rather than conversation behavior.	Seed a minimal versioned public-context fixture in the test or assert a typed unavailable source state. Run the same test from a clean Git checkout with network disabled and with a populated cache; conversation-state assertions must be identical.

Specification corrections to keep synchronized

These observations refer to the pre-audit specifications. The documentation update now incorporates the corrections; GAME implementation gaps remain future work.

These are documentation corrections, not claims that the code should necessarily grow to match every older requirement.

1. Build specification lines 527, 717 and 833 describe two challenge rounds. Current mission and conversation Councils have no challenge-round state. Specify seven sequential reviews and a final chair synthesis, or design the missing versioned rounds.
2. Build specification line 523 lists terminal states including `MISSING_INPUT`, `SOURCE_UNAVAILABLE` and `BUDGET_EXCEEDED`; the planning-run state machine uses warnings/council status for those conditions and terminal run states `STALE_INPUT` and `ACCEPTED_FOR_SIMULATION`, which that list omits.
3. Build specification lines 630 and 693, and DeepSeek runtime specification DS-20, refer to idempotent simulated work. The implementation has idempotent plan/branch/job creation and a worklist, but no simulated work execution API or work-event table. Narrow the claim or implement GAME-01.
4. Build specification line 735 calls the CSV a practical worklist. Clarify that it is a stale-version-guarded simulation projection with no completion/update path.
5. Current product copy should reserve “Council” qualifiers: **mission Council** is seven provider calls, **conversation Council** is a seven-turn provider chain, and **Council research** is scripted with zero provider calls unless the user opens and sends the separate optional actual-advisor question.
6. “Plant growth simulation” should mean dated yield endpoints and inventory flow unless and until GAME-01 and an empirical growth model exist. No backend transition models daily morphology, environment response or stage observation.

Reproducible verification boundary

The existing automated suite covers mission acceptance/replay/replan and stale worklists (security API tests), scenario isolation, branching, comparison, quests and restart recovery (scenario tests), conversation idempotency, reply graphs, validation, partial failure, tenant isolation and PostgreSQL persistence (conversation tests), and research controls/session behavior (research tests). PostgreSQL tests also cover concurrent planning/replan idempotency and event sequencing (concurrency tests).

The coordinated fresh offline run used:

```
pytest tests/council_research/test_sessions.py \  
  tests/council_research/test_independent_review.py \  
  tests/gameplay/test_scenarios.py \  
  tests/gameplay/test_conversations.py -k 'not real_postgres'
```

It reported **66 passed, 1 failed, 1 deselected in 36.06 seconds**. The failure was `test_seven_current_advisors_and_frozen_create_idempotency` at its assertion that `source:D04.summary` exists. The ignored public-data cache was absent; no source fetch was run. This is GAME-18, not evidence that conversation idempotency failed. No database service, external provider or inference was started for this chapter.

The highest-value next verification is a deterministic lifecycle test harness that queues all four job classes, injects response loss and process restarts at each claim/commit boundary, and proves the state diagrams above. That would validate workflow correctness. It would not establish agronomic validity, real farm performance or safe operational execution.

System, data, and evidence architecture

Historical baseline: This chapter records the pre-remediation audit of source a96025e and published v7. Its implementation findings and measured results describe that baseline. Read the V8 remediation report for the current simulation engine, synthetic model evaluation, Council contracts, tests and remaining limitations. Source links below are navigation aids into the maintained repository; use the frozen baseline to reproduce the original inspection.

Audit date: 11 September 2026 UTC

Repository baseline: commit a96025e (the working tree may contain the documentation changes prepared for the next release)

Published application covered: immutable v7, with v1–v6 retained as separate editions

This chapter describes what the FarmTact source and recorded release evidence implement. It does not promote roadmap items in the v1.2 specification to shipped capabilities. FarmTact is a synthetic planning demonstration: its farm records, recipes, demand histories, schedules, yields, prices, costs, scenarios, and simulated outcomes are fictional. Real public weather, forecast, trade, news, and scholarly records supply attributed context. They do **not** automatically change demand, yield, recipe, price, or solver inputs.


The key separation is therefore:

Real public information can help a person or advisor ask a better question. The numerical planner only changes when a validated, tenant-owned farm snapshot or an explicit bounded scenario changes.

The implementation evidence for that statement is the empty `public_features_used` array in the committed feature manifest, the planner's Farm-only input in `engine.py`, the News context's explicit `numerical_effect="none"` in `news.py`, and the bounded scenario controls in `scenarios.py`.

Implemented system

FarmTact is a React/TypeScript client served by a Python FastAPI application. PostgreSQL stores anonymous tenant state, immutable farm versions, planning jobs, scenario branches, conversations, explorer snapshots, and research sessions. A small in-process worker claims persisted jobs. Local Python code performs forecasting, constraint solving, simulation, comparison, and validation. The DeepSeek boundary is optional for explicit advisor interpretation and is not the source of numerical values.



Primary implementation links:

- Application assembly, lifecycle, API boundaries, worker, and bootstrap: `services/api/app.py`
- Persistent core store and tenant predicates: `services/api/store.py`
- Browser API adapter: `apps/web/src/lib/api.ts`
- Board and farm-world views: `Board.tsx`, `World.tsx`
- Data, crop, and evidence views: `DataExplorer.tsx`, `PublicExplorer.tsx`, `Rooms.tsx`
- Public API response contracts: `services/api/view_contracts.py`

The client uses edition-scoped browser storage for presentation preferences and continuation pointers such as the selected scenario or explorer snapshot, the current research/advisor identifiers, and a cached decision-journey view. Authoritative farm, plan, scenario, conversation, and research records remain server-side and are reloaded through owned APIs. API calls use a random session cookie; the database stores its SHA-256 hash rather than the bearer value. The session is an anonymous workspace identity, not a verified user account.

Data domains and authority

Domain	Current contents	Origin	Can directly drive the planner?	Persistence and authority
Private farm snapshot	16 beds, four recipes, batches, orders, demand history, opening inventory, labour, nursery capacity, and cash	Synthetic fixture or schema-valid JSON import; current contract still restricts all records to <code>synthetic_demo</code>	Yes	Versioned per tenant in <code>farm_versions</code> ; this is the planner's numerical authority
Derived demand and harvest baseline	Confirmed orders, residual weekly EWMA, scheduled harvests	Deterministic local calculation over the frozen farm snapshot	Yes	Frozen into planning, explorer, and scenario results with <code>model/configuration/input</code> hashes


Domain	Current contents	Origin	Can directly drive the planner?	Persistence and authority
Plan and simulation	Lean, Balanced, Resilient allocations; daily inventory ledger; scenario metrics	Deterministic CP-SAT plus local simulation	Yes	Persisted with calculation version, strategy IDs, constraints, hashes, and simulation-only acceptance
Scenario/playground	Delay, yield, demand, labour, cash, bed-reservation, order-status, and EWMA-alpha changes	Explicit bounded synthetic controls	Yes, inside the frozen branch only	Tenant-owned immutable roots and versioned results; no automatic main-farm mutation
Public observations	NEA station weather, NEA forecasts, SingStat trade volume, NASA POWER history	Real public APIs	No	Content-addressed raw snapshots and normalized cache; displayed with source/time/licence limits
News	RSS metadata from four allowlisted feeds; three additional sources visibly unconnected or blocked	Real public headline metadata	No	Separate bounded cache; copied into a decision snapshot at its cutoff
Crop knowledge	Twelve crop concepts linked to 24 reviewed publication records	Curated public/government and scholarly metadata	No	Version-controlled JSON; recipe parameters remain separate synthetic engineering assumptions
Community/market reactions	Validated supplied-record adapter with no connected feed	No runtime records in the published demonstration	No	Status is <code>not_connected</code> ; it cannot manufacture demand or sentiment
AI advisor prose	Qualitative interpretation of frozen numerical and evidence references	Actual DeepSeek only when explicitly invoked; scripted v7 dialogue otherwise	No	Persisted and locally validated; unsupported text cannot authorize a change

The private contract in `contracts.py` is intentionally narrower than the future operational schema proposed by the build specification. It supports four synthetic crops—caixin, pak choi, kailan, and lettuce—and forces recipe and row origin to synthetic. The broader catalogue contains twelve knowledge profiles, but the other eight crops do not have numerical recipes. This is a useful safety property: a literature record cannot silently become a production coefficient.

Private synthetic data lineage

The reference generator is `synthetic-farm-v1` in `fixtures.py`. It uses a fixed default cutoff (2026-09-08T00:00:00Z) and seed (20260908). It creates four demo recipes, sixteen 20 m² sheltered-hydroponic beds, eight already-executed batches, eight future orders per crop, twelve historical weeks per crop, one opening inventory lot, and declared farm resources. Every record that can carry an origin is labelled synthetic.

The Data Explorer may vary five bounded generator settings: historical multiplier, historical trend, repeating-pattern amplitude, order multiplier, and price multiplier. Its demand baseline uses `recipe-ewma-v1`; alpha is bounded from 0.05 to 0.95 and defaults to 0.35. For each crop, the forecast starts from the first point-in-time-eligible historical value, folds later values through an exponentially weighted moving average, retains booked orders by due date, and adds only the positive residual on the week end. This is an engineering baseline, not a learned production model.



Availability rules are applied before demand calculation. Historical values require `available_at <= cutoff` and a week before the planning date. Orders require `booked_at <= cutoff`; the private contract rejects an order booked after the farm cutoff. Dates are converted to Singapore civil dates for the planning horizon. The feature builder records dependency IDs for each forecast row and produces a Parquet artifact plus a manifest. The committed feature manifest records 32 synthetic features and the exact input, artifact, and private-snapshot hashes.

The planner does not mutate executed batches. It creates whole-bed candidates by recipe and date, reserves nursery sites from sowing to transplant, reserves a bed from transplant through harvest and sanitation, and budgets weekly labour and cash. A daily FIFO lot ledger keeps opening stock, harvest, delivery, disposal, and closing stock in mass balance; later harvest cannot fill an earlier delivery. The implementation and its exact policy weights are in `packages/planner/engine.py`. A companion numerical chapter should be used for formula-level interpretation.

Snapshot and scenario lifecycle

1. `/api/v1/bootstrap` creates an anonymous tenant and its first synthetic farm version when no valid session exists.
2. `/api/v1/imports` validates either the named reference fixture or a strict Farm payload, then appends a farm version. It never overwrites the prior version.
3. `/api/v1/planning-runs` freezes the latest farm payload, input hash, farm version, public source-card snapshot, evidence-register hash, market scaffold, and News context. An idempotency key prevents changed inputs from reusing the request identity.
4. The worker calculates all three policies against one forecast and scenario set. Persisted, sequenced events support server-sent-event reconnect.
5. A versioned backend policy can mark a feasible result `ACCEPTED_FOR_SIMULATION`. The result includes the synthetic ledger and simulated work; it does not authorize a farm operation.
6. Replay returns persisted results and events with `execution_mode=replay`. It does not calculate again or create a provider request.
7. A replan appends a new synthetic farm version and child run. If the farm changes during a run, acceptance becomes stale and a numerical-only refresh starts against the current version.
8. Scenario branches copy a farm or saved-playground root, apply bounded controls, and persist baseline and result. Comparisons reject branches with different roots, baseline hashes, forecast settings, or calculation versions.

Explorer snapshots add another integrity layer. A saved playground freezes the actual generated dataset, forecast, settings, generator/forecast versions, cutoff, reference snapshot, and hashes. Reload validates stored JSON against the recorded hashes. Unsupported old numerical versions fail explicitly rather than being recomputed under current code. Relevant implementation is in `data_explorer.py`; verification is recorded in `reports/data_explorer.md`.

Real public-data ingestion


The source registry identifies 23 candidate datasets, but only seven connectors produced normalized observations in the recorded dataset build. “Listed” and “integrated” therefore mean different things.

Source	Implemented record type	Recorded release snapshot	Point-in-time/planning status
D01 NEA rainfall	Station rainfall, provider interval retained	88 rows	Observation available_at is uncertain; context only

Source	Implemented record type	Recorded release snapshot	Point-in-time/planning status
D02 NEA air temperature	Station air temperature	18 rows	Context only; not root-zone or on-farm temperature
D03 NEA relative humidity	Station relative humidity	18 rows	Context only; station microclimate differs from farm conditions
D04 NEA 24-hour forecast	Issued and valid forecast fields	20 rows	Eligible metadata exists, but no planner feature consumes it
D05 NEA four-day outlook	Four daily forecast groups	20 rows	Valid horizon is preserved and cannot be stretched over a crop cycle
D06 SingStat T010002	Monthly HS8 merchandise trade volume	15 rows	Exact release time unknown; broad codes are not crop demand; export blocked while reuse terms remain unverified
D11 NASA POWER	Seven days of gridded historical temperature, humidity, and precipitation	21 rows	Release timing unknown; historical grid context, not a farm sensor or forecast

The counts above come from the retained dataset manifest: 145 weather observations, 40 weather-forecast fields, 15 trade observations, 200 normalized rows total, and eight raw snapshots because D06 required an initial period probe and a larger response. The manifest was built on 8 September 2026 and is historical evidence, not a claim that the same observations are current now.

The implementation stores raw responses by SHA-256 under a source-specific directory with mode 0600, then records request URL, sorted query, retrieval time, byte count, media type, safe response headers, licence state, and relative path. It normalizes rows with stable IDs and raw locators. Quality validation rejects duplicate IDs, wrong units, non-finite values, naive timestamps, uncertain availability marked as point-in-time eligible, untraceable rows, and orphaned snapshot references.



Source-specific code is split among NEA normalization, SingStat normalization, NASA POWER normalization, snapshot storage, context orchestration, and quality validation. The machine-readable dependency graph is retained in data/manifests/lineage.json.

The public explorer only exposes curated fields. Callers cannot provide file paths. Lists, pages, queries, and exports are bounded; CSV cells beginning with spreadsheet formula characters are escaped. A source-level reuse policy enables export only when licence state is verified. Weather rows from D01–D05 are exportable under the recorded Singapore Open Data licence status;

D06 and D11 rows are view-only in the current configuration. Registry-only candidates appear as metadata without invented observations.

Rebuild semantics and an important reproducibility limit

`python scripts/build_dataset.py --data-dir data` performs a current live refresh. `--offline` rebuilds from raw snapshot files named in the manifest. Raw, normalized, and runtime data directories are intentionally excluded by `.gitignore`, while manifests and quality reports are committed. As a result:

- a live rebuild is reproducible at the code and transformation level, but it will fetch the provider's current responses and need not reproduce the 8 September row values or hashes;
- an offline byte-for-byte rebuild requires an authorized copy of the original raw snapshot directory, which a clean Git checkout does not contain;
- the committed manifest permits audit of hashes, row counts, source queries, transforms, and provenance, but does not independently permit recomputing every retained public value;
- public data licences and terms still govern redistribution even when a checksum is public.

This distinction should accompany any “reproducible dataset” claim.

News and current-information lifecycle

News is a deterministic supporting collector, not an eighth council agent and not a numerical feature pipeline. `packages/news.py` contains seven declared sources. Four are allowlisted RSS metadata collectors: SFA Newsroom, SFA Circulars, AFSIS agricultural-disaster reports, and ASEAN Agri-food. Reddit, SAFE, and the Central Singapore CDC calendar remain visibly unavailable, link-only, or blocked. The v5 release evidence records 346 valid metadata records (154 SFA Newsroom, 123 SFA Circulars, 64 AFSIS, and 5 ASEAN Agri-food); see the News implementation report.

The collector runs separately from the application and provider inference. It follows no redirects, refuses compressed responses, limits each response to 1 MiB and 200 items, validates a fixed HTTPS host, accepts XML media types, strips HTML to bounded text, and never fetches linked articles, images, or social profiles. It retains old records on a refresh failure and exposes the source failure state. A six-hour minimum attempt interval prevents repeated collection.

Each record preserves the exact raw publication string, parsed publication time, retrieval time, canonical URL, publisher geography, keyword-derived crop/topic tags, optional explicit structured event dates, reuse mode, and content hash. A date mentioned in a title is never inferred as an event date. Headlines are treated as untrusted context.

For a new decision, the API chooses a cutoff and copies a bounded News pack into the run or branch payload:

- historical replay uses the farm cutoff;
- a current synthetic decision uses its creation time while retaining the separate farm cutoff;

- both publication and retrieval must be at or before the chosen cutoff;
- continuation and child branches inherit the parent's frozen News context;
- replay reads the copied context;
- the pack declares `numerical_effect: none`.

There is no code path from a News record, headline keyword, social observation, public weather row, or trade row into EWMA demand, recipe yield, order quantities, or CP-SAT constraints. An advisor can cite a recognized frozen key, but any proposed numerical change must become an explicit bounded scenario and pass local validation.

Crop knowledge and scientific evidence

The version-controlled crop catalogue contains twelve unranked Singapore-relevant leafy-herb concepts. The evidence register contains 24 publication records (P01–P24). The catalogue-to-publication relation is many-to-many and preserves taxonomy status, aliases, harvested part, system relevance, evidence IDs, scope boundaries, and parameter gaps.

Crop knowledge profile	Linked evidence	Numerical demo recipe?	Major unresolved boundary
Caixin	P01, P02	Yes, synthetic	Local commercial mature marketable yield and cultivar cycle
Pak choi / xiao bai cai	P03–P05	Yes, synthetic	Species/cultivar-to-SKU mapping and commercial packout
Kailan	P05–P07	Yes, synthetic	Authority-resolved name/cultivar and commercial yield
Bayam	P10	No	Portfolio species, full-size yield, and multi-cut response
Kangkong	P12	No	Singapore commercial yield, cultivar, and regrowth
Lettuce	P08, P09, P16, P17, P19	Yes, synthetic	Cultivar/SKU, Singapore cycle yield, energy, and shelf life
Kale	P15	No	Taxon and cultivar remain blocked pending methods confirmation
Mustard greens	P11	No	Evidence taxon does not resolve all portfolio SKUs

Crop knowledge profile	Linked evidence	Numerical demo recipe?	Major unresolved boundary
Malabar spinach	P14	No	Singapore evidence, trellis capacity, and cut-specific yield
Sweet-potato leaves	P13, P18	No	Fresh-market leaf yield and local regrowth schedule
Garlic chives	P21, P22	No	Local accession, commercial yield, and cut-specific regrowth
Sawtooth coriander	P23, P24	No	Local seed source, flowering/regrowth, and buyer quality

The two v6 additions and their primary-study interpretation are documented in docs/research/v6_crop_evidence.md. The catalogue also links public government identifiers and names through source_registry.json, while crop_hs_mappings.json distinguishes exact product references from unresolved broad customs codes.

The evidence register is a targeted research set, not a systematic review or a training corpus. It records access-review status, study setting, finding, design implication, transfer limit, licence state, and a parameter_use decision. Most entries are restricted to context, candidate evaluation, or citation discovery. None of the 24 papers supplies the four numerical demo recipes. In particular:

- a non-Singapore controlled experiment does not validate Singapore commercial performance;
- biomass, dry weight, physiology, metabolomics, and nutritional quality are not interchangeable with fresh marketable kilograms;
- study measurements available only at harvest would leak future information into an earlier yield prediction;
- one review and its cited trial are not independent replications;
- a title, abstract, or inaccessible full text cannot authorize a numeric coefficient;
- broad HS trade categories cannot become named-crop customer demand;
- an illustration is representative artwork and does not resolve a cultivar, SKU, disease, or growth measurement.

The Crop Atlas returns the exact linked evidence records for a selected profile and displays the record's source URL and transfer limit. It also shows a planning recipe only for the four simulated crops and labels it demo_only. This makes the scientific layer inspectable without presenting it as numerical authority.

Persistence, tenancy, and lifecycle controls

The application uses one SQLAlchemy metadata collection and creates additive tables at service startup. The important persisted grains are:

Tables	Grain and lifecycle role
tenants, farm_versions	Anonymous session identity and append-only tenant farm versions
planning_runs, run_events	One idempotent planning request and its ordered event stream
scenario_branches, quest_progress	Isolated numerical experiment and per-tenant game progress
explorer_snapshots	Integrity-checked synthetic playground snapshot
conversations, conversation_messages, conversation_requests, conversation_events	Frozen advisor context, messages, idempotent provider request, and replayable events
council_research_sessions, council_research_jobs, council_research_actions	V7 isolated research state, one result per input version, and idempotent actions
inference_budget, rate_counters, security_settings	Durable global provider budget and abuse-control state

All user-state queries include `tenant_id`. Planning events and conversation children additionally use composite foreign keys that bind child identity to the same tenant as the parent. Unknown or unowned IDs return not found. Idempotency keys are scoped to the tenant or parent resource and reject reuse with changed request hashes. Tenant-row locking serializes farm-version allocation and competing creates.

Worker recovery reflects side-effect risk. Numerical scenarios return from `RUNNING` to `QUEUED` because recalculation is local and deterministic. An abandoned inference-bearing planning or conversation request is marked failed or interrupted so a restart does not silently repeat a paid call. Reconnect and replay consume stored events.

The v7 research route is a separately persisted study. Free text is parsed only into a small allowlisted set of proposals. An explicit apply action creates a new research input version, and calculation uses a deep copy of the synthetic farm. The main farm and its missions are not modified. The optional actual-advisor action is separately labelled and persists evidence-validation status. Implementation: `council_research.py` and `planner/research.py`.

API-to-interface traceability

The API is the only bridge from persisted or curated data into the React client. The public OpenAPI/Swagger endpoints are disabled, so source contracts and generated TypeScript types are the reviewable interface definition.


User surface	API source	What the UI can truthfully establish
Farm board and planning room	/bootstrap, /planning-runs, /events, /replay, /worklist.csv	Synthetic current farm, frozen inputs, locally calculated alternatives, event status, and simulation-only work
Crop Atlas	/crops, /crops/{id}/evidence	Catalogue identity, unresolved boundaries, demo recipe status, and linked evidence records
Data Explorer	/data-explorer/*	Selected tenant snapshot, records, charts, EWMA preview, exact hashes, same-root comparisons, and bounded exports
Public Explorer	/data-explorer/public*, /sources	Cached public rows, source time, retrieval time, freshness, licence state, and metadata-only candidates
News panel	/news	Frozen or current cache metadata with explicit time/relevance limits and no numerical effect
Scenario Lab	/scenarios, /scenarios/compare, /quests	Explicit synthetic control, immutable baseline, local result, and affected-resource pointers
Advisor conversations	/conversations/*	Frozen snapshot/evidence context, provider-request status, validated references, and stored replay
V7 council research	/council-research/*	Versioned scripted interaction, explicit proposal/apply lifecycle, local plans, challenges, and optional provider interpretation

The UI renders provider prose as text, opens external evidence links with safe browser attributes, and exposes structured provenance rather than provider internals. Static routes are fixed build directories; there is no route that accepts a filesystem path. Request bodies use strict Pydantic models with unknown fields rejected, and writes are same-origin checked and capped at 1 MiB.

Edition and deployment isolation

The current release registry marks v7 as latest and records a full source commit and immutable OCI image digest for every v1-v7 edition. Since v6, one Singapore Fly Machine and one encrypted volume host the gateway and all edition containers. Sharing a host is not sharing application state: the host adapter binds a fixed edition subtree to /data, removes access to the common parent, and each container maintains its own PostgreSQL cluster, cache, session state,

and game records. The gateway owns shared abuse counters and the 48-call daily inference ceiling because those are operational safeguards rather than game progress.



Cross-edition cookies are rejected and published source/image identities are immutable. A new edition begins with fresh reference state. The single Machine and volume remain a shared availability and capacity boundary: a host restart can interrupt every edition, and the v7 verification is a bounded two-job smoke test rather than a saturation or long-duration study. Details and recovery procedure are in docs/deployment/editions.md; v7 preservation and capacity evidence are in reports/v7/implementation.md.

Evidence strength and claim limits

Claim	Best supporting artifact	What the artifact does not establish
The public pipeline produced 200 traceable rows from seven source IDs	dataset manifest, lineage graph, quality report	Current freshness, farm relevance, commercial prediction accuracy, or redistribution permission for every source
The feature builder excludes public context from the demand baseline	feature manifest, build_features.py	That public signals would be useless after future validation; only that they are unused now
Twelve crop profiles link to 24 publications	crop catalogue, evidence register	Exhaustive literature coverage, independent replication, local transferability, or approved production coefficients
The News collector retained 346 valid metadata records for v5	News report	Full-article verification, sentiment, causal impact, representative public opinion, or current availability
Explorer snapshots and scenarios persist and replay	Data Explorer report, v7 persistence evidence	Durability under arbitrary infrastructure failure or migration from unsupported future model versions

Claim	Best supporting artifact	What the artifact does not establish
Prior edition records survived v7 publication	v7 release report	Multi-region availability or unlimited shared-host capacity
An actual v7 advisor request completed	v7 advisor probe	Correct farming advice; the recorded response had no verified evidence references and was marked unsupported

Automated checks provide strong evidence for schema enforcement, deterministic calculations, idempotency, tenant isolation, browser flows, hashes, and bounded release preservation. They are weaker evidence for scientific validity, human usability, and real operations. Agent walkthroughs are not human studies. Headless mobile geometry is not physical touch or software-keyboard testing. Digital audio checks are not proof of audibility on a farmer's device. Synthetic held-out weeks generated by the same fixture process do not measure out-of-sample commercial accuracy.

No real farm outcome dataset is present. There is no trained supervised demand model, harvest/yield model, disease classifier, satellite predictor, calibrated probabilistic risk model, production retraining pipeline, drift monitor, causal estimator, or measured SME benefit. EWMA is a statistical baseline, CP-SAT is an optimizer, and scenario weights are declared assumptions. "AI" in this release refers to optional qualitative advisor calls; it does not replace data collection, model validation, or numerical computation.

Reproduction and audit route

From a dependency-complete checkout, the main read-only or local-data reproduction commands are:

```

# Validate the committed synthetic fixture and rebuild its point-in-time features.
.venv/bin/python scripts/build_features.py

# Fetch current bounded public observations and write a new manifest.
.venv/bin/python scripts/build_dataset.py --data-dir data --with-power

# Rebuild from an authorized local copy of the raw snapshots named in the manifest.
.venv/bin/python scripts/build_dataset.py --data-dir data --offline

# Refresh the bounded RSS metadata cache; no inference is involved.
.venv/bin/python scripts/refresh_news.py

# Run focused data, private-contract, explorer, and isolation verification.
.venv/bin/python -m pytest -q \
  tests/contracts/test_private_data.py \
  tests/contracts/test_features.py \
  tests/data/test_ingestion.py \
  tests/data/test_explorer_public.py \
  tests/explorer/test_data_explorer.py \
  tests/review/test_postgres_concurrency.py

# Verify the generated browser contracts and production client build.
.venv/bin/python scripts/generate_web_contracts.py --check
npm run build --prefix apps/web

```

The public-data commands access changing upstream services, so record retrieval time, response status, row counts, failures, and new hashes rather than expecting the historical manifest's values. PostgreSQL concurrency tests require a dedicated test database and must never target a live application database. Running the commands above does not require or authorize a provider call.

Material gaps found in this audit

The following discrepancies were found against the audit baseline. This documentation update corrects stale prose and specifies code fixes; it does not claim the listed code gaps are repaired:

1. `packages/ingestion/validation.py` still hard-codes `catalogue_profiles: 10` and `scientific_evidence_documents: 20`. The current registries contain 12 profiles and 24 documents. Consequently, the committed data quality report is historically accurate for the 8 September build but stale as a current coverage statement. Coverage should be derived from the registries or clearly versioned as build-time metadata.
2. `FarmTact_Build_Specification.md` labels its v6 section “Current implementation,” says the shared host runs v1–v6, and says the evidence register has 20 publications. V7 is the current published edition and the register now has 24 publications. Historical descriptions should be labelled historical and the current amendment consolidated at the top.
3. `docs/execution-plan.md` correctly names v7 as current, but its unqualified “Acceptance scope” still says ten crop profiles and its early build summary says ten profiles/twenty

records/six roles. Those statements describe the original milestone and need an explicit “historical v1 scope” qualifier where they appear.

4. docs/deployment/fly.md describes v6 and v1–v6 as the current topology. It should defer to the v7 editions runbook or identify itself as a v6 historical snapshot.
5. docs/v5-engagement-plan.md says v6 is current. Its historical-plan label is useful, but the “current” sentence should point readers to v7.
6. docs/research/v5_news_sources.md refers to coverage broader than “the ten crops.” The source was a v5 research brief, but the current shared crop vocabulary has twelve identifiers; either freeze the statement as v5-era scope or update it.
7. The build specification's recommended architecture mentions Next.js, DuckDB, an object-store abstraction, and broader operational tables. The implementation uses Vite React, FastAPI, PostgreSQL JSON payloads, local content-addressed files, and a much narrower synthetic Farm contract. This is a roadmap-versus-implementation difference and should be stated where the architecture recommendation appears.
8. The committed public manifest points to ignored raw and normalized files. The documented offline rebuild is therefore not self-contained in a clean checkout. The docs should say that an authorized raw-snapshot bundle is required for exact historical reproduction.

These gaps do not invalidate the implemented synthetic demonstration. They narrow what can be claimed from its current artifacts and make clear which records are historical evidence, which are active contracts, and which remain future design requirements.

The specification/documentation wording findings above refer to the audit baseline. The integrated documentation update corrects current pointers and labels historical requirements; numerical, data-generation and runtime gaps remain future work.

Implementation gaps and next-iteration acceptance criteria

Audit: 11 September 2026, source baseline a96025e; deployed v7 remains frozen. This register preserves the 48 findings observed at that audit baseline and adds an evidence-backed disposition for the current V8 local tree. The “Observed gap” columns below remain historical; they have not been rewritten to make the original audit look better. Fixed means the specified technical mechanism is implemented and locally verified, while Partial means an empirical, operational or product acceptance condition remains open. P0, P1 and P2 retain their original audit meaning. Actual farm operations remain disabled.

Source detail: Council/provider audit, numerical audit, game backend audit, and data architecture audit.

1. Council, evidence and provider behavior

ID / priority	Observed gap and source	Required next behavior / acceptance evidence
AI-01 / P0	Both v7 actual adviser probes remain unsupported; transport/replay success is not answer-quality success. Local, public.	Build frozen, expert-reviewable cases with expected references and permissible conclusions. Evaluate role relevance, unsupported quantities, contradictions, abstention and useful answers separately from HTTP/schema success. Preserve all failed attempts and per-case usage. Do not hide unsupported replies or use them to accept a plan.
AI-02 / P0	Reference membership and numeric-value matching do not prove that cited evidence entails a sentence. Claims, conversation validation.	Represent quantitative assertions as typed reference/value/unit/entity/period relations and render those values from code. Test wrong crop/date/unit/sign, a correct number attached to the wrong meaning, fabricated causal claims and citation laundering. Qualitative statements retain an explicit unverified status unless independently supported.

ID / priority	Observed gap and source	Required next behavior / acceptance evidence
AI-03 / P1	Three mechanisms share “council” language: precomputed planning findings, bounded conversation turns and scripted research dialogue.	Expose workflow type, inference origin, model/call status and whether a new calculation occurred. Tests must distinguish a script, fixture, actual completion, stored actual replay, failed provider call and numerical-only result. No invented tool or agent activity.
AI-04 / P0	Capability verified is derived from archived trial JSON plus key presence for text; vision uses archived success. views.capabilities.	Return historical probe date/model/source, configured/not-configured state and last observed execution separately. Missing/aged/account-mismatched evidence must not claim current verified access. Test old reports, no key, missing model and failed last execution without triggering paid background probes.
AI-05 / P1	Registered helper, evaluator and vision role names exceed actual integrated call sites. Runtime manifest.	Maintain a machine-readable route-to-caller/feature inventory. Distinguish available server capability from active product integration. A new helper/evaluator requires caller, prompt/schema version, finite budget, test, egress review and visible failure semantics.
AI-06 / P1	Per-run request and reserved-output limits plus shared call counts do not establish a hard USD spending cap or strict end-to-end cancellation. Gateway.	Add conservative total-cost reservation with dated price metadata if a monetary guarantee is required; report input/output/cache/reasoning accounting. Test timeout/cancellation during transport and between turns, midnight reconciliation and multi-edition concurrency. Do not automatically retry potentially billed interrupted generations.
AI-07 / P1	Crop-harvest forecast emits marketable_kg; the conversation reference builder expects expected_kg, omitting that forecast mass reference. Forecast, context builder.	Adopt one shared typed harvest record and reference generator. Assert that every forecast batch/date/marketable-mass field is available with the same value, unit and frozen snapshot hash in adviser context. Preserve existing batch-level references during migration.

ID / priority	Observed gap and source	Required next behavior / acceptance evidence
AI-08 / P1	Mission specialists see the same frozen tools but not one another's findings; only the chair sees previous validated/rejected findings. No two-round challenge protocol exists.	Preserve this truthful sequential contract or implement explicit challenge targets, turn/round budgets, stop states and evidence updates. Test which prior findings and rejection reasons each role receives; measure usefulness against the simpler existing chain.
AI-09 / P1	Conversation prompts request tighter response limits than AdvisorReply enforces (400 versus 900 characters and lower reference/action counts).	Put one bounded response contract in the schema and generate prompt instructions from it. Test limits at and just beyond every boundary; distinguish structural repair from unsupported-claim handling.
AI-10 / P1	Conversational chair sees prior validation_status but its context projection omits detailed validation_errors; mission chair receives rejection reasons.	Define a bounded shared validated-message projection with explicit error categories and test that later roles cannot treat rejected/unsupported statements as evidence.
AI-11 / P1	Per-call records lack a complete explicit prompt/schema/validator version set; roster version alone cannot identify prompt behavior.	Persist prompt-template, output-schema, validator, numerical/context and source versions with every completion and replay export. Demonstrate reconstruction of the sent public context without persisting private reasoning.
AI-12 / P1	A persistent conversation request may be COMPLETED while every reply is unsupported; a valid mission claim can withhold selection through its recommendation field.	Expose execution status, evidence status and decision influence separately. Test all-successful transport with unsupported answers, partial failures and a single veto; compare with a numerical-only run on identical inputs.

Council-specific chain, validation and budget gaps are analyzed in the AI chapter. The game's challenge/proposal/acceptance state gaps are separately identified as GAME- * in the backend chapter; those are part of this register's next-iteration scope and must not be lost in a provider-only assessment.

2. Crop development, optimization and simulation

ID / priority	Observed gap and source	Required next behavior / acceptance evidence
NUM-01 / P0	Backend marks batches ready at/after harvest; standalone frontend preview makes beds empty at/after harvest. Neither exposes sanitation occupancy. views, World.	Define one versioned crop-cycle state contract for nursery, growing, ready, harvested, sanitation and available, distinguishing planned versus recorded events. Test every day around transplant, harvest and sanitation end, with/without accepted overlays and delayed batches. UI availability must agree with planner occupancy.
NUM-02 / P1	Preview slider includes <code>planning_date + horizon_days</code> , one day beyond the final planner index.	Use the same inclusive horizon contract everywhere. Test a 7-, 56- and 84-day horizon, date labels, exports and final-day inventory.
NUM-03 / P1	Three scenario weights are stored but not consumed by the solver; seed produces no random draws. models, engine.	Either enforce equal-weight-only configuration explicitly or consume validated normalized weights. Demonstrate changed weights affect a constructed optimum; report scenario provenance. Rename/remove unused random-seed semantics until stochastic sampling exists.
NUM-04 / P1	“Resilient” increases shortage penalties; it is not maximin, CVaR or a low-yield service guarantee. The objective is policy-weighted utility, not literal expected net profit.	Publish objective terms and units. If downside protection is the desired product contract, implement an explicit risk/service constraint with infeasibility reporting. Test trade-offs on small enumerated cases and compare all policies on the same baseline.
NUM-05 / P0	CP-SAT inventory age buckets do not explicitly require FIFO; the reporting simulator uses FIFO. This is an unproven equivalence, not a demonstrated failing numerical case.	Prove equivalence for the supported domain or align both allocation rules. Enumerate small perishability cases and reconcile delivered mass, expiry, revenue and selected objective against replay. Keep this classified as a validation gap until a counterexample or proof exists.
NUM-06 / P1	General terminal stock value/penalty is absent; horizon truncation can affect the preference for unconsumed stock.	Declare terminal treatment and test horizon extension, late harvest, salvage assumptions and expiry beyond the horizon. Never classify all closing stock as physical waste.

ID / priority	Observed gap and source	Required next behavior / acceptance evidence
NUM-07 / P1	Residual demand with no booked line in its week is assigned zero price. Fulfilment is by crop/date, not individual buyer/order/grade.	Make missing prices explicit and distinguish requested, booked and delivered quantities. Add buyer-level allocation only with corresponding schema/priority rules; test scarce stock with multiple prices and cancellation states.
NUM-08 / P2	No fitted maturity/yield/growth/quality/weather-response model; delay/yield shocks are direct inputs without biological coupling.	Collect independent crop-cycle outcomes and availability timestamps. Compare recipe baselines with candidate models on whole-batch temporal holdouts; report timing/yield errors and interval coverage. No model promotion from synthetic accuracy alone.
NUM-09 / P2	EWMA is constant-level and tested only on patterned fictional history; no calibrated demand distribution.	Rolling-origin multi-horizon real-order evaluation with last-week and seasonal baselines, leakage checks and stockout/cancellation handling. Separate farm/buyer cohorts and report uncertainty in validation estimates.
NUM-10 / P2	Whole-bed, single-harvest assumptions omit water, energy, HVAC, seed, packaging, grade, rotation, multi-cut, partner supply and buyer-level contracts.	Maintain a supported-constraint registry. Extend schema, validator, solver, replay and UI together for each capability. In particular, multi-cut crops require persistent occupied capacity and regrowth states, not another single-cut recipe.

3. Evidence, provenance and engineering reproduction

ID / priority	Observed gap and source	Required next behavior / acceptance evidence
DATA-01 / P0	Quality builder hard-codes 10 profiles/20 publications; current registries contain 12/24. validation.py.	Derive coverage counts from validated input registries, version manifests and fail on contradictory totals. Preserve old release reports as dated artifacts rather than rewriting them as current evidence.

ID / priority	Observed gap and source	Required next behavior / acceptance evidence
DATA-02 / P1	integrated_demo.py still expects 10 crops, while the catalogue contains 12.	Derive expected coverage from a versioned manifest and distinguish recipe count from knowledge count. An integrated run must pass without weakening tests to accept arbitrary missing crops.
DATA-03 / P1	The checked-in numerical report has evaluation_scope metadata that its current generator does not emit.	Add a versioned report schema and generator check. Regeneration must preserve cohort, settings, fixture hash, cutoff, model versions and explicit synthetic-only conclusions. Archive before replacing evidence; report solver nondeterminism where applicable.
DATA-04 / P1	Ignored raw/public cache files prevent an exact historical rebuild from Git alone.	Produce a licensed, integrity-checked snapshot bundle or clearly offer transform-only reproduction. Compare source byte hashes and availability cutoffs; a new live fetch is a new dataset version.
DATA-06 / P1	Fresh offline game checks fail because test_seven_current_advisors_and_frozen_create_idempotency assumes an ignored NEA D04 cache is present (source:D04.summary).	Give offline tests explicit bounded source fixtures or assert the documented missing-source state; keep real-source integration checks separate. A clean checkout must pass its declared offline group without fetching current data to satisfy a fixture assertion.
DATA-05 / P1	Citation/source presence does not validate numerical crop coefficients; public sources are contextual only.	Preserve public_features_used=[] until an explicit feature path and point-in-time validation exist. Any new weather/News/market numerical effect needs units, availability, exposure mapping, ablation and evidence of predictive benefit.
OPS-01 / P1	Python egress policy is not an OS/network enforcement boundary; a single host shares availability across editions.	Before operational promotion, validate infrastructure egress controls and recovery on an isolated environment. Exercise provider denial and database restore, not only static string scans or volume existence.

ID / priority	Observed gap and source	Required next behavior / acceptance evidence
UX-01 / P1	AI/headless walkthroughs do not establish human usability or agronomic decision benefit.	Pre-register a small task-based human study with novice and relevant practitioner groups, clear synthetic framing, task success, time, error and comprehension outcomes. Do not infer benefits from animated dialogue or agent agreement.

4. V8 remediation disposition in the current local tree

The machine-readable source for this table is gap-status.json. The complete PostgreSQL-backed regression reports 589 passed, 1 skipped and 0 failed in full-regression.xml. The newest focused root checks report 23 passed and 0 failed in final-root-focused.xml. These local results do not establish current AI answer quality or a published V8 release. Those two gates remain PENDING_ROOT_AI_QUALITY_FINAL and PENDING_ROOT_V8_DEPLOYMENT.

ID	V8 status	Evidence-backed disposition	Remaining limitation
AI-01	Partial	A frozen zero-inference quality scorer and bounded actual-output harness now assess role relevance, expected references, unsupported quantities, forbidden conclusions, abstention and usefulness. Evidence: council.md, ai_quality.py, deepseek_quality_e2e.py.	Final stored direct/invite/Council/research quality run is pending; human expert agronomic review is absent.
AI-02	Fixed	Server-owned typed facts bind reference, value, unit, entity, period and snapshot hash; unknown references and model-authored quantities fail local validation. Qualitative entailment is not established. Evidence: council.md, ai_contracts.py, test_ai_injection_boundaries.py.	Qualitative entailment remains explicitly unverified rather than promoted.

ID	V8 status	Evidence-backed disposition	Remaining limitation
AI-03	Fixed	Workflow type, inference origin, model-call state and calculation state distinguish numerical, scripted, actual, failed and replayed paths. Evidence: council.md, conversations.py.	Older records use conservative compatibility values.
AI-04	Partial	Configured state, dated archived evidence and tenant-scoped last observed execution are separate and reads trigger no probe. Evidence: council.md, test_v8_integration_policy.py, live-call-ledger.json.	Current canonical provider quality/availability remains pending explicit root-owned evidence.
AI-05	Fixed	The fail-closed runtime manifest now inventories route callers, integrations, prompt/schema versions and bounded roles. Evidence: council.md, deepseek_runtime.json, test_provider_isolation.py.	It is reviewed configuration rather than automatic call-graph discovery.
AI-06	Partial	Exact request ceilings, token accounting, cooperative cancellation and no automatic billed retry are implemented and tested. Evidence: council.md, deepseek_gateway.py, game-reliability.md.	No dated-price USD cap, strict mid-transport abort, or completed deployment-wide midnight accounting proof exists.
AI-07	Fixed	A typed harvest forecast and shared reference builder expose marketable mass with batch/date/endpoint/origin /value status and units. Evidence: council.md, test_synthetic_evaluation.py.	Frozen older records remain unchanged and are interpreted conservatively.

ID	V8 status	Evidence-backed disposition	Remaining limitation
AI-08	Partial	The product now states the actual sequence: six independent specialists followed by a chair that receives bounded validated and rejected findings. The malformed blank-table break in the audit register is removed. Evidence: council.md, test_council.py.	No two-round challenge protocol exists and no evidence shows it would improve usefulness.
AI-09	Fixed	One RESPONSE_LIMITS contract drives schema and prompt bounds, with exact-boundary and one-over tests. Evidence: council.md, conversations.py.	Sentence count remains nonbinding guidance.
AI-10	Fixed	Later-role projections include bounded validation issue codes/messages, evidence state and eligibility; rejected turns cannot become validated evidence. Evidence: council.md, conversations.py.	Provider compliance with instructions is not assumed; the local gate remains authoritative.
AI-11	Fixed	New completions persist prompt, schema, validator, numerical/context and source versions plus a public-context hash. Evidence: council.md, deepseek_gateway.py.	Private provider reasoning is intentionally neither persisted nor reconstructed.
AI-12	Fixed	Execution, evidence and decision-influence states are separate for missions and conversations, including all-unsupported successful transport. Evidence: council.md, conversations.py.	The legacy validation_status field remains only as a compatibility alias.
DATA-01	Fixed	Registry coverage is derived from validated inputs and contradictory manifest totals fail; archived reports remain dated artifacts. Evidence: data-ml.md, validation.py.	This validates registry accounting, not crop coefficients.

ID	V8 status	Evidence-backed disposition	Remaining limitation
DATA-02	Fixed	Integrated coverage checks derive expected catalogue size and distinguish knowledge profiles from planner recipes. Evidence: data-ml.md, integrated_demo.py.	Only four crops currently have planner recipes.
DATA-03	Fixed	Versioned numerical-report generation preserves fixture, cohort, settings, cutoff, model and synthetic-only metadata and checks regeneration. Evidence: data-ml.md, numerical_evaluation.json.	Time-limited CP-SAT objective/bounds can vary and are disclosed.
DATA-04	Partial	A licensed project-authored offline provider-contract fixture is byte-hashed and reproduces the current offline pipeline. Evidence: data-ml.md, data_pipeline_verification.json.	The ignored historical public response bodies remain an archive and cannot be reconstructed exactly from Git.
DATA-05	Fixed	Public context is excluded from numerical features and every crop/model artifact remains synthetic/demo-only with production promotion blocked. Evidence: data-ml.md, synthetic_model_evaluation.json.	Real coefficient validity remains unevaluated by design.
DATA-06	Fixed	Clean-checkout tests use bounded explicit source fixtures or typed missing-source state and do not require a network fetch. Evidence: data-ml.md, test_ingestion.py.	Live-source checks remain separate and date-specific.

ID	V8 status	Evidence-backed disposition	Remaining limitation
GAME-01	Fixed	Versioned tenant-owned synthetic worlds now advance a civil-date clock, record tasks/order service/inventory/cash, replan future work and replay idempotently. Evidence: execution-validation.md, test_simulation_execution.py, test_v8_simulation_postgres.py.	Execution remains synthetic and cannot operate a farm.
GAME-02	Fixed	Provider and numerical work use independent worker lanes; a blocked provider job no longer blocks another tenant local calculation. Evidence: test_v8_integration_policy.py, app.py.	Each lane is still bounded single-process capacity, not a distributed queue.
GAME-03	Fixed	Mission requests and acceptance records persist explicit required/advisory Council policy, with missing/partial Council outcomes handled accordingly. Evidence: council.md, test_mission_gate.py.	No Council outcome itself authorizes real operations.
GAME-04	Partial	The implemented topology is now represented truthfully as sequential specialist reviews plus chair synthesis with explicit visibility. Evidence: council.md, v8-human-study-protocol.md.	Challenge rounds were not implemented and comparative usefulness is unmeasured.
GAME-05	Fixed	Research challenges retain history, revalidate on changed inputs, and action retries return the original revision response. Evidence: test_v8_history.py, council_research.py.	The workflow remains scripted unless an explicit separate advisor request is made.

ID	V8 status	Evidence-backed disposition	Remaining limitation
GAME-06	Fixed	A shared browser mutation transport persists edition/path/body-bound keys before submission; imports and other mutations have backend receipts. Evidence: ui.md, v8_simulation_ui.mjs, test_v8_integration_policy.py.	Session storage is browser-session scoped rather than cross-device synchronization.
GAME-07	Fixed	Conversation transport completion and evidence validity are distinct, and unsupported output remains advisory-only. Evidence: council.md, test_ai_quality.py.	Answer quality still requires the pending AI-01 evaluation.
GAME-08	Fixed	Failed/cancelled immutable scenarios can retry with the original run key, bounded attempt history and atomic no-duplicate accounting. Evidence: game-reliability.md, test_v8_reliability_postgres.py.	Retries remain local numerical recomputation only.
GAME-09	Partial	Published editions retain exact pinned images and isolated state, so their frozen calculation code remains addressable. Evidence: editions.md, operations-review.md.	Development upgrades still lack a general per-object runner registry; V8 deployment verification is pending.
GAME-10	Fixed	Council research retains append-only revision history beyond bounded current message/event projections and returns original action receipts. Evidence: test_v8_history.py, council_research.py.	The current view remains intentionally bounded.

ID	V8 status	Evidence-backed disposition	Remaining limitation
GAME-11	Partial	Scenario and conversation admission/listing now enforce tenant quotas and bounded cursor queries before expensive work. Evidence: game-reliability.md, test_v8_reliability.py.	An explicit retention command now previews or deletes bounded inactive expired tenants with shared budgets protected. Operator scheduling, backup-retention alignment and deletion-request procedures remain unestablished.
GAME-12	Partial	Scenario, conversation and research cancellation endpoints and safe-boundary checks prevent stale acceptance/new provider turns and reconcile unused reservations. Evidence: game-reliability.md, conversations.py, scenarios.py.	An active CP-SAT call and synchronous provider transport are cooperative rather than immediately preemptible.
GAME-13	Fixed	Main disruption replans accept bounded batch ID, delay, yield and input version, reject stale/unknown input and preserve parent provenance. Evidence: app.py, test_v8_integration_policy.py.	The supported disruption remains one synthetic crop-delay/yield action, not a general observation model.
GAME-14	Fixed	Automatic selection persists a deterministic Balanced/service/margin/I D ranking independent of input list order. Evidence: app.py, test_mission_gate.py.	The ranking is a declared development policy, not a learned preference.
GAME-15	Fixed	Planner and scenario impacts now expose exact demand-line and lot-to-order allocations with expiry, delivery and shortfall conservation. Evidence: planner.md, game-reliability.md, test_planner_v8.py.	Buyer, grade and contractual priority remain outside the current schema.

ID	V8 status	Evidence-backed disposition	Remaining limitation
GAME-16	Partial	Mission identity and pending mutation state survive reload in edition-scoped browser storage; current hashes/revisions refresh backend truth. Evidence: ui.md, v8_simulation_ui.mjs.	Derived navigation/persona/sound state is still local and does not synchronize across devices or durably acknowledge sound.
GAME-17	Fixed	Cancelling a terminal mission returns its actual status, cancelled=false and an explicit already_terminal reason without mutation. Evidence: test_v8_integration_policy.py, app.py.	Active cancellation remains cooperative.
GAME-18	Fixed	Offline conversation and data tests use explicit bounded source fixtures or typed unavailable states; clean-checkout behavior no longer depends on ignored D04 cache. Evidence: data-ml.md, data_pipeline_verification.json.	Live public-source freshness is evaluated separately.
NUM-01	Fixed	One schedule-state contract distinguishes nursery, growing, ready, recorded harvest, sanitation and availability across backend, UI and execution. Evidence: ui.md, growth.py, test_simulation_execution.py.	Elapsed schedule progress is not physiological growth.
NUM-02	Fixed	The UI and planner use the same inclusive horizon ending at planning_date + horizon_days - 1 for 7, 56 and 84 days. Evidence: ui.md, v8_simulation_ui.mjs.	No dates beyond the frozen horizon are inferred.

ID	V8 status	Evidence-backed disposition	Remaining limitation
NUM-03	Fixed	Scenario weights are validated, normalized to integer units and consumed by the objective; constructed tests show weight-sensitive choice and provenance. Evidence: planner.md, test_planner_v8.py.	No stochastic draws or calibrated scenario probabilities are claimed.
NUM-04	Fixed	Objective terms/units are explicit and Resilient performs lexicographic maximin fill before weighted utility when the primary solve is proven optimal. Evidence: planner.md, test_planner_v8.py.	Time-limited solver metadata does not claim a global proof.
NUM-05	Fixed	Solver expiry buckets and deterministic FEFO/FIFO replay share earliest-expiry behavior; tiny enumerated perishability regressions reconcile selection and replay. Evidence: planner.md, accounting.py, test_planner_v8.py.	The proof applies only to the registered supported domain.
NUM-06	Fixed	Closing lots remain explicit physical terminal inventory with zero salvage and a separate policy penalty; they are not mislabeled as waste. Evidence: planner.md, test_planner_v8.py.	No empirical terminal price or general salvage model exists.
NUM-07	Fixed	Missing price has an explicit state and creates no revenue; booked/residual/order delivery and lot attribution are separate. Evidence: planner.md, game-reliability.md.	Buyer, grade and contractual-priority allocation requires future schema expansion.

ID	V8 status	Evidence-backed disposition	Remaining limitation
NUM-08	Partial	Independent typed synthetic whole-batch outcomes, temporal/external-style splits, baseline evaluation and hard promotion blocks are implemented; a real-outcome protocol is preregistered. Evidence: data-ml.md, v8-outcome-collection-protocol.md, synthetic_model_evaluation.json.	No authorized observed farm outcomes or fitted production biology model exist.
NUM-09	Partial	Synthetic rolling-origin 1/2/4-week farm/buyer evaluation includes cancellations, availability checks, baselines and bootstrap intervals; real evaluation is preregistered. Evidence: data-ml.md, v8-outcome-collection-protocol.md, test_v8_availability_adversarial.py.	No authorized real-order holdout or calibrated demand distribution exists.
NUM-10	Partial	A machine-readable registry now enumerates implemented constraints, bounds, unsupported capabilities and the contract/validator/solver/replay/UI work needed for each expansion. Evidence: simulation_capabilities.json, operations-review.md.	Water, energy, HVAC, consumables, grade, rotation, multi-cut, partner supply and richer buyer/logistics constraints remain unsupported.
OPS-01	Partial	A local isolated PostgreSQL dump/restore preserved hashes/counts; shared namespace and Python-process egress mechanisms have focused tests. Evidence: operations-review.md, test_shared_transfer_operator.py, test_process_egress.py.	OS/network egress, Fly snapshot restore, multi-host failover, sustained availability and recovery objectives remain unproven.

ID	V8 status	Evidence-backed disposition	Remaining limitation
UX-01	Partial	A novice/practitioner task study is preregistered with fixed tasks, outcomes, thresholds, coding and privacy handling. Evidence: v8-human-study-protocol.md, ui.md.	No participants have been enrolled; human usability and agronomic benefit remain unknown.

The unresolved items cluster around evidence that software tests cannot create: observed farm outcomes (NUM-08), real-order forecasting (NUM-09), participant results (UX-01), and operator recovery/availability controls (OPS-01). Other partial items retain bounded product limitations such as immediate in-flight cancellation, retention, cross-device client state, challenge rounds and unsupported resource/biology constraints. AI quality and immutable V8 deployment remain the explicit pending release gates above.

5. Original bounded implementation sequence (historical)

1. **Restore semantic consistency:** AI-04, AI-07, NUM-01/02, DATA-01/02/03 and the concrete game-state inconsistencies in GAME-*. These improve truthful status, synchronized state and reproducible evidence without pretending to improve biology.
2. **Make the council useful and inspectable:** AI-01/02/03 and the Council backend gaps. Test interpretations against fixed quantities and explicit unresolved assumptions; make proposal → apply → calculation → compare → selection observable.
3. **Strengthen numerical acceptance:** NUM-03/04/05/06/07 with enumerated small examples, explicit policy semantics and objective/replay reconciliation.
4. **Evaluate people and data:** UX-01 and NUM-08/09/10 only with appropriate observed data and research design. Keep operational integrations disabled.

This sequence records the order proposed at the audit baseline. The current status and remaining acceptance work are recorded in Section 4. Any published V8 iteration must still use the next unused immutable edition and retain previous editions' source, image and state.

Reproducing the implementation report

Historical baseline: This chapter records the pre-remediation audit of source a96025e and published v7. Its implementation findings and measured results describe that baseline. Read the V8 remediation report for the current simulation engine, synthetic model evaluation, Council contracts, tests and remaining limitations. Source links below are navigation aids into the maintained repository; use the frozen baseline to reproduce the original inspection.

The report audits source baseline a96025e; released v7 uses the source/image in config/releases/v7.json. Documentation and code can evolve after an edition is frozen. Record which source is being checked.

1. Distinguish the kinds of reproduction

Activity	Dependencies	Changes / external calls
Read source and archived JSON	Git checkout	None
Regenerate report figures	Python application requirements; optional CairoSVG + native Cairo for PNG	Writes documentation figures/data; no database, solver or network call
Offline provider/contracts tests	Locked Python requirements	In-process HTTP mocks; no real inference
Local numerical checks	Locked Python requirements	Real local solver; no inference; timings/feasible schedules can vary
PostgreSQL lifecycle/concurrency tests	Explicit dedicated local test databases	Writes isolated test state; never target a live farm database
Browser journeys	Built frontend, running dev service, Playwright Chromium	Real API/numerical work; some scripts intercept adviser fixtures; check each script
Public-data rebuild	Approved public source access or archived raw bytes	Live fetch creates new snapshots; offline needs raw files absent from Git
Authenticated DeepSeek trial	Explicit server credential, registered model access, finite budget	Metered calls; not needed for documentation reproduction
Edition publication	GitHub/Fly access and clean committed source	Updates hosting and immutable release registry; not part of this audit

2. Local application prerequisites

Use Python 3.13+, Node 24, PostgreSQL 18, the committed `requirements.lock.txt` and web lockfile. From repository root:

```
python -m venv .venv
.venv/bin/pip install -r requirements.lock.txt
npm ci --prefix apps/web
npm run build --prefix apps/web
```

Provision a **development** database and supply `FARMTACT_DATABASE_URL` through the process environment. `Store()` defaults to the Sprite-specific Unix socket `/tmp/farmtact-pg` and database `farmtact`; it does not create the PostgreSQL server or database. `.env.example` lists settings; the entrypoint does not automatically read a `.env` file. The application otherwise defaults to test execution and synthetic data; a key must not switch it to real farm operations.

```
.venv/bin/python scripts/initialize_database.py
.venv/bin/python -m packages.fixtures
.venv/bin/python scripts/build_dataset.py --with-power
.venv/bin/python scripts/build_features.py
```

The fixture command writes the fixture and shared schema. Feature/dataset commands write artifacts; do not mistake regeneration for a read-only check. Raw and normalized public files are ignored. `build_dataset.py --offline` requires already retrieved bytes and cannot reconstruct historical source contents from a manifest alone.

Outside Sprite, run `.venv/bin/python scripts/serve.py`. Inside Sprite, register it through `sprite-env services create` using the platform skill; do not create an unmanaged background server. For example, after provisioning the database service:

```
sprite-env services create farmtact-doc-dev \
  --cmd /home/sprite/workspaces/Code2/.venv/bin/python \
  --args /home/sprite/workspaces/Code2/scripts/serve.py \
  --http-port 8080
```

Use absolute paths appropriate to the actual checkout. This example does not provision the database or credentials. Inspect service CLI help for dependency registration and logs. The public Sprite URL must serve only the application, never repository roots, dumps or private files. No local HTTP service was required or started during this documentation audit.

3. Offline and numerical verification

These representative groups avoid a running PostgreSQL server and metered calls. The separate broader conversation test module has a known implicit D04 public-cache dependency

(DATA-06); its failed fresh-checkout assertion is preserved in the audit:

```
.venv/bin/python -m pytest -q \  
  tests/deepseek/test_gateway.py \  
  tests/contracts/test_private_data.py \  
  tests/contracts/test_features.py \  
  tests/models/test_explorer_numerics.py \  
  tests/seven_agents/test_council.py \  
  tests/review/test_provider_isolation.py  
  
.venv/bin/python -m pytest -q \  
  tests/numerical/test_planner.py \  
  tests/planner/test_research_constraints.py  
  
.venv/bin/python -m pytest -q \  
  tests/council_research/test_sessions.py \  
  tests/council_research/test_independent_review.py \  
  tests/gameplay/test_scenarios.py  
  
.venv/bin/python scripts/generate_web_contracts.py --check
```

The full `pytest -q` suite also includes PostgreSQL recovery/concurrency tests. Inspect each database fixture first. Research tests explicitly require a separate `farmtact_research_test` database; other historical test modules use their documented local test configuration and may recover queues. **Do not run them against a running application database.** SQLite tests establish local behavior but do not prove PostgreSQL lock/concurrency semantics.

`node tests/browser/run.mjs` is the original journey. Later features have separate scripts under `tests/browser/` and their release reports record invocation details. Install Chromium with `cd apps/web && npx playwright install chromium` when needed. An intercepted adviser fixture verifies UI behavior, not DeepSeek response quality.

`scripts/evaluate_numerical.py` rewrites `reports/numerical_evaluation.json`. Run it only when intentionally producing a new versioned evaluation artifact. Its four-second per-policy solver limit means exact allocations, bounds and timing are not guaranteed identical between environments. Preserve the archived report and compare constraints/semantics, not only JSON equality. The metadata gap DATA-03 must be fixed before treating regeneration as complete evaluation provenance.

4. Report figures and rendering

The committed SVG/PNG figures can be read without Python. To regenerate SVGs and source data from the fixture:

```
.venv/bin/python docs/technical/figures/generate.py
```

For PNGs, install the optional documentation tools in a separate environment using `requirements-docs.txt`, plus the native Cairo library and DejaVu fonts supplied by your operating system. These packages are not runtime application dependencies. Then run the generator with `--png`. The figure generator imports the locked application's `fixture/contracts/forecast` modules and makes no provider, source-ingestion, database or optimization request.

The figure data records the fixture hash, crop parameters, caixin observations and one-step EWMA values. Architecture arrows are authored annotations from source inspection. SVG text includes descriptions for accessibility. The figures were visually inspected as PNGs during the documentation audit.

The committed diagram manifest records the 20 rendered Mermaid SVGs and hashes of their chapter source blocks. To regenerate those images, use a separate documentation tool directory:

```
npm install --prefix /tmp/farmtact-doc-tools \
  @mermaid-js/mermaid-cli@11.17.0 mathjax-full@3.2.1
node /tmp/farmtact-doc-tools/node_modules/puppeteer/install.mjs
.venv/bin/python docs/technical/render_diagrams.py \
  --mmdc /tmp/farmtact-doc-tools/node_modules/.bin/mmdc
```

Chromium requires its native system libraries. A restricted local VM may require an explicit Puppeteer launch configuration; `--puppeteer-config` passes that file to Mermaid CLI. This audit used a local `--no-sandbox` browser solely for trusted authored documentation, with no public HTTP service. Mermaid parsing caught one sequence-label semicolon error, which was corrected before all 20 diagrams rendered.

Render the report with Python Markdown and the optional local MathJax module to produce a single printable HTML document:

```
.venv/bin/python docs/technical/render_report.py \
  --mathjax-module /tmp/farmtact-doc-tools/node_modules/mathjax-full \
  --output /tmp/farmtact-technical-report.html
python docs/technical/check_docs.py
```

The renderer embeds scientific figures, Mermaid SVGs, existing screenshot evidence and MathJax-rendered equations. It includes all seven chapters and makes no network request. Without `--mathjax-module`, TeX remains source notation; without rendered diagram files, Mermaid remains source notation. The committed reading copies use both renderers. Source/evidence links point to GitHub; there are no external scripts, fonts or CDN dependencies in the HTML. The PDF exporter can print that HTML using local Puppeteer and also checks image loading, internal links and page overflow. See the audit record for the actual invocation and limits.

5. Authenticated trials and archival discipline

The DeepSeek runbook describes the bounded real trial. - -live means authenticated transport, not operational farm mode. Never print a credential or include it in command arguments. Provider failure must remain blocked/failed; do not substitute another provider.

Keep separate records for source inspection, offline mocks, fresh authenticated calls and archived replay. A full suite pass does not validate crop biology or human decision benefit. The exact checks run for this audit, including counts, limitations and generated files, are in the audit record.